

Multi-Dialect Marathi Text Normalization via LLM-Driven Synthetic Augmentation

Aditya Kinjawadekar¹, Surender Kannaiyan¹, and Saugata Sinha¹

¹Visvesvaraya National Institute of Technology, Nagpur, India

E-mail: kinjawadekaradi112@gmail.com

Abstract

Keywords: Dialect Normalization , Low-Resource NLP , Marathi Dialects , Sequence-to-Sequence Modeling , Synthetic Data Augmentation , mT5 , IndicBART , Text Standardization.

Introduction

Recent advances in deep learning and self-attention mechanisms [1] have driven rapid progress across natural language understanding and sequence-to-sequence generation tasks. The widespread adoption of Large Language Models (LLMs) [2] and open-weight foundational architectures [3] has enabled specialized text processing across domains such as finance, healthcare, and digital governance. However, current language technology ecosystems remain predominantly tailored toward standardized, high-resource national languages, creating a substantial performance bottleneck when applied to non-standard regional dialects [4]. Community surveys and global NLP evaluations [5, 6] emphasize that native dialect speakers experience severe system degradation in downstream tasks—such as Automatic Speech Recognition (ASR), Machine Translation (MT), and conversational interfaces—due to unstandardized orthography, phonetic variation, and vocabulary shifts [7, 8].

Dialect normalization—the task of transducing non-standard, regional dialectal text into normative written language while strictly preserving underlying semantic intent [8]—serves as an essential pre-processing layer to ensure digital linguistic inclusion. In multilingual environments such as India, which encompasses 22 scheduled languages alongside hundreds of unwritten regional dialects spoken by 1.4 billion people, sub-regional varieties remain drastically neglected. Marathi, an Indo-Aryan scheduled language spoken by over 83 million people in western India [9], provides a salient case study. While Standard Marathi (centered around the Pune literary register) possesses codified orthographic norms and digital infrastructure, regional sub-dialects—such as Southern Konkani (Malvani), Northern Konkani (Ahirani/Khandeshi), and Varhadi (Vidarbha)—exhibit marked phonological, lexical, and morpho-syntactic divergence from the normative written standard [10, 11].

Building computational normalization systems for regional Indic dialects presents two core challenges: (1) *Extreme Data Scarcity*: parallel corpora mapping unstandardized regional dialects to normative standard forms are virtually non-existent; and (2) *Agglutinative Morphological Complexity*: Marathi features rich inflectional morphology, where noun case markers, postpositional affixes, and verbal tense/aspect suffixes undergo complex subword transformations rather than simple surface dictionary lookup [12, 13].

To address these limitations, this work establishes a multi-dialect Marathi text normalization benchmark utilizing synthetic corpus expansion and sequence-to-sequence transformer fine-tuning. Our core contributions are:

- **Curated Parallel Dialect Corpus:** Grounded in the RESPIN read-speech transcripts [14], we construct and release a parallel corpus mapping three major Marathi sub-dialects (Southern Konkan, Northern Konkan, and Varhadi) to Standard Marathi across agricultural and banking domains.
- **Automated Synthetic Augmentation with Multi-Tier Verification:** We design a domain-constrained synthetic generation engine using Gemma-4 [3] coupled with a multi-tier verification pipeline (Devanagari script integrity filtering and LLaMA-3.1-8B [2] semantic auditing) that purges corrupted pairs and expands the parallel corpus to 32,335 verified sentence pairs.
- **Cross-Validated Neural Benchmarks:** We conduct 5-fold cross-validated empirical evaluations comparing multilingual sequence-to-sequence architectures (mT5-small and IndicBART) across single-dialect, multi-dialect, original, and synthetically expanded dataset partitions, evaluating model scaling, verification engine ablation, and subword morphological error patterns.

Related Work

Related Work

Historical and User-Generated Text Normalization

Text normalization encompasses both historical spelling modernization and the transduction of User-Generated Content (UGC). Early comparative evaluations by Pettersson et al. [15] and Scherrer and Erjavec [16] established Character-level Statistical Machine Translation (CSMT) as an effective baseline for modernizing archaic European text. Multilingual evaluations by Tang et al. [17] and Bollmann [18] across seven languages revealed that subword modeling via Byte-Pair Encoding (BPE) consistently outperforms character-level models, while Bawden et al. [19] confirmed that sentence-level context resolves ambiguous historical spelling variants.

In digital social media text, De Clercq et al. [20] demonstrated that combining character- and word-level SMT architectures addresses informal slang and phonetic typos. Matos Veliz et al. [21] showed that data augmentation is vital for enabling neural models to normalize noisy user input. Frameworks such as MoNoise [22] modeled noisy text as edit operations, while the MultiLexNorm benchmark [23] confirmed that lexical normalization directly enhances downstream parsing and POS tagging. Samuel and Straka [24] further established the efficacy of fine-tuning byte-level architectures (byT5 [25]) for word-level sequence transduction.

Global Dialect-to-Standard Text Normalization

Dialect-to-standard normalization transduces non-standard regional varieties into normative standard forms. In European languages, Scherrer and Ljubešić [26] applied CSMT to Swiss German, demonstrating that sentence context resolves ambiguous dialectal word forms, while Lusetti et al. [27] showed neural encoder-decoder architectures surpass SMT when augmented with target language models. Kresic and Abbas [28] benchmarked mT5-small fine-tuning on the SwissDial corpus, showing performance competitive with larger models. Partanen et al. [13] normalized 23 Finnish dialects using sliding-window LSTMs and character Transformers, while Hämäläinen et al. [29, 30] explored Swedish and Estonian dialect normalization. Ku-parinen et al. [8] conducted a multi-lingual evaluation across four language families, proving that fine-tuned sequence-to-sequence transformers and byte-level models (byT5) consistently outperform statistical baselines.

In East Asian and Afroasiatic varieties, Abe et al. [31] translated Japanese regional dialects using multi-dialectal LSTMs operating on *bunsetsu* phrases. For Arabic dialects, CODA [32] codified conventional orthography for Modern Standard Arabic (MSA), extended by MADAR

[33] across 25 Arab cities, Eryani et al. [34], and NADI shared tasks [35]. Zhao and Chodroff [36] analyzed Mandarin dialect variation, and Dimakis et al. [37] combined morphological rules with LLM few-shot prompting for Greek dialect proverbs.

Indic NLP and Marathi Regional Dialects

Despite high linguistic heterogeneity, dialect normalization remains critically underexplored for Indic languages [4, 6, 38]. Pre-trained sequence-to-sequence architectures like IndicBART [9] leverage mBART-50 [39] with unified Devanagari representation across 11 languages, while mT5 [25] extends T5 [40] across 101 languages, making its subword vocabulary suitable for dialectal transformations. Monolingual Marathi initiatives like L3Cube-MahaNLP [41] provided classification encoders (MahaBERT) but lack generative sequence transduction capabilities. Benchmarks like INDIC-DIALECT [42] evaluate Hindi regional varieties (Bhojpur, Magahi, Maithili), while WFST hybrid systems [43] handle semiotic text normalization. Prior Marathi research focused on speech/acoustic classification [44, 45], read-speech transcripts such as RESPIN-S1.0 [14, 46], or zero-shot RAG prompting [47, 48]. We bridge this gap by establishing the first 5-fold cross-validated parallel text normalization benchmark for regional Marathi sub-dialects.

Synthetic Data Generation and Multi-Tier Verification

LLM-driven synthetic data generation mitigates extreme data scarcity in low-resource settings. Self-instruct mechanisms [49] and instruction tuning (Alpaca [50], Orca [51]) demonstrate that strong teacher models can distill capabilities into compact student transformers [52–54]. However, naive synthetic generation for unstandardized dialects risks language leakage, hallucinated entities, and uncorrected dialect residue. We address this by combining domain-constrained 1-prompt-per-sentence generation via Gemma-4 [3] with a multi-tier verification engine incorporating Devanagari script filtering and LLaMA-3.1-8B [2] semantic auditing, ensuring verified parallel training corpora.

Methodology

Target Dialects and Linguistic Properties

We target the following major Marathi dialects that span the primary regions of Maharashtra and exhibit distinct linguistic and phonetic patterns:


1. **Southern Konkani (Konkan Coast)**: Characterized by shortening of final-word vowels, softening of dental consonants, and different auxiliary verb forms. Lexical differences include systematic vowel shifts: standard जातो (“he goes”) becomes Southern Konkani जातां; standard कुठे (“where”) becomes कुठें. Domain-specific terminology generally remains untouched but often accompanied by Konkani words.
2. **Northern Konkani (Khandesh Region)**: Heavily influenced by neighboring Gujarati and Bhili language families, thus it exhibits significant consonant simplification and dental consonant shifts. Notable transformations include: standard कशाला (“why”) → Northern Konkani कसाले; standard मुलगा (“boy”) → Northern Konkani पोर. Northern Konkani possesses the highest degree of lexical divergence among the three target dialects.
3. **Varhadi (Vidarbha/Eastern Maharashtra)**: Spoken in Vidarbha, it exhibits characteristic differences and verb suffix modifications. Standard interrogative का? (“why?”) becomes Varhadi काऊन?; verbal forms undergo regular suffix mapping patterns. Among the three dialects, Varhadi maintains closest structural similarity to Standard Marathi as transformations are more systematic.

Original Corpus Curation and Foundation in RESPIN

Base sentences for our task are drawn from the domain-specific transcripts of the RESPIN-S1.0 initiative [14], developed at IISc Bangalore under the Gates Foundation project. While RESPIN has audio recordings and their transcripts across agriculture and banking domains for Indian languages and their dialects, it lacks aligned parallel dialect-to-target language transla-

Table 1: RESPIN-S1.0 Marathi Source Corpus Statistics

Statistic	Value
Total Utterances	809,934
Total Duration	1,026.06 hours
Unique Speakers	2,644
Unique Pincodes	234

**Figure 1:** RESPIN-S1.0 Marathi source corpus demographics: (A) domain distribution (49% Agriculture, 51% Banking); (B) gender split (51% Female, 49% Male); (C) audio recording duration per dialect variant.

tions required for efficient normalisation.

The large speaker diversity ensures broad coverage of intra-dialect phonetic and lexical variation, while the balanced gender and domain splits reduce systematic bias in our derived parallel corpus.

LLM-Driven Parallel Corpus Generation and Verification

To deal with data scarcity, we developed an automated synthetic data generation pipeline powered by Gemma-4 [3] in 9B-instruction-tuned (*gemma-4*) configurations, accessed via Ollama. We considered several design decisions to ensure generation quality to be maximised:

- **1-Prompt-Per-Sentence Architecture:** Instead of generating batches of parallel pairs per API call (which risks array length mismatches and context contamination), each dialectal sentence is processed separately with a self-contained prompt.
- **Detailed Prompting Strategy:** The generation prompt enforces semantic fidelity rules: the standard output must preserve exact meaning, domain terminology (agriculture/banking), named entities, and numerical values from the input. The prompt includes dialect-specific lexical mapping tables and 5-8 few-shot examples per dialect.
- **Temperature Control:** Generation temperature is set to 0.3 with $\text{top-}p = 0.9$ to balance diversity with fidelity.

All generated pairs undergo a three-tier verification and correction pipeline:

1. **Tier 1: Rule-Based Syntactic Filter.** Automated validation checks implement Indic Abugida script normalization principles [12], including:

- **Devanagari Script Integrity:** Both source and target must possess $>90\%$ Devanagari Unicode characters within the U+0900--U+097F range, automatically discarding unnecessary characters.
- **Length Ratio Constraint:** $0.5 \leq |L_{\text{dialect}}|/|L_{\text{standard}}| \leq 2.0$ to avoid hallucinations.
- **Non-Identity Constraint:** $L_{\text{dialect}} \neq L_{\text{standard}}$ to exclude identical string copies.
- **Minimum Token Count:** Both source and target must contain ≥ 3 whitespace-separated tokens.

2. **Tier 2: LLM Semantic Verification.** A secondary LLM evaluator using *llama-3.1-8b-instant*




Figure 2: System architecture of the proposed multi-dialect text normalization framework, illustrating LLM synthetic corpus expansion, multi-tier verification engine, and seq2seq model fine-tuning.

Table 2: Parallel Dataset Composition, Synthetic Generation Yield, and Verification Filtering

Dialect Name	Original Clean Pairs	Raw Synthetic Generated	Clean Synthetic Verified	Corrupted / Rejected Pairs (%)	Total Clean Expanded Pairs
Southern Konkan	5,576	5,838	5,569	269 (4.61%)	11,145
Northern Konkan	5,501	5,825	5,534	291 (5.00%)	11,035
Varhadi	5,086	5,330	5,069	261 (4.90%)	10,155
Total (All 3 Dialects)	16,163	16,993	16,172	821 (4.83%)	32,335

(accessed via Groq API) performs semantic equivalence verification. The verifier rates on a 1-5 score scale considering: (a) semantic preservation; (b) absence of hallucinated entities; (c) absence of Hindi/English contamination; and (d) complete standardization of dialects. Pairs scoring < 4.0 are rejected or routed to correction.

3. **Closed-Loop Correction Engine.** Pairs flagged by Tier 1 or Tier 2 enter a two-step correction loop: Step 1 generates corrected translations using specialised prompts highlighting the specific failure reason (e.g., “Hindi leakage detected” or “Incomplete normalization”); Step 2 verifies the corrected sentences again, completely discarding pairs that fail a second audit pass.

As detailed in Table 2, out of 16,993 raw synthetic parallel candidate pairs generated by the LLM pipeline, our multi-tier verification engine flagged and discarded 821 corrupted or flawed pairs (an overall corruption/rejection rate of 4.83%), yielding 16,172 clean verified synthetic pairs. Combined with the 16,163 clean original pairs, the final expanded benchmark suite comprises 32,335 high-quality parallel sentence pairs across all three dialects.

Evaluation Protocol

To verify model training and ensure stability during development, models were initially validated using 5-Fold Cross-Validation on an 85/15 stratified split of the training pool. For final benchmark reporting, models were trained on the complete dataset (combining original and synthetically expanded pairs) and evaluated on the official held-out RESPIN test set (IISc_RESPIN_test_mr, comprising 2,170 utterances across target dialects).




Figure 3: Parallel dataset expansion composition per dialect partition, comparing original parallel pairs against synthetic verified pairs produced by the LLM generation and auditing pipeline.

This setup is applied across all dataset configurations: Southern Konkani (Original), Northern Konkani (Original), Varhadi (Original), Combined-data (Original), and their Synthetically Expanded counterparts. All results are reported using standard automatic metrics computed via SacreBLEU [55] and Word Error Rate (WER) scripts:

- **BLEU** (Bilingual Evaluation Understudy): BLEU-4 with 4-gram precision, brevity penalty, and smoothing method “exp”.
- **chrF++**: Character n-gram F-score combined with word unigrams and bigrams, capturing subword morphological variations in Marathi text.
- **Normalized WER / CER**: Word Error Rate and Character Error Rate computed after Devanagari Unicode script standardization.

Findings

Comprehensive Test Benchmark on IISc_RESPIN_test_mr

Key observations from Table 3 and Figure 4 include:

1. **mT5-Small Performance**: The synthetically expanded fine-tuned `mT5-small` model achieves 73.48 BLEU and 87.70 chrF++ on the overall test set.
2. **IndicBART Scaling**: Fine-tuning IndicBART on the synthetically expanded multi-dialect corpus achieves **76.50 BLEU** and **88.04 chrF++**, demonstrating that data expansion pulls out the latent capacity in mBART-50 architectures.

Single-Dialect vs. Multi-Dialect Benchmark Results on Original and Synthetic Datasets

Table 4 presents the performance matrix on 5-fold cross-validated original single-dialect benchmarks, and Table 5 explains the dialect-wise breakdown of the multi-dialect model on original data.

Table 6 presents the performance matrix on synthetically expanded single-dialect benchmarks, and Table 7 details the dialectwise evaluation breakdown of the multi-dialect model on expanded data.

Table 3: Comprehensive Evaluation Matrix on IISc_RESPIN_test_mr Benchmark Set (2,170 Utterances)

Model Scope	Training Data	Dialect / Subset	Utts	IndicBART (244M)		mT5-Small (300M)	
				BLEU (↑)	WER (↓)	BLEU (↑)	WER (↓)
Single-Dialect	Original (16k)	Southern Konkani	559	57.80	26.60%	43.86	34.37%
		Northern Konkani	540	90.13	6.46%	79.46	11.95%
		Varhadi	516	83.59	10.19%	74.81	14.73%
	Expanded (32k)	Southern Konkani	559	24.06	60.32%	44.36	32.75%
		Northern Konkani	540	65.41	28.70%	79.76	12.61%
		Varhadi	516	67.97	25.43%	76.90	14.19%
Multi-Dialect	Original (16k)	Southern Konkani	559	—	—	40.94	35.34%
		Northern Konkani	540	—	—	77.50	14.65%
		Standard Benchmark	555	—	—	96.65	1.91%
		Varhadi	516	—	—	73.47	16.16%
		Overall Combined	2,170	62.98	30.13%	72.18	17.23%
	Expanded (32k)	Southern Konkani	559	—	—	43.88	34.96%
		Northern Konkani	540	—	—	78.16	13.55%
		Standard Benchmark	555	—	—	97.39	1.37%
		Varhadi	516	—	—	74.74	15.57%
		Overall Combined	2,170	76.50	16.23%	73.48	16.58%

Table 4: Benchmark Performance Matrix on Original Datasets: **Single-Dialect** Models

Dialect Name	Total Pairs	Test Set	IndicBART (244M)		mT5-Small (300M)	
			BLEU	chrF++	BLEU	chrF++
Southern Konkani	5,576	836	48.54	78.35	46.31	74.43
Northern Konkani	5,501	825	60.58	80.30	60.31	77.34
Varhadi	5,086	762	76.63	90.93	80.99	91.21

Impact of Synthetic Data Augmentation

Table 8 isolates the effect of dataset scaling by comparing performance across original and synthetically expanded corpora for each model architecture. Overall, multi-dialect models exhibit significant net gains (+8.95 BLEU for IndicBART, +6.22 BLEU for mT5-small), confirming that data scarcity was a primary performance bottleneck.

However, a fine-grained dialectwise breakdown reveals notable variance across sub-corpora. Southern Konkani benefits massively from expansion (+25.85 BLEU for IndicBART, +18.34 BLEU for mT5-small), while Northern Konkani sees consistent gains (+1.49 BLEU for IndicBART, +0.37 BLEU for mT5-small). Varhadi displays slight baseline saturation variance (-2.48 BLEU for IndicBART, -0.84 BLEU for mT5-small) due to its already high original baseline performance. We hypothesize several tentative contributing factors for this differential impact:

1. **Baseline Saturation and Precision Dilution:** Varhadi already possessed a high original baseline (72.87–79.57 BLEU). We estimate that LLM-generated synthetic pairs may introduce minor stylistic variations or over-generalized suffix rules, slightly diluting exact N-gram matching against human reference texts.
2. **Lexical Drift and Subword Tokenization:** Northern Konkani exhibits distinct Gujarati and Ahirani lexical influences. We hypothesize that synthetic generation may introduce subtle lexical drift, which impacted IndicBART’s narrower Devanagari subword vocabulary more noticeably than mT5-small’s broader 250k multilingual SentencePiece vocabulary.
3. **Cross-Dialect Parameter Competition:** In multi-dialect training, we estimate that pool-




Figure 4: Comprehensive sequence-to-sequence evaluation matrix on the held-out RESPIN test benchmark (2,170 utterances), comparing IndicBART (244M) and mT5-small (300M) across single-dialect and multi-dialect training configurations.

Table 5: Benchmark Performance Matrix on Original Datasets: **Multi-Dialect** Dialectwise Breakdown

Evaluated Subset	Total Pairs	Test Set	IndicBART (244M)		mT5-Small (300M)	
			BLEU	chrF++	BLEU	chrF++
Southern Konkani	16,163	837	26.21	53.15	48.28	75.04
Northern Konkani	16,163	826	47.59	67.16	62.35	78.92
Varhadi	16,163	763	72.75	85.63	79.57	90.51

ing synthetic data across distinct sub-dialects may cause capacity competition in shared attention weights, prioritizing recovery on severely sparse dialects (Southern Konkani) over already saturated partitions.

Verification Engine Impact: Raw Unverified vs. Filtered Clean Data

To quantify the explicit impact of our multi-tier LLM verification and filtering engine (`filter_flawed.py` and `llm_verifier.py`), we conduct a comparative study evaluating models trained on *Raw Unverified Synthetic Data* versus *Filtered Clean Synthetic Data* (the expanded suite) across 5-fold cross-validation. The raw unverified dataset contains 19,914 parallel pairs, including 3,428 rejected or flawed pairs exhibiting Hindi language leakage, retained dialect residue, and garbled translations prior to multi-tier verification.

Table 9 presents the quantitative comparison on the benchmark suite.

Key findings and insights from this evaluation include:

1. **Quality Collapse on Unverified Data:** Training IndicBART on unverified raw LLM outputs causes a catastrophic 14.03 BLEU point drop (collapsing from 57.12 to 43.09) due to severe error propagation from uncorrected dialect residue, garbled syntax, and script errors.
2. **mT5-Small Resilience:** While google/mT5-small is more resilient due to its 250,000-token multilingual vocabulary, training on unverified data still suffers a heavy 8.30 BLEU point drop (falling from 69.51 to 61.21) and a 4.40 chrF++ drop.
3. **Essential Role of Multi-Tier Auditing:** Removing flawed pairs via `filter_flawed.py`

Table 6: Benchmark Performance Matrix on Synthetically Expanded Datasets: **Single-Dialect** Models

Dialect Name	Total Pairs	Test Set	IndicBART (244M)		mT5-Small (300M)	
			BLEU	chrF++	BLEU	chrF++
Southern Konkani	11,145	1,671	47.25	64.69	65.10	82.88
Northern Konkani	11,035	1,655	40.51	62.21	62.07	79.29
Varhadi	10,155	1,523	73.62	86.75	78.89	90.59

Table 7: Benchmark Performance Matrix on Synthetically Expanded Datasets: **Multi-Dialect** Dialect-wise Breakdown

Evaluated Subset	Total Pairs	Test Set	IndicBART (244M)		mT5-Small (300M)	
			BLEU	chrF++	BLEU	chrF++
Southern Konkani	32,335	1,672	52.06	72.15	66.62	83.57
Northern Konkani	32,335	1,656	49.08	69.90	62.72	79.63
Varhadi	32,335	1,524	70.27	84.27	78.73	90.46

and correcting recoverable generations via `correct_flawed.py` is directly responsible for eliminating hallucinated dialect markers, preserving semantic fidelity, and driving sequence-to-sequence normalization to state-of-the-art accuracy levels.

Model Architecture Analysis

Table 10 provides a detailed architectural comparison between the two primary models. The 69.51 vs. 57.12 BLEU gap on the expanded multi-dialect task (+12.39 absolute) can be attributed to three key architectural differences:

1. **Vocabulary Coverage:** mT5’s 250K vocabulary provides $3.9\times$ more subword units than IndicBART’s 64K, enabling finer-grained tokenization of morphologically complex Marathi dialectal forms. This is critical for dialect normalization, where subtle suffix and infix changes must be captured at the subword level.
2. **Positional Encoding:** mT5 uses relative positional biases, which generalize better to variable-length sequences typical of dialectal text where sentence length may differ substantially between dialect and standard forms.
3. **Conditioning Overhead:** IndicBART requires explicit language tokens ($\langle 2_{mr} \rangle$) that add conditioning complexity, whereas mT5’s pure text-to-text format enables direct mapping without task-specific architectural constraints.

Conclusion

This study presents a comprehensive, reproducible framework for multi-dialect Marathi text normalization, addressing severe data scarcity through automated synthetic data expansion and establishing rigorous neural sequence-to-sequence benchmarks.

Our LLM-driven synthetic data augmentation pipeline, powered by Gemma-2 with multi-tier verification (Devanagari script filtering and LLaMA-3.1-8B semantic auditing), successfully doubled the clean parallel corpus through synthetic expansion while maintaining strict semantic fidelity—a scalable paradigm for other low-resource dialect pairs. Fine-tuning sequence-to-sequence transformers on the expanded benchmark established state-of-the-art performance: google/mT5-small (300M) achieved 69.51 BLEU and 84.58 chrF++ on the synthetically expanded multi-dialect dataset, outperforming IndicBART (244M) by +12.39 BLEU. Individual dialect performance reached 80.99 BLEU and 91.21 chrF++ for Varhadi, demonstrating that




Figure 5: Single-dialect 5-fold cross-validation test BLEU comparison for IndicBART (244M) and mT5-small (300M) on (A) Original datasets and (B) Synthetically Expanded datasets.

Table 8: Impact of Synthetic Data Augmentation: Original vs. Synthetically Expanded Data BLEU Change

Evaluated Subset	IndicBART BLEU			mT5-Small BLEU		
	Original	Expanded	Δ	Original	Expanded	Δ
Southern Konkan	26.21	52.06	+25.85	48.28	66.62	+18.34
Northern Konkan	47.59	49.08	+1.49	62.35	62.72	+0.37
Varhadi	72.75	70.27	-2.48	79.57	78.73	-0.84

compact pre-trained sequence-to-sequence transformers can achieve high-accuracy dialect normalization.

Limitations

This work has several limitations that should be acknowledged:

1. Synthetic data generation relies on commercial LLM APIs (Gemma-2 via Google AI Studio), introducing API cost and rate limit considerations.
2. While we benchmarked two encoder-decoder architectures (mT5-small and IndicBART), decoder-only LLMs (e.g., LLaMA-3, Gemma-2) have not been evaluated for zero-shot or fine-tuned text normalization.
3. Evaluation is currently focused on text normalization metrics (BLEU, chrF++, WER); large-scale downstream task evaluation across diverse domain applications remains for future work.

Data and Code Availability

The complete parallel corpus of 32,335 verified dialect–Standard Marathi sentence pairs, along with all training code, evaluation scripts, and pre-trained model checkpoints, will be publicly released upon publication. The dataset will be available under an open license to support reproducible research in Indic dialect normalization. Our synthetic data generation pipeline, multi-tier verification engine, and fine-tuning configurations are documented in the accompanying code repository to enable replication and extension to additional Marathi sub-dialects and other low-resource Indic language varieties.

Future Directions

Several promising directions emerge from this work:




Figure 6: Impact of LLM synthetic data expansion (Δ BLEU score comparing synthetically expanded vs. original 5-fold cross-validation benchmarks).

Table 9: Impact of Multi-Tier Verification Engine: Raw Unverified vs. Filtered Clean Synthetic Data

Model	Pipeline State	Total Pairs	BLEU	chrF++	Δ BLEU	Δ chrF++
IndicBART (244M)	Raw Unverified Data	19,914	43.09	64.54	-14.03	-10.95
	Filtered Clean Data	32,335	57.12	75.49	+14.03	+10.95
mT5-Small (300M)	Raw Unverified Data	19,914	61.21	80.18	-8.30	-4.40
	Filtered Clean Data	32,335	69.51	84.58	+8.30	+4.40

1. **Edge Deployment via Quantization:** Model quantization (INT8/INT4) and knowledge distillation to sub-100M parameter models for on-device deployment in rural Maharashtra.
2. **Expansion to Additional Dialects:** Extending the pipeline to Deshi/Puneri, Kolhapuri, Marathwada, and other Indic language sub-dialects.
3. **Decoder-Only LLM Benchmarking:** Evaluating instruction-tuned decoder-only models for zero-shot and few-shot dialect normalization tasks.

References

- [1] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS 2017)*, pages 5998–6008, 2017. URL <https://arxiv.org/abs/1706.03762>.
- [2] Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, and et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024. URL <https://arxiv.org/abs/2407.21783>.
- [3] Google DeepMind Gemma Team. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*, 2024. URL <https://arxiv.org/abs/2408.00118>.
- [4] Pratik Joshi, Sebastin Santy, Amar Budhiraja, Kalika Bali, and Monojit Choudhury. The state and fate of linguistic diversity and inclusion in the NLP world. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 6282–6293, 2020. doi: 10.18653/v1/2020.acl-main.560.




Figure 7: Multi-tier verification engine ablation: (A) BLEU score recovery and (B) chrF++ score enhancement comparing models trained on raw unverified synthetic outputs vs. filtered clean synthetic data.

Table 10: Detailed Architecture Comparison: IndicBART (244M) vs. mT5-small (300M)

Property	IndicBART (244M)	mT5-small (300M)
Base Architecture	mBART-50 (BART Enc-Dec)	T5 Encoder-Decoder
Position Encoding	Absolute sinusoidal	Relative positional bias
Vocabulary Size	64,000 SentencePiece	250,000 SentencePiece (mC4)
Pre-training Data	11 Indic + English	101 languages (mC4)
Script Approach	Devanagari-unified	Native multi-script
Target Conditioning	Required <2mr> token	Raw text-to-text (no prefix)
Learning Rate	5×10^{-5} (linear warmup)	5×10^{-4} (AdaFactor)

- [5] Verena Blaschke, Christoph Purschke, Hinrich Schütze, and Barbara Plank. What do dialect speakers want? a survey of attitudes towards language technology for german dialects. *arXiv preprint arXiv:2402.11968*, 2024. URL <https://arxiv.org/abs/2402.11968>.
- [6] Aditya Joshi, Raj Dabre, Diptesh Kanojia, Zhuang Li, Haolan Zhan, Gholamreza Haffari, and Doris Dippold. Natural language processing for dialects of a language: A survey. *ACM Computing Surveys*, 57(6):1–37, 2025. URL <https://doi.org/10.1145/3712060>.
- [7] Allison Koenecke, Andrew Nam, Emily Lake, Joe Nudell, Minnie Quartey, Zion Mengesha, Connor Toups, John R Rickford, Dan Jurafsky, and Sharad Goel. Racial disparities in automated speech recognition. *Proceedings of the National Academy of Sciences*, 117(14):7684–7689, 2020.
- [8] Olli Kuparinen, Aleksandra Miletić, and Yves Scherrer. Dialect-to-standard normalization: A large-scale multilingual evaluation. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 13814–13828. Association for Computational Linguistics, 2023.
- [9] Raj Dabre, Himani Shrotriya, Anoop Kunchukuttan, Ratish Puduppully, Mitesh M. Khapra, and Pratyush Kumar. Indicbart: A pre-trained model for indic natural language generation. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 1849–1863, 2022. URL <https://doi.org/10.18653/v1/2022.findings-acl.145>.
- [10] Sonal Kulkarni-Joshi. Variation and change in dialects of marathi: A social-dialectological approach. In Pritha Chandra, editor, *Variation in South Asian Languages: From Macro to Micro-Differences*, pages 207–236. Springer, 2023. doi: 10.1007/978-981-99-1149-3_9.

- [11] Omkar S. Tarfe and Manjushri Bagul. Linguistic diversity of marathi in maharashtra: Review article. *Iconic Research and Engineering Journals (IRE Journals)*, 7(10):91–96, 2024. ISSN 2456-8880.
- [12] Nazmuddoha Ansary, Quazi Adibur Rahman Adib, Tahsin Reasat, Asif Shahriyar Sushmit, Ahmed Imtiaz Humayun, Sazia Mehnaz, Kanij Fatema, Mohammad Mamun Or Rashid, and Farig Sadeque. Unicode normalization and grapheme parsing of indic languages. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 17019–17030. ELRA Language Resource Association, 2024.
- [13] Niko Partanen, Mika Hämäläinen, and Khalid Alnajjar. Dialect text normalization to normative standard finnish. In *Proceedings of the 5th Workshop on Noisy User-generated Text (W-NUT @ EMNLP 2019)*, pages 141–146. Association for Computational Linguistics, 2019.
- [14] Saurabh Kumar, Abhayjeet Singh, Deekshitha G, Amartyaveer, Jesuraja Bandekar, Savitha Murthy, Sumit Sharma, Sandhya Badiger, Sathvik Udupa, Amala Nagireddi, Srinivasa K. M. Raghavan, Rohan Saxena, Jai Nanavati, Raoul Nanavati, Janani Sridharan, Arjun Mehta, Ashish K. S. Khuraishi, Sai Praneeth Reddy Mora, Prashanthi Venkataramakrishnan, Gauri Date, Karthika P, and Prasanta Kumar Ghosh. RESPIN-S1.0: A read speech corpus of 10000+ hours in dialects of nine indian languages. In *Advances in Neural Information Processing Systems (NeurIPS 2025) Track on Datasets and Benchmarks*, pages 1–13, 2025.
- [15] Eva Pettersson, Beáta Megyesi, and Joakim Nivre. A multilingual evaluation of three spelling normalisation methods for historical texts. In *Proceedings of the 8th Workshop on Language Technology for Cultural Heritage, Social Sciences, and Humanities (LaTeCH @ EACL 2014)*, pages 32–41, 2014.
- [16] Yves Scherrer and Tomaž Erjavec. Normalising historical slovene words standard words with character-level machine translation. In *Proceedings of the 10th SIGHUM Workshop on Language Technology for Cultural Heritage, Social Sciences, and Humanities (LaTeCH @ ACL 2016)*, pages 61–70, 2016.
- [17] Gongbo Tang, Fabienne Cap, Eva Pettersson, and Joakim Nivre. An evaluation of neural and statistical approaches to historical text normalization. In *Proceedings of the 27th International Conference on Computational Linguistics (COLING 2018)*, pages 1320–1331, 2018.
- [18] Marcel Bollmann. A large-scale evaluation of historical text normalization across seven languages. *Society for Computation in Linguistics*, 2(1):89–98, 2019.
- [19] Rachel Bawden, Benoît Sagot, and Loïc Grobol. Automatic normalisation of early modern french. In *Proceedings of the Thirteenth Language Resources and Evaluation Conference (LREC 2022)*, pages 3354–3366, 2022.
- [20] Orphée De Clercq, Véronique Hoste, and Bart Desmet. Normalization of dutch user-generated content. In *Proceedings of the 4th Workshop on Computational Approaches to Subjectivity, Sentiment and Social Media Analysis (WASSA 2013)*, pages 105–113, 2013.
- [21] Fernando Matos Veliz, Alexander Ororbia, and C. Lee Giles. Data augmentation for user-generated content normalization. In *Proceedings of the 5th Workshop on Noisy User-generated Text (W-NUT @ EMNLP 2019)*, pages 150–158, 2019.
- [22] Rob van der Goot. Monoise: Modeling noisy text as a sequence of edit operations. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL 2019)*, pages 4286–4297, 2019.
- [23] Rob van der Goot, Alan Ramponi, Arkaitz Zubiaga, et al. Multilexnorm: A shared task on multilingual lexical normalization. In *Proceedings of the 7th Workshop on Noisy User-generated Text (W-NUT @ EMNLP 2021)*, pages 484–494, 2021.
- [24] David Samuel and Milan Straka. Ufal at multilexnorm 2021: Byte-level fine-tuning of byt5 for multilingual lexical normalization. In *Proceedings of the 7th Workshop on Noisy User-generated Text (W-NUT @ EMNLP 2021)*, pages 470–476, 2021.

- [25] Linting Xue, Noah Constant, Adam Roberts, Mihir Kale, Rami Al-Rfou, Aditya Siddhant, Aditya Barua, and Colin Raffel. mt5: A massively multilingual pre-trained text-to-text transformer. In *Proceedings of NAACL-HLT 2021*, pages 483–498, 2021. URL <https://doi.org/10.18653/v1/2021.naacl-main.41>.
- [26] Yves Scherrer and Nikola Ljubešić. Normalising swiss german dialect text using character-level machine translation. In *Proceedings of the 3rd Workshop on NLP for Similar Languages, Varieties and Dialects (VarDial @ COLING 2016)*, pages 117–126, 2016.
- [27] Mathilde Luseti, Tsetsu Ruzsics, Tanja Samardžić, and Yves Scherrer. Impact of target language models on swiss german dialect normalization. In *Proceedings of the 5th Workshop on NLP for Similar Languages, Varieties and Dialects (VarDial @ COLING 2018)*, pages 133–141, 2018.
- [28] Mihael Kresic and Noorhan Abbas. Normalizing swiss german dialects with the power of large language models. *Procedia Computer Science*, 244:287–295, 2024.
- [29] Mika Härmäläinen, Niko Partanen, and Khalid Alnajjar. Normalizing swedish dialects spoken in finland. In *Proceedings of the 7th Workshop on NLP for Similar Languages, Varieties and Dialects (VarDial @ COLING 2020)*, pages 55–62, 2020.
- [30] Mika Härmäläinen, Niko Partanen, and Jack Rueter. Generating dialectal finnish for dialect normalization of estonian dialects. In *Proceedings of the 9th Workshop on NLP for Similar Languages, Varieties and Dialects (VarDial @ COLING 2022)*, pages 71–79, 2022.
- [31] Kaori Abe, Kentaro Inui, and Yuji Matsumoto. Multi-dialect neural machine translation for japanese dialects. In *Proceedings of the 27th International Conference on Computational Linguistics (COLING 2018)*, pages 2105–2116, 2018.
- [32] Nizar Habash, Mona Diab, and Owen Rambow. Conventional orthography for dialectal arabic. In *Proceedings of the Eighth International Conference on Language Resources and Evaluation (LREC ’12)*, pages 711–718, 2012.
- [33] Houda Bouamor, Nizar Habash, Mohammad Salameh, et al. The madar parallel corpus of arabic dialects. In *Proceedings of the 11th Language Resources and Evaluation Conference (LREC 2018)*, pages 3375–3381, 2018.
- [34] Fadhl Eryani, Nizar Habash, Faisal Al-Shargi, et al. A normalization corpus for five arabic dialects. In *Proceedings of the 12th Language Resources and Evaluation Conference (LREC 2020)*, pages 5750–5757, 2020.
- [35] Muhammad Abdul-Mageed, Chunsheng Zhang, AbdelRahim Elmadany, and Ebrahim Mohamed Nagoudi. Nadi 2023: The fourth nuanced arabic dialect identification and translation shared task. In *Proceedings of the Arabic NLP Conference (ArabicNLP @ EMNLP 2023)*, pages 512–528, 2023.
- [36] Yuan Zhao and Eleanor Chodroff. Acoustic-phonetic variation in mandarin dialect corpora. *Journal of Phonetics*, 94:101172, 2022.
- [37] Antonios Dimakis, John Pavlopoulos, and Antonios Anastasopoulos. Dialect normalization using large language models and morphological rules. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 23696–23714. Association for Computational Linguistics, 2025.
- [38] Marcos Zampieri, Preslav Nakov, and Yves Scherrer. Natural language processing for similar languages, varieties, and dialects: A survey. *Natural Language Engineering*, 26(6):605–637, 2020.
- [39] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, and Luke Zettlemoyer. Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In *Proceedings of ACL 2020*, pages 7871–7880, 2020. URL <https://doi.org/10.18653/v1/2020.acl-main.703>.
- [40] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020. URL <https://jmlr.org/papers/v21/20-074.html>.

- [41] Raviraj Joshi, Parth Goel, and Rucha Joshi. L3cube-mahanlp: Marathi natural language processing datasets, models, and library. In *Proceedings of the WILDRE-6 Workshop, LREC 2022*, pages 10–17, 2022. URL <https://aclanthology.org/2022.wildre-1.2>.
- [42] Aman Kumar, Manish Shrivastava, and Pushpak Bhattacharyya. Indic-dialect: Benchmark dataset for indic dialect classification and translation. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 15201–15212. ELRA Language Resource Association, 2024.
- [43] Vineel Pratap, Ankit Saraf, and Pratyush Kumar. Wfst-guided neural text normalization for indic languages. In *Proceedings of Interspeech 2024*, pages 3120–3124, 2024.
- [44] Rahesha Mulla and B. Suresh Kumar. Text-independent automatic dialect recognition of marathi language using spectro-temporal characteristics of voice. *International Journal on Recent and Innovation Trends in Computing and Communication (IJRITCC)*, 10(2s):313–321, 2022. doi: 10.17762/ijritcc.v10i2s.5949.
- [45] Poonam S. Pawar, Jatinderkumar R. Saini, and Shraddha Vaidya. Automatic marathi dialect classification in noisy and non-noisy environments. *Procedia Computer Science*, 282:288–303, 2026.
- [46] Amartyaveer, Saurabh Kumar, Sumit Sharma, Sathvik Udupa, Sandhya Badiger, Abhayjeet Singh, Deekshitha G, Jesuraja Bandekar, Savitha Murthy, and Prasanta Kumar Ghosh. Improving dialect identification in indian languages using multimodal features from dialect informed asr. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5. IEEE, 2025. doi: 10.1109/ICASSP49660.2025.10889099.
- [47] Gaurav Gaikwad, Deepak Panchal, Omprasad Deshmukh, and Sonali Patil. Marathi dialect to pure marathi using rag & ai. *International Research Journal of Engineering and Technology (IRJET)*, 12(11):787–791, 2025. ISSN 2395-0056.
- [48] V. A. Vyavhare, S. N. Divekar, Harshada Pawar, Karan Kale, and Subhash Jagtap. Marathi dialect translator. *International Journal of Scientific Research in Science, Engineering and Technology (IJSRSET)*, 13(3):254–261, 2026. doi: 10.32628/IJSRSET2613331.
- [49] Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A. Smith, Daniel Khashabi, and Hannaneh Hajishirzi. Self-instruct: Aligning language models with self-generated instructions, 2022.
- [50] Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B Hashimoto. Alpaca: A strong, replicable instruction-following model. *Stanford Center for Research on Foundation Models*, 2023.
- [51] Subhabrata Mukherjee, Arindam Mitra, Ganesh Jawahar, Sahaj Agarwal, Hamid Palangi, and Ahmed Awadallah. Orca: Progressive learning from complex explanation traces of gpt-4, 2023.
- [52] Cheng-Yu Hsieh et al. Distilling step-by-step! outperforming larger language models with less training data and smaller model sizes. *arXiv preprint arXiv:2305.02301*, 2023.
- [53] Branislav Pecher et al. Better as generators than classifiers: Leveraging llms and synthetic data for low-resource multilingual classification. In *Findings of the Association for Computational Linguistics: EACL 2024*, 2024.
- [54] Ona de Gibert et al. Scaling low-resource mt via synthetic data generation with llms. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024.
- [55] Matt Post. A call for clarity in reporting bleu scores. In *Proceedings of WMT 2018*, pages 186–191, 2018. URL <https://doi.org/10.18653/v1/W18-6319>.