

From Regional Dialects to Standard Written Marathi: A Normalization pipeline based on sequence-to-sequence fine-tuning

Aditya Kinjawadekar¹, Surender Kannaiyan¹, and Saugata Sinha¹

¹Visvesvaraya National Institute of Technology, Nagpur, India

E-mail: kinjawadekaradi112@gmail.com

Abstract

Scarcity of resources and absence of benchmarks for regional dialects of Indic languages is an established problem. In this work, we present an end-to-end neural fine-tuning pipeline for normalization of dialectal transcripts from the Marathi dialects of Southern Konkani, Northern Konkani, and Varhadi into Standard Marathi. To address the issues regarding data, we introduce a synthetic data generation pipeline with a multi-tier verification engine for generating the aligned corpus and expanding it further. We evaluate IndicBART (244M) and mT5-small (300M) across varied configurations. Experimental metrics on the held-out test benchmark show that mT5-small achieves 65.59 BLEU across the regional dialects under multi-dialect training, outperforming IndicBART by 6.77 BLEU, while IndicBART achieves its peak multi-dialect gain of +26.07 BLEU on Southern Konkani via synthetic expansion. Also, the multi-tier verification prevents the quality drop (+15.73 BLEU recovery over unverified data), suggesting the importance of verification of data during synthetic augmentation for dialect normalization.

Keywords: Dialect Normalization, Low-Resource NLP, Marathi Dialects, Sequence-to-Sequence Modelling, Synthetic Data Augmentation, mT5, IndicBART.

1 Introduction

Rapid development of deep learning architectures and self-attention mechanisms (Vaswani et al., 2017) has made it possible to develop systems capable of understanding and generating natural language. Open-weight models have gained popularity to process text in various domains as they provide a finer control over system architecture, runtime environment, and training data (Dubey et al., 2024). Yet current language systems tend to work efficiently with standard high-resource languages. Queries based on non-standard input yield unexpected and undesired outputs (Joshi et al., 2020). Speakers of regional dialects face many issues in systems that use these models across automated speech recognition, machine translation and other language processing systems (Koenecke et al., 2020; Zampieri et al., 2020).

Dialect normalization is the task of translating non-standard regional dialects into standard written language. The main challenge to be addressed while designing such systems is the preservation of the meaning of original text (Blaschke et al., 2024). It serves to be an essential layer for digital linguistic inclusion. In countries having large and diverse speech communities, such as India, which has 22 official languages and hundreds of regional dialects spoken by 1.4 billion people, regional varieties have remained unattended in computational linguistic research. An example of such linguistic disparity is Marathi, an Indo-Aryan language spoken by over 83 million people in western India (Kulkarni-Joshi, 2023). Development of reliable pipelines for such regional Indic dialects has faced these fundamental challenges:

- *Resource Scarcity:* Parallel corpora mapping the regional dialects to standard written language do not exist. Curation of such sets by experts fluent in multiple dialects is both expensive and cannot be scaled (Blaschke et al., 2024; Dhasmana et al., 2026).
- *Grammatical Variance:* Marathi exhibits rich morphology where different words undergo

varied complex transformations across different regions where the dialects are spoken (Ansary et al., 2024)(Joshi et al., 2025).

- *Degraded performance of Multilingual Models*: Performance of pre-trained models like IndicBART (Dabre et al., 2022), IndicNLP Suite (Kakwani et al., 2020), and mT5 (Xue et al., 2021) seems to drop drastically on non-standard language variants due to fragmentation during tokenisation and out-of-vocabulary phonological shifts (Joshi et al., 2020; Kumar et al., 2024).

In this work, we investigate the normalization performance of IndicBART and mT5 across regional Marathi dialects and evaluate how efficiently they adapt to dialect-to-standard mapping under task-specific conditioning, and check whether synthetic data can solve severe data scarcity while maintaining considerable performance.

Our core contributions are:

- *Parallel Corpus Generation*: Based on transcripts from the RESPIN dataset (Kumar et al., 2025), we construct a parallel corpus for three of the Marathi dialects (Southern Konkan, Northern Konkan, and Varhadi) mapped to Standard Marathi.
- *Automated Synthetic Augmentation with Verification*: We design a data expansion pipeline paired with a multi-tier verification engine, which has rule-based script filtration and LLM-based semantic inspection.

2 Related Work

2.1 Regional Dialect Standardization across Language Families

Early approaches in the context of European languages used Character-Level Statistical Machine Translation (CSMT) and phrase-based models for Swiss German Dialects (Scherer and Ljubešić, 2016) which were later replaced by pre-trained language models (Lusetti et al., 2018; Kresic and Abbas, 2024). Similar approaches have been developed for Finnish, Swedish, and Estonian dialects (Partanen et al., 2019; Hämäläinen et al., 2020, 2022). A study of four language families by Kuparinen et al. (Kuparinen et al., 2023) suggested that fine-tuned sequence-to-sequence transformers surpassed the rule-based and SMT baselines. Other than European countries, East Asian and African-Asian varieties have also shown similar patterns. Japanese regional dialects have been processed using phrase-level LSTM modules (Abe et al., 2018), while Mandarin scripts relied on acoustic-phonetic mapping (Zhao and Chodroff, 2022). The Conventional Orthography for Dialectal Arabic (CODA) framework (Habash et al., 2012) in the landscape of Arabic languages gave guidelines for Modern Standard Arabic (MSA) which led to multi-city translation benchmarks (MADAR (Bouamor et al., 2018), NADI (Abdul-Mageed et al., 2023)) and morphology-oriented prompt design for LLMs (Eryani et al., 2020; Dimakis et al., 2025).

2.2 Indic Dialect Landscape and Marathi NLP

Even though pre-trained models such as IndicBART (Dabre et al., 2022) and mT5 (Xue et al., 2021), which are trained on multilingual corpora cover major languages, they lack understanding of regional sub-dialects. Models based on a single language like L3Cube-MahaNLP (Joshi et al., 2022) offer pre-trained Marathi models for classification, but do not support generative text standardization. INDIC-DIALECT (Kumar et al., 2024) considers Hindi varieties (Bhojpuri, Magahi, Maithili), for evaluation via WFST frameworks for text normalization while preserving the underlying semantics (Pratap et al., 2024). However, Marathi dialects were primarily dealt for acoustic/ASR classification (Mulla and Kumar, 2022; Pawar et al., 2026), or zero-shot LLM retrieval (Gaikwad et al., 2025; Vyavhare et al., 2026). Our work fills this gap by presenting an end-to-end neural fine-tuning pipeline with evaluations and data augmentation strategies.

2.3 LLM Distillation and Multi-Tier Verification

LLMs can be used to suffice for data scarcity in various tasks (Wang et al., 2022; Taori et al., 2023; Mukherjee et al., 2023). Knowledge distillation strategies can be employed to train compact and efficient models based on the outputs generated by a larger teacher model (Hsieh et al., 2023; Pecher et al., 2024; de Gibert et al., 2024). But the use of such techniques for low-resource regional dialects involves significant amount of hallucinations and noise (de Gibert et al., 2024; Ji et al., 2023). We solve this by utilising Gemma-4 (Google DeepMind Gemma Team, 2026) with a robust morphologically aware system prompt and a multi-tier verification system that ensures integrity of the data generated.

3 Methodology

3.1 Original Corpus Curation and Foundation in RESPIN

Base sentences for our task are drawn from the domain-specific transcripts of the RESPIN-S1.0 initiative (Kumar et al., 2025), developed at IISc Bangalore under the Gates Foundation project. While RESPIN has audio recordings and their transcripts across agriculture and banking domains for Indian languages and their dialects, it lacks aligned parallel dialect-to-target language translations required for efficient normalisation. The large speaker diversity keeps a sufficiently diverse coverage of intra-dialect variations, while the balanced domain splits reduce systematic bias in our derived parallel corpus. Detailed corpus statistics are summarized in Table 1 and Figure 1.

Table 1: RESPIN-S1.0 Marathi Source Dialect Corpus Summary

Code	Dialect Name	Utterances (%)	Duration (h)	Speakers	Prompts
D1	Southern Konkan (Malvani)	203,651 (33.88%)	247.31	732	5,838
D2	Northern Konkan (Ahirani)	198,560 (33.03%)	265.25	557	5,825
D4	Varhadi (Vidarbha)	198,873 (33.09%)	257.70	747	5,330
Total	Source Dialects	601,084 (100%)	770.26	2,036	16,993




Figure 1: Corpus demographics: (A) domain split, (B) gender distribution, and (C) audio duration per dialect.

3.2 Target Dialects and Linguistic Properties

We target the following Marathi dialects based on the RESPIN corpus:

- **Southern Konkani:** Known for its shortening of final-word vowels, softening of dental consonants, and different auxiliary verb forms. Differences include systematic vowel shifts: standard जातो ("he goes") becomes Southern Konkani जातां; standard कुठे ("where") becomes कुठेंय.
- **Northern Konkani:** Heavily influenced by neighbouring Gujarati and Bhili language families, thus it has consonant simplification and dental consonant shifts. Some variations include: standard कशाला ("why") → Northern Konkani कसाले; standard मुलगा ("boy") → Northern Konkani पोर.
- **Varhadi:** Spoken in the regions of Vidarbha, it shows very specific verb suffix modifications. Standard interrogative का? ("why?") becomes Varhadi काऊन?.

3.3 Synthetic Parallel Data Generation Pipeline

To deal with data scarcity, we developed an automated synthetic data translation and generation pipeline powered by Gemma-4 (Google DeepMind Gemma Team, 2026) with 12B parameters configuration, accessed via Ollama. The complete end-to-end normalization architecture is illustrated in Figure 2. We considered several design decisions to ensure generation quality to be maximised:

- **1-Call-Per-Sentence Generation:** We observed batch generation caused hallucinations in context. Therefore, a single API call was utilized for a single sentence, flushing the KV-cache at every instant that gave fresh context to the LLM every time.
- **Preservation of Meaning:** System prompts were designed in a way that actually instructed the model that while the dialectal text is to be translated into standard written Marathi, the standard output must preserve the exact semantic meaning, domain-specific terminology (agriculture/banking), and proper nouns without any change.




Figure 2: Dialect normalization pipeline: synthetic data generation, two-tier verification, and model training.

3.4 Verification and Quality Filtering Engine

LLMs can generate flawed target sequences when dealing with regional sub-dialects. We implemented a two-tier verification engine to remove them before model training:

- **Tier 1: Rule-Based Syntactic Filter.** Automated validation checks implement Indic Abugida script normalization principles (Ansary et al., 2024), including:
 - **Script Consistency:** Discards samples containing non-Devanagari characters.
 - **Length Ratio Constraint:** Discards pairs where sentence length ratio $\frac{\text{len}(\text{Target})}{\text{len}(\text{Source})} \notin [0.5, 2.0]$.
 - **Minimum Token Count:** Both source and target must contain ≥ 3 whitespace-separated tokens.
- **Tier 2: LLM Semantic Verification.** A secondary LLM evaluator using `llama-3.1-8b-instant` (Dubey et al., 2024) evaluates the accuracy on a scale of 1-5 under redesigned system prompts. Pairs with semantic scores < 3 are discarded.

Through this two-step process, 3,751 low-quality sentence pairs were completely discarded from the corpus (19,914 pairs), preventing further error propagation. The same synthetic generation and verification pipeline was again used to expand the parallel dataset from 16,163 original pairs to 32,335 verified clean pairs across all the dialects. The dataset composition across final clean partitions is displayed in Figure 3.




Figure 3: Dataset sizes per dialect before and after synthetic expansion.

3.5 Model Fine-Tuning and Conditioning

We evaluate two pre-trained sequence-to-sequence models across different configurations: IndicBART (244M) (Dabre et al., 2022) and mT5-small (300M) (Xue et al., 2021).

To enable unified multi-dialect training where a single model is used for all regional variants, we design special task-based conditioning prefixes:

- **Task Normalization Prefixes:** A normalisation tag (`<normalize>`) is attached that instructs the encoder to perform the dialect-to-standard text transformation.
- **Dialect-Source Identification Tags:** Source sentences have dialect markers, `<D1>` for Southern Konkan, `<D2>` for Northern Konkan, and `<D4>` for Varhadi for explicit conditioning of the models in multi-dialect training and testing scenarios.
- **Target Standardization Tags:** Decoder target is conditioned using standard language tags (`<standard_mr>`), enforcing the output alignment into Standard Written Marathi.

Table 2 considers the underlying properties based on the architecture and hyperparameter configurations for the fine-tuning of IndicBART (244M) and mT5-small (300M).

Table 2: Architecture comparison: IndicBART (244M) vs. mT5-small (300M)

Property	IndicBART (244M)	mT5-small (300M)
Base Architecture	mBART-50 (BART Enc-Dec)	T5 Encoder-Decoder
Position Encoding	Absolute sinusoidal	Relative positional bias
Vocabulary Size	64,000 SentencePiece	250,000 SentencePiece (mC4)
Pre-training Data	11 Indic + English	101 languages (mC4)
Script Approach	Devanagari-unified	Native multi-script
Target Conditioning	Required <code><2mr></code> token	Raw text-to-text (no prefix)
Learning Rate	5×10^{-5} (linear warmup)	5×10^{-4} (AdaFactor)

There is a fundamental trade-off present in the models. IndicBART is tailored for Indic language scripts with a unified 64k Devanagari vocabulary and language-specific target tokens (`<2mr>`), providing stronger performance in single-dialect tasks. However, mT5-small features a substantially larger 250k multilingual vocabulary which reduces subword fragmentation on

rare dialectal forms, and relative positional embeddings capture local suffix and verbal shifts efficiently. These inherent structural advantages allow mT5-small to scale its performance in multi-dialect scenarios.

3.6 Evaluation Protocol

We report the model performance across the following two evaluation settings:

- *Mean 5-Fold Cross-Validation*: To evaluate the stability of models during training, they are trained and evaluated under an 85/15 stratified 5-fold cross-validation scheme.
- *Held-Out RESPIN Test Benchmark*: To consider the performance on unseen data, models were trained on the complete dataset and later evaluated on the official held-out set (IISc_RESPIN_test_mr), consisting of 1,615 unseen dialectal utterances across the three target regional dialects (Southern Konkani, Northern Konkani, and Varhadi).

Results are reported using standard text normalization evaluation metrics via SacreBLEU (Post, 2018):

- *BLEU* (Bilingual Evaluation Understudy): BLEU-4 with 4-gram precision, brevity penalty, with an exponential smoothing method.
- *chrF++*: Character n-gram F-score combined with word unigrams and bigrams, capturing subword morphological variations in Devanagari.

4 Findings

4.1 Single-Dialect vs. Multi-Dialect 5-Fold Cross-Validation Results

Cross-validation performance on original single-dialect datasets is reported in Table 3.

Table 3: Single-dialect model performance on original datasets (5-Fold CV)

Dialect Name	IndicBART (244M)		mT5-Small (300M)	
	BLEU	chrF++	BLEU	chrF++
Southern Konkani (D1)	48.54	78.35	46.31	74.43
Northern Konkani (D2)	60.58	80.30	60.31	77.34
Varhadi (D4)	76.63	90.93	81.00	91.21
Overall Combined	61.92	83.19	62.54	80.99

Baseline performance based on Table 3 highlights the structural differences in Marathi dialects. Both models achieve their highest scores on Varhadi (76.63-81.00 BLEU and 90.93-91.21 chrF++), whereas Southern and Northern Konkani have shown lower baseline scores (46.31-60.58 BLEU). This indicates the higher order of overlap between Varhadi dialect with standard Marathi when it comes to grammar and lexicography. Table 4 presents the dialect-wise breakdown of metrics when training on multi-dialect data.

Table 4: Multi-dialect model performance on original datasets (5-Fold CV)

Evaluated Subset	IndicBART (244M)		mT5-Small (300M)	
	BLEU	chrF++	BLEU	chrF++
Southern Konkani (D1)	26.21	53.15	48.28	75.04
Northern Konkani (D2)	47.59	67.16	62.35	78.92
Varhadi (D4)	72.75	85.63	79.57	90.51
Overall Combined	48.85	68.65	63.40	81.49

Combined training showed a drop in IndicBART metrics on Southern Konkani, from 48.54 to 26.21 BLEU (and chrF++ fell to 53.15), possibly due to cross-dialectal interference. But mT5-small retained performance, maintaining 48.28 BLEU on Southern Konkani and achieving

an overall average score of 63.40 BLEU and 81.49 chrF++. Table 5 evaluates single-dialect models trained on synthetically expanded corpora.

Table 5: Single-dialect model performance on expanded datasets (5-Fold CV)

Dialect Name	IndicBART (244M)		mT5-Small (300M)	
	BLEU	chrF++	BLEU	chrF++
Southern Konkani (D1)	47.25	64.69	65.10	82.88
Northern Konkani (D2)	40.51	62.21	62.07	79.29
Varhadi (D4)	73.62	86.75	78.89	90.59
Overall Combined	53.79	71.22	68.69	84.25

IndicBART experiences slight saturation of performance across single-dialect configurations with a score of 47.25 BLEU on Southern Konkani and 40.51 BLEU on Northern Konkani, while mT5-small has utilised the synthetic data more effectively, spiking the scores to 65.10 BLEU (+18.79) on Southern Konkani and 62.07 BLEU on Northern Konkani with chrF++ exceeding 79.29-90.59. Table 6 presents the evaluation of the multi-dialect model trained on the synthetically expanded corpus.

Table 6: Multi-dialect model performance on expanded datasets (5-Fold CV)

Evaluated Subset	IndicBART (244M)		mT5-Small (300M)	
	BLEU	chrF++	BLEU	chrF++
Southern Konkani (D1)	52.06	72.15	66.62	83.57
Northern Konkani (D2)	49.08	69.90	62.72	79.63
Varhadi (D4)	70.27	84.27	78.73	90.46
Overall Combined	57.14	75.44	69.36	84.55

IndicBART recovers performance from the cross-dialect interference to score 52.06 BLEU on Southern Konkani and 57.14 overall BLEU, with a +25.85 BLEU gain on Southern Konkani over the original multi-dialect configuration. mT5-small reaches peak performance across all dialects (66.62 BLEU on Southern Konkani, 62.72 BLEU on Northern Konkani, and an overall average of 69.36 BLEU and 84.55 chrF++). Comparative single-dialect 5-fold cross-validation performance shifts are illustrated in Figure 4.




Figure 4: Single-dialect cross-validation scores before and after synthetic expansion.

The multi-dialect performance matrix is shown in Figure 5.




Figure 5: Multi-dialect cross-validation performance across dialects and models.

4.2 Comprehensive Benchmark on the Held-Out Test Set

Subsequent evaluations are conducted on the official held-out set after training on the complete datasets without cross-validation. Single-dialect training shows that synthetic data augmentation substantially reduces IndicBART performance across all three dialects, while mT5-small remains comparatively stable and shows slight improvements on Northern Konkani and Varhadi. Multi-dialect training produces substantial gains for IndicBART on Southern Konkani (+26.07 BLEU) and Northern Konkani (+6.43 BLEU), while Varhadi decreases slightly (−3.08 BLEU), yielding an overall multi-dialect gain of +9.81 BLEU (from 49.01 to 58.82 BLEU). mT5-small shows consistently positive gains across all three dialects (+2.94, +0.66, and +1.27 BLEU for Southern Konkani, Northern Konkani, and Varhadi, respectively), reaching an overall multi-dialect score of 65.59 BLEU. The full performance evaluation on the held-out test set is shown in Table 7 and Figure 6.

Table 7: Model performance on the held-out RESPIN test set (1,615 Dialectal Utterances)

Model Scope	Training Data	Dialect / Subset	IndicBART BLEU	mT5-Small BLEU
Single	Original	Southern Konkani (D1)	57.80	43.86
		Northern Konkani (D2)	90.13	79.46
		Varhadi (D4)	83.59	74.81
	Expanded	Southern Konkani (D1)	24.06	44.36
		Northern Konkani (D2)	65.41	79.76
		Varhadi (D4)	67.97	76.90
Multi	Original	Southern Konkani (D1)	26.08	40.94
		Northern Konkani (D2)	47.92	77.50
		Varhadi (D4)	73.04	73.47
		Overall Combined	49.01	63.97
	Expanded	Southern Konkani (D1)	52.15	43.88
		Northern Konkani (D2)	54.35	78.16
		Varhadi (D4)	69.96	74.74
		Overall Combined	58.82	65.59




Figure 6: Held-out test set performance across training setups.

Key observations from Table 7 and Figure 6 include:

- mT5-Small Performance:** mT5-small fine-tuned on the synthetically expanded multi-dialect dataset achieves **65.59 BLEU** overall across all three target dialects on the held-out test set, outperforming IndicBART (58.82 BLEU) by +6.77 BLEU, with notable normalization on Northern Konkani (78.16 BLEU) and Varhadi (74.74 BLEU).
- IndicBART Scaling:** Fine-tuning IndicBART on the synthetically expanded multi-dialect corpus achieves **58.82 BLEU** (up from 49.01 BLEU on original data), showing an overall gain of +9.81 BLEU and a substantial improvement of **+26.07 BLEU** on Southern Konkani (52.15 BLEU).

4.3 Impact of Synthetic Data Augmentation on Test Set

Table 8 and Figure 7 show the overall impact on metrics by synthetic data augmentation directly on the held-out test set under multi-dialect training schemes. IndicBART showed improvement for Southern Konkani (+26.07 BLEU) and Northern Konkani (+6.43 BLEU), but its Varhadi score dropped by 3.08 BLEU. mT5-small showed consistent improvements: +2.94 BLEU for Southern Konkani, +0.66 BLEU for Northern Konkani, and +1.27 BLEU for Varhadi.

Table 8: Effect of data augmentation on held-out test BLEU

Evaluated Subset	IndicBART BLEU			mT5-Small BLEU		
	Original	Expanded	Δ	Original	Expanded	Δ
Southern Konkani (D1)	26.08	52.15	+26.07	40.94	43.88	+2.94
Northern Konkani (D2)	47.92	54.35	+6.43	77.50	78.16	+0.66
Varhadi (D4)	73.04	69.96	-3.08	73.47	74.74	+1.27




Figure 7: BLEU score changes on the held-out test set after data augmentation.

4.4 Verification Engine Impact: Raw Unverified vs. Filtered Clean Data

We also analyzed the performance of our verification and filtering engine on the test set by training on the contaminated versus the clean dataset. The raw unverified dataset contains 19,914 parallel pairs, which included 3,751 rejected pairs that contained flawed or corrupted translations. Table 9 and Figure 8 show the BLEU and chrF++ score differences due to multi-tier verification.

Table 9: Impact of data verification on normalization scores

Model	Pipeline State	Pairs	BLEU	chrF++	Δ BLEU	Δ chrF++
IndicBART (244M)	Raw Unverified Data	19,914	43.09	64.54	-15.73	-10.95
	Filtered Clean Data	32,335	58.82	75.49	+15.73	+10.95
mT5-Small (300M)	Raw Unverified Data	19,914	61.21	80.18	-4.38	-4.40
	Filtered Clean Data	32,335	65.59	84.58	+4.38	+4.40




Figure 8: Effect of verification filtering on BLEU (A) and chrF++ (B).

Training on unverified data led to a notable drop in performance (−15.73 BLEU for IndicBART, −4.38 BLEU for mT5-small). Unfiltered LLM outputs had hallucinated sentences which corrupted the gradients during training. Multi-tier verification removes the flawed pairs successfully, elevating test BLEU scores to 58.82 (IndicBART) and 65.59 (mT5-small).

IndicBART shows high sensitivity to noise and corrupted data, with a 26.7% relative BLEU drop. mT5-small showed a smaller relative drop (6.7%). This could be directly attributed to mT5’s massive pre-training corpus. This highlights the importance of the quality audit of data while using LLMs for generating or correcting the data.

5 Conclusion

This study builds and presents a reproducible framework for text normalization across three Marathi dialects. By the efficient use of LLMs, it solves the problem of not having enough parallel data for low-resource Indic language varieties instead of relying on manual annotations. Also, the verification engine proves to be effective since its correction ability prevented the models from performance degradation.

Our experiments on the held-out test benchmark demonstrate that mT5-small achieves 65.59 BLEU on the expanded multi-dialect dataset, outperforming IndicBART (58.82 BLEU) by 6.77 BLEU overall, with strong individual dialect performance reaching 78.16 BLEU on Northern Konkan and 74.74 BLEU on Varhadi. IndicBART shows significant benefit from synthetic

expansion on Southern Konkani (+26.07 BLEU gain). Furthermore, our multi-tier verification engine proves essential in filtering out flawed synthetic pairs, preventing error propagation and recovering +15.73 BLEU for IndicBART.

These results offer a decent starting point for normalizing other Indic sub-dialects, but a few limitations remain. Firstly, LLMs add up to the extra cost of infrastructure and compute. Also, synthetic expansion comes at a cost of degraded performance for dialects that already perform well (Varhadi dialect in this case). Future work should test the performance around speech-based systems, not just text-based metrics. Overall, our verified pipeline offers a scalable way to expand digital language support for low-resource regional dialects.

References

- Abdul-Mageed, M., Zhang, C., Elmadany, A., and Nagoudi, E. M. (2023). Nadi 2023: The fourth nuanced arabic dialect identification and translation shared task. In *Proceedings of the Arabic NLP Conference (ArabicNLP @ EMNLP 2023)*, pages 512–528.
- Abe, K., Inui, K., and Matsumoto, Y. (2018). Multi-dialect neural machine translation for japanese dialects. In *Proceedings of the 27th International Conference on Computational Linguistics (COLING 2018)*, pages 2105–2116.
- Ansary, N., Adib, Q. A. R., Reasat, T., Sushmit, A. S., Humayun, A. I., Mehnaz, S., Fatema, K., Or Rashid, M. M., and Sadeque, F. (2024). Unicode normalization and grapheme parsing of indic languages. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 17019–17030. ELRA Language Resource Association.
- Blaschke, V., Purschke, C., Schütze, H., and Plank, B. (2024). What do dialect speakers want? a survey of attitudes towards language technology for german dialects. *arXiv preprint arXiv:2402.11968*.
- Bouamor, H., Habash, N., Salameh, M., et al. (2018). The madar parallel corpus of arabic dialects. In *Proceedings of the 11th Language Resources and Evaluation Conference (LREC 2018)*, pages 3375–3381.
- Dabre, R., Shrotriya, H., Kunchukuttan, A., Puduppully, R., Khapra, M. M., and Kumar, P. (2022). Indicbart: A pre-trained model for indic natural language generation. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 1849–1863.
- de Gibert, O. et al. (2024). Scaling low-resource mt via synthetic data generation with llms. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*.
- Dhasmana, A., Srivastava, A., and Chiang, D. (2026). Dialect matters: Cross-lingual asr transfer for low-resource indic language varieties. In *Proceedings of the Thirteenth Workshop on NLP for Similar Languages, Varieties and Dialects (VarDial @ ACL 2026)*, pages 145–156. Association for Computational Linguistics.
- Dimakis, A., Pavlopoulos, J., and Anastasopoulos, A. (2025). Dialect normalization using large language models and morphological rules. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 23696–23714. Association for Computational Linguistics.
- Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Yang, A., Fan, A., and et al. (2024). The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.
- Eryani, F., Habash, N., Al-Shargi, F., et al. (2020). A normalization corpus for five arabic dialects. In *Proceedings of the 12th Language Resources and Evaluation Conference (LREC 2020)*, pages 5750–5757.
- Gaikwad, G., Panchal, D., Deshmukh, O., and Patil, S. (2025). Marathi dialect to pure marathi using rag & ai. *International Research Journal of Engineering and Technology (IRJET)*, 12(11):787–791.
- Google DeepMind Gemma Team (2026). Gemma 4: Open language models. *arXiv preprint arXiv:2603.00118*.
- Habash, N., Diab, M., and Rambow, O. (2012). Conventional orthography for dialectal arabic. In *Proceedings of the Eighth International Conference on Language Resources and Evaluation (LREC '12)*, pages 711–718.
- Hämäläinen, M., Partanen, N., and Alnajjar, K. (2020). Normalizing swedish dialects spoken in finland. In *Proceedings of the 7th Workshop on NLP for Similar Languages, Varieties and Dialects (VarDial @ COLING 2020)*, pages 55–62.
- Hämäläinen, M., Partanen, N., and Rueter, J. (2022). Generating dialectal finnish for dialect normalization of estonian dialects. In *Proceedings of the 9th Workshop on NLP for Similar Languages, Varieties and Dialects (VarDial @ COLING 2022)*, pages 71–79.
- Hsieh, C.-Y. et al. (2023). Distilling step-by-step! outperforming larger language models with less training data and smaller model sizes. *arXiv preprint arXiv:2305.02301*.

- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y. J., Madotto, A., and Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12):1–38.
- Joshi, A., Dabre, R., Kanojia, D., Li, Z., Zhan, H., Haffari, G., and Dippold, D. (2025). Natural language processing for dialects of a language: A survey. *ACM Computing Surveys*, 57(6):1–37.
- Joshi, P., Santy, S., Budhiraja, A., Bali, K., and Choudhury, M. (2020). The state and fate of linguistic diversity and inclusion in the NLP world. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 6282–6293.
- Joshi, R., Goel, P., and Joshi, R. (2022). L3cube-mahanlp: Marathi natural language processing datasets, models, and library. In *Proceedings of the WILDRE-6 Workshop, LREC 2022*, pages 10–17.
- Kakwani, D., Kunchukuttan, A., Golla, S., Gokul, N. C., Bhattacharyya, P., Khapra, M. M., and Kumar, P. (2020). IndicNlpSuite: Monolingual corpora, evaluation benchmarks and pre-trained multilingual language models for indic languages. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 4948–4961. Association for Computational Linguistics.
- Koenecke, A., Nam, A., Lake, E., Nudell, J., Quartey, M., Mengesha, Z., Touts, C., Rickford, J. R., Jurafsky, D., and Goel, S. (2020). Racial disparities in automated speech recognition. *Proceedings of the National Academy of Sciences*, 117(14):7684–7689.
- Kresic, M. and Abbas, N. (2024). Normalizing swiss german dialects with the power of large language models. *Procedia Computer Science*, 244:287–295.
- Kulkarni-Joshi, S. (2023). Variation and change in dialects of marathi: A social-dialectological approach. In Chandra, P., editor, *Variation in South Asian Languages: From Macro to Micro-Differences*, pages 207–236. Springer.
- Kumar, A., Shrivastava, M., and Bhattacharyya, P. (2024). Indic-dialect: Benchmark dataset for indic dialect classification and translation. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 15201–15212. ELRA Language Resource Association.
- Kumar, S., Singh, A., G, D., Amartyaveer, Bandekar, J., Murthy, S., Sharma, S., Badiger, S., Udupa, S., Nagireddi, A., Raghavan, S. K. M., Saxena, R., Nanavati, J., Nanavati, R., Sridharan, J., Mehta, A., Khuraishi, A. K. S., Mora, S. P. R., Venkataramakrishnan, P., Date, G., P. K., and Ghosh, P. K. (2025). RESPIN-S1.0: A read speech corpus of 10000+ hours in dialects of nine indian languages. In *Advances in Neural Information Processing Systems (NeurIPS 2025) Track on Datasets and Benchmarks*, pages 1–13.
- Kuparinen, O., Milić, A., and Scherrer, Y. (2023). Dialect-to-standard normalization: A large-scale multilingual evaluation. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 13814–13828. Association for Computational Linguistics.
- Lusetti, M., Ruzsics, T., Samardžić, T., and Scherrer, Y. (2018). Impact of target language models on swiss german dialect normalization. In *Proceedings of the 5th Workshop on NLP for Similar Languages, Varieties and Dialects (VarDial @ COLING 2018)*, pages 133–141.
- Mukherjee, S., Mitra, A., Jawahar, G., Agarwal, S., Palangi, H., and Awadallah, A. (2023). Orca: Progressive learning from complex explanation traces of gpt-4.
- Mulla, R. and Kumar, B. S. (2022). Text-independent automatic dialect recognition of marathi language using spectro-temporal characteristics of voice. *International Journal on Recent and Innovation Trends in Computing and Communication (IJRITCC)*, 10(2s):313–321.
- Partanen, N., Härmäläinen, M., and Alnajjar, K. (2019). Dialect text normalization to normative standard finnish. In *Proceedings of the 5th Workshop on Noisy User-generated Text (W-NUT @ EMNLP 2019)*, pages 141–146. Association for Computational Linguistics.
- Pawar, P. S., Saini, J. R., and Vaidya, S. (2026). Automatic marathi dialect classification in noisy and non-noisy environments. *Procedia Computer Science*, 282:288–303.
- Pecher, B. et al. (2024). Better as generators than classifiers: Leveraging llms and synthetic data for low-resource multilingual classification. In *Findings of the Association for Computational Linguistics: EACL 2024*.
- Post, M. (2018). A call for clarity in reporting bleu scores. In *Proceedings of WMT 2018*, pages 186–191.
- Pratap, V., Saraf, A., and Kumar, P. (2024). Wfst-guided neural text normalization for indic languages. In *Proceedings of Interspeech 2024*, pages 3120–3124.
- Scherrer, Y. and Ljubešić, N. (2016). Normalising swiss german dialect text using character-level machine translation. In *Proceedings of the 3rd Workshop on NLP for Similar Languages, Varieties and Dialects (VarDial @ COLING 2016)*, pages 117–126.
- Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., Liang, P., and Hashimoto, T. B. (2023). Alpaca: A strong, replicable instruction-following model. *Stanford Center for Research on Foundation Models*.

- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. (2017). Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS 2017)*, pages 5998–6008.
- Vyavhare, V. A., Divekar, S. N., Pawar, H., Kale, K., and Jagtap, S. (2026). Marathi dialect translator. *International Journal of Scientific Research in Science, Engineering and Technology (IJSRSET)*, 13(3):254–261.
- Wang, Y., Kordi, Y., Mishra, S., Liu, A., Smith, N. A., Khashabi, D., and Hajishirzi, H. (2022). Self-instruct: Aligning language models with self-generated instructions.
- Xue, L., Constant, N., Roberts, A., Kale, M., Al-Rfou, R., Siddhant, A., Barua, A., and Raffel, C. (2021). mt5: A massively multilingual pre-trained text-to-text transformer. In *Proceedings of NAACL-HLT 2021*, pages 483–498.
- Zampieri, M., Nakov, P., and Scherrer, Y. (2020). Natural language processing for similar languages, varieties, and dialects: A survey. *Natural Language Engineering*, 26(6):605–637.
- Zhao, Y. and Chodroff, E. (2022). Acoustic-phonetic variation in mandarin dialect corpora. *Journal of Phonetics*, 94:101172.