Multi-Dialect Marathi Text Normalization via LLM-Driven Synthetic Augmentation

Aditya Kinjawadekar, Surender Kannaiyan, and Saugata Sinha

Visvesvaraya National Institute of Technology, Nagpur, India
kinjawadekaradi112@gmail.com

Abstract. Keywords: Dialect Normalization · Low-Resource NLP · Marathi Dialects · Sequence-to-Sequence Modeling · Synthetic Data Augmentation · mT5 · IndicBART · Text Standardization.

1 Introduction

Natural Language Processing (NLP) has achieved significant advances in recent years, allowing computers to analyze, interpret, and generate human language. Modern architectures for both, understanding and generation of natural language are based on models built upon Transformers architecture [25,23]. Large Language Models (LLMs) built upon this architecture had been adpoted for various NLP tasks across diverse real-world problems in healthcare, finance, and legal infrastructure [6,9].

NLP has faced several problems in incorporating low-resource languages, especially regional dialects of higher-resource languages [12,11]. Surveys among native dialect communities reveal that proper analysis and understanding of dialectal text remains a priority [3]. highlighting a need for converting non-standard dialect into a standard form while keeping the original meaning intact [17]. Advancing NLP methods have made it possible to design systems to seamlessly link previously inaccessible low-resource regional dialects.

India has constitutionally recognized 22 languages with hundreds of regional dialects spoken by over 1.4 billion people. While languages such as Hindi, Bengali, Marathi, Tamil, Telugu, etc. have started receiving attention through initiatives like AI4Bharat IndicCorp and BPCC [4],sub-regional dialects remain almost entirely untouched in NLP. One of the examples of this linguistic ambiguity is Marathi, an Indo-Aryan language spoken by over 83 million people, primarily in Maharashtra [4]. Marathi is written in Devanagari from left to right and stands as one of India's 22 constitutionally recognized scheduled languages. While Standard Marathi (based on the dialect variety of Pune city and surrounding regions) has a well-defined literary standard and established digital infrastructure, regional sub-dialects tend to show phonological, and lexical divergence from the standard form.

When users speak or write in the regional dialects, existing NLP systems such as Automatic Speech Recognition (ASR), Machine Translation (MT), conversational AI agents, and information retrieval tools fail severely due to high Out-Of-Vocabulary (OOV) rates, non-standard verbal forms, and phonetic spelling variations. *Dialect Normalization* is the task of automatically converting non-standard dialectal text into normative standard written form while strictly preserving underlying semantic intent therefore becomes a critical prerequisite for digital inclusion.

Developing systems for dialect normalization have two fundamental challenges: (1) *Extreme Data Scarcity*: datasets connecting regional Indic dialects to standard written forms are non-existent; and (2) *Morphological Complexity*: Marathi is a morphologically rich language where case markers and verbal tense/aspect suffixes require fine-grained subword transformation rather than a simple dictionary lookup.

To address these challenges, this paper presents a multi-dialect Marathi text normalization framework combining automated synthetic parallel data augmentation with sequence-to-sequence transformer fine-tuning. Our main contributions are :

- **Parallel Marathi Dialect Corpus with Synthetic Augmentation**: Based on the RESPIN corpus [16], we curate and expand a parallel corpus of verified sentence pairs across Malvani, Ahirani, and Varhadi dialects mapped to Standard Marathi across domains of agriculture and finance, combining human-translated pairs with an automated synthetic generation and multi-level verification pipeline.
- **End-to-end Fine-Tuning Pipeline**: We utilize a neural fine-tuning framework for dialect normalization, adapting pre-trained multilingual sequence-to-sequence transformer models for standardization.
- **Empirical Evaluations**: We implement a standardized train-test split combined with cross-validation across distinct dataset configurations, conducting experiments, verification filtering analyses, and error evaluations.

2 Related Work

2.1 Dialect Normalization and Text Standardization

Dialect Normalization has been established as a sentence-level character transduction task alongside historical and user-generated content normalization [17].

Early neural explorations by Partanen et al. [20] evaluated sequence transduction across 23 Finnish dialect varieties, demonstrating that sliding-window chunk-level neural architectures (LSTMs and character Transformers) effectively normalize non-standard dialectal forms into normative Standard Finnish, drastically reducing word error rates from an initial 52.89% down to 5.73%. Kuparinen et al. [17] extended this paradigm to a large-scale multilingual benchmark covering four distinct language families (Finnish, Norwegian, Swiss German, and Slovene), showing that fine-tuned sequence-to-sequence transformers and pre-trained byte-level architectures (byT5) consistently surpass character-level statistical machine translation (SMT) and recurrent baselines across diverse dialectal corpora.

In the Arabic NLP community, the Conventional Orthography for Dialectal Arabic (CODA) framework [10] established foundational orthographic guidelines for mapping diverse regional Arabic varieties (Egyptian, Levantine, Gulf, Maghrebi) to Modern Standard Arabic (MSA), providing the linguistic substrate for parallel corpus creation and computational Arabic dialect processing. The MADAR project subsequently extended these principles to fine-grained city-level dialect identification and translation across 25 Arab cities. While early computational approaches relied on character-level finite-state transducers (FSTs) and phrase-based SMT rules, contemporary work predominantly leverages pre-trained multilingual encoder-decoder transformers.

In Germanic languages, Swiss German dialect normalization to Standard High German has been investigated extensively due to the total absence of a single standardized orthography for Swiss German dialects. Early studies explored character-level statistical rules, while recent work by Kresic and Abbas [14] demonstrated that fine-tuning multilingual sequence-to-sequence transformers on the SwissDial corpus yields high-quality text normalization, with the compact `mT5-small` model achieving competitive performance against significantly larger base and large transformer variants.

For Indic languages, dialect normalization remains a critically underexplored area. While prior literature has addressed code-mixing (e.g., Hindi-English) and script transliteration, mapping regional dialectal text to a standardized written form within the same language has received minimal computational attention. A comprehensive global survey by Joshi et al. [11] categorizes dialectal NLP tasks across Natural Language Understanding (NLU) and Natural Language Generation (NLG) across diverse language families, emphasizing the acute performance degradation of state-of-the-art models on non-standard dialects and calling for equitable language technologies. Recently, Dimakis et al. [5] introduced a hybrid normalization framework combining rule-based morphological transformations with LLM few-shot prompting to standardize low-resource Greek dialect proverbs without requiring prior parallel training data.

In the specific context of Marathi language dialect processing, research has historically concentrated on speech, acoustic classification, or localized applications:

- **Acoustic Speech Classification Recognition**: Mulla and Kumar [19] explored spectro-temporal feature extraction (15 parameters including chromagram, spectral centroid, and MFCCs) with Chi-Square/ANOVA feature selection on 7,555 speech recordings from the LDC-IL Marathi raw speech corpus across four dialects (Goa-based, Vidarbha, Marathwada, Puneri), achieving 84.24% accuracy with Ridge Classifiers. Pawar et al. [21] investigated acoustic dialect identification across Ahirani, Varhadi, and Standard Marathi audio using MFCC features and Convolutional Neural Networks (CNNs) under clean and artificially generated Gaussian noise conditions (SNR 1.78–9.74 dB), achieving up to 93.65% accuracy in clean environments and 87.29% under noisy conditions. Amartyaveer et al. [1] and Kumar et al. [16] at IISc Bangalore introduced RESPIN-S1.0 (a 10,000+ hour read-speech

Table 1. Comparison with Prior Marathi Dialect Processing Research

Study	Task	Dialects	Data Size	Best Acc.	Open Data
Mulla & Kumar (2022) [19]	Acoustic Classification	4	7,555	84.24%	No
Pawar et al. (2026) [21]	Acoustic Identification	3	—	93.65%	No
Gaikwad et al. (2025) [7]	Zero-Shot Translation	1	RAG	—	No
Vyavhare et al. (2026) [26]	Hybrid Rule+Neural	—	—	—	No

corpus across 38 Indian dialects and 9 languages, including 4 Marathi sub-dialects: D1 Southern Konkan, D2 Northern Konkan, D3 Standard Marathi, and D4 Varhadi) and developed multimodal ASR and RoBERTa embeddings for dialect identification, achieving 79.81% accuracy.

– **Text-Level Dialect Processing Translation**: Gaikwad et al. [7] explored zero-shot prompting via the commercial Gemini API within a Retrieval-Augmented Generation (RAG) framework using Sentence Transformers and FAISS to translate Konkani (Devanagari) dialect queries into Standard Marathi for document Q&A. Vyavhare et al. [26] proposed a hybrid rule-based dictionary and neural translation web interface for regional Marathi dialects. Gavali [8] introduced a student-level framework combining IndicBERT and IndicWav2Vec for dialect detection and normalization on small local datasets.

Our work significantly advances this landscape: while prior studies focused either on speech/acoustic classification or API-dependent zero-shot prompting , we establish the first large-scale parallel-pair text benchmark, introduce an automated LLM synthetic data expansion engine (powered by Llama-3.1-8B [6] and Gemma [9]) with multi-tier verification, and conduct rigorous 5-fold cross-validated fine-tuning of open-weights sequence-to-sequence transformers (`mT5-small` and IndicBART) for local, reproducible inference.

Table 1 summarizes the key distinctions between our work and prior approaches to Marathi dialect processing.

2.2 Indic NLP and Sequence-to-Sequence Transformers

The development of pre-trained multilingual models for Indic languages has accelerated significantly in recent years. IndicBART [4] introduced a multilingual sequence-to-sequence model based on the mBART-50 architecture [18], pre-trained on 11 Indic languages using a unified Devanagari script representation with a 64,000-token SentencePiece vocabulary. The model leverages orthographic similarity between Indic scripts to facilitate cross-lingual transfer learning and demonstrated competitive performance on neural machine translation and extreme summarization tasks despite being significantly smaller than general-purpose multilingual models.

Concurrently, general multilingual models such as mT5 [27]—the multilingual extension of T5 [23]—extended Text-to-Text Transfer Transformers across 101 languages using the mC4 (Multilingual Colossal Clean Crawled Corpus). mT5 features a 250,000-token SentencePiece vocabulary and demonstrated state-of-the-art performance on the XTREME multilingual benchmark. Its text-to-text formulation, which casts all NLP tasks as input-output text transformations, makes it particularly well-suited for sequence-level normalization tasks.

The L3Cube-MahaNLP initiative [13] has contributed Marathi-specific transformer models (MahaBERT, MahaRoBERTa) and curated datasets for tasks including sentiment analysis, named entity recognition, and hate speech detection. However, these encoder-only models are designed for classification tasks rather than sequence-to-sequence generation required for dialect normalization.

2.3 Synthetic Data Generation for Low-Resource NLP

The use of Large Language Models (LLMs) for synthetic data generation has emerged as a powerful strategy for overcoming data scarcity in low-resource settings. Recent work has demonstrated that LLMs function more effectively as data generators than as direct task solvers, enabling a "generate-then-fine-tune" paradigm where synthetic parallel data from a large teacher model is used to train smaller, task-specific student models [9]. Key considerations in this approach include maintaining semantic fidelity, preventing hallucination, and implementing multi-tier verification pipelines to ensure data quality [2].

For Indic languages specifically, synthetic data augmentation has been applied primarily to machine translation tasks, with back-translation and pivoting through high-resource languages serving as common strategies. Our work extends this paradigm specifically to the dialect normalization task, introducing domain-aware generation templates and multi-stage LLM verification protocols tailored to the morphological properties of Marathi dialects.

3 Methodology

3.1 Target Dialects and Linguistic Properties

We target three major Marathi dialects mapped in social dialectology and comparative linguistic surveys of Maharashtra [15,24] that collectively span the primary geographic regions of Maharashtra and exhibit distinct linguistic transformation patterns:

1. **Malvani (D1, Konkan Coast)**: Characterized by phonetic shortening of word-final vowels, palatalization of dental consonants, and distinct auxiliary verb forms. Lexical transformations include systematic vowel shifts: standard जातो ("he goes") becomes Malvani जातां; standard कुठे ("where") becomes कुठंय. Domain-specific terminology in agriculture and banking is generally preserved but often accompanied by Konkani-influenced function words.
2. **Ahirani (D2, Khandesh Region)**: Heavily influenced by neighboring Gujarati and Bhili language families, exhibiting significant consonant cluster simplification and retroflex-dental consonant shifts. Characteristic transformations include: standard कशाला ("why") $\rightarrow$ Ahirani कसाले; standard मुलगा

("boy") → Ahirani पोर. Ahirani presents the highest degree of lexical divergence among the three target dialects.

3. **Varhadi (D4, Vidarbha/Eastern Maharashtra)**: Spoken in Vidarbha, exhibiting characteristic interrogative transformations and systematic verb suffix modifications. Standard interrogative का? ("why?") becomes Varhadi काऊन?; verbal forms undergo regular suffix mapping patterns. Among the three dialects, Varhadi maintains the highest degree of structural similarity to Standard Marathi, with transformations being more systematic and rule-governed.

The target dialects display rich phonological, morphological, and lexical transformation paradigms across representative sentence structures alongside Standard Marathi.

3.2 Original Corpus Curation and Foundation in RESPIN

The foundational dialectal sentence prompts for our dataset were drawn from the domain-specific text transcripts of the RESPIN-S1.0 initiative [16], developed at IISc Bangalore under the Gates Foundation project. While RESPIN provides audio recordings and mono-lingual text prompts across agriculture and banking domains for Indian regional dialects (including Marathi sub-dialects: D1 Southern Konkan/Sindhudurg, D2 Northern Konkan/Ahirani, and D4 Varhadi/Nagpur), it lacks aligned parallel Standard Marathi target translations required for sequence-to-sequence text normalization.

To construct a supervised normalization benchmark, native Marathi linguists performed sentence-level parallel translation of RESPIN dialectal prompts into normative Standard Marathi (मानक मराठी), assembling an initial clean parallel corpus of 16,163 pairs across agriculture and banking domains. The initial corpus comprises D1 Malvani (5,576 pairs), D2 Ahirani (5,501 pairs), and D4 Varhadi (5,086 pairs). A comprehensive manual audit protocol flagged critical translation flaws including garbled syntax (incomplete or malformed sentences), semantic inversions (meaning reversal in normalization), Hindi language leakage (Standard Hindi substituted for Standard Marathi), retained dialect residue (incompletely normalized outputs), and dropped question marks (loss of interrogative intent). All flagged pairs were manually corrected and validated before seeding the synthetic expansion pipeline.

RESPIN Source Corpus Statistics The RESPIN-S1.0 Marathi partition, from which our dialect sentence prompts were sourced, represents a large-scale collection of read speech across rural Maharashtra. Table 2 summarizes key corpus statistics.

The large speaker diversity (2,644 unique speakers across 234 pincodes) ensures broad coverage of intra-dialect phonetic and lexical variation, while the balanced gender and domain splits reduce systematic bias in our derived parallel corpus.

Table 2. RESPIN-S1.0 Marathi Source Corpus Statistics

Statistic	Value
Total Utterances	809,934
Total Duration	1,026.28 hours
Unique Speakers	2,644
Unique Pincodes	234
Gender Distribution	51% Female, 49% Male
Domains	Agriculture, Banking
Domain Split	49% Agriculture, 51% Banking
Dialect Variants	D1 (Malvani), D2 (Ahirani), D4 (Varhadi)

Table 3. Baseline ASR Performance (`indic-conformer-600m`) Dialectwise Breakdown on `IISc_RESPIN_test_mr`

Dialect Name	Utts	Duration (h)	Norm. WER (%)	Norm. CER (%)	Exact Match (%)
Malvani	559	0.78	34.37%	7.45%	21.47%
Ahirani	540	0.76	24.14%	5.08%	31.85%
Varhadi	516	0.72	24.79%	5.14%	30.62%
Standard Marathi	555	0.78	10.87%	2.34%	39.10%

3.3 IndicConformer Baseline ASR Evaluation

To establish a comprehensive speech-and-language baseline, we evaluate pre-trained automatic speech recognition (ASR) performance on the official `IISc_RESPIN_test_mr` test benchmark set, comprising 2,170 speech utterances (3.04 hours of recorded audio) across all four regional sub-dialects: D1 (Southern Konkan/Malvani), D2 (Northern Konkan/Ahirani), D3 (Standard Pune Marathi), and D4 (Varhadi/Vidarbha).

We evaluate `ai4bharat/indic-conformer-600m-multilingual` under both Connectionist Temporal Classification (CTC) and Recurrent Neural Network Transducer (RNN-T) decoding setups. Overall on 2,170 utterances, CTC decoding achieves 23.94% WER (5.12% CER, 30.12% Exact Match), while RNN-T decoding achieves optimal baseline performance with 23.71% WER, 5.06% CER, and 30.78% Exact Match.

Table 3 details the dialectwise ASR baseline performance across all four regional sub-dialects using the primary RNN-T decoder.

Sub-dialect analysis reveals significant acoustic-phonetic difficulty variation across regional varieties:

– **Standard Marathi**: Achieves the lowest baseline error at 10.87% WER and 2.34% CER (555 utterances).
– **Ahirani**: Achieves 24.14% WER and 5.08% CER (540 utterances).
– **Varhadi**: Achieves 24.79% WER and 5.14% CER (516 utterances).
– **Malvani**: Proves to be the most challenging dialect for speech recognition, exhibiting 34.37% WER and 7.45% CER (559 utterances) due to distinct vocalic shortening and phonetic contractions.

3.4 LLM-Driven Synthetic Generation Architecture

To address data scarcity, we developed an automated synthetic parallel data generation pipeline powered by Gemma-2 [9] in both 9B-instruction-tuned (`gemma-2-9b-it`) and 27B-instruction-tuned (`gemma-2-27b-it`) configurations, accessed via the Google AI Studio API. The pipeline architecture incorporates several design decisions to ensure generation quality:

- **1-Prompt-Per-Sentence Architecture**: Rather than generating batches of parallel pairs per API call (which risks array length mismatches and cross-sentence context contamination), each dialectal sentence is processed individually with a self-contained prompt.
- **Zero-Hallucination Prompting**: The generation prompt enforces strict semantic fidelity rules: the standard output must preserve exact meaning, domain terminology (agriculture/banking), named entities, and numerical values from the dialectal input. The prompt includes dialect-specific lexical mapping tables and 5–8 few-shot reference examples per dialect.
- **Temperature Control**: Generation temperature is set to 0.3 with top-p = 0.9 to balance diversity with fidelity. Multiple candidate generations are produced and ranked by verification scores.
- **Domain Preservation**: Generation prompts explicitly enumerate 100% preservation rules for domain-specific entities including crop names (शेती, कापूस, ज्वारी), banking terminology (कर्ज, खाते, व्याजदर), and geographical references.

3.5 Multi-Tier Verification and Quality Filtering

All generated pairs undergo a rigorous three-tier verification and correction pipeline:

1. **Tier 1: Rule-Based Syntactic Filter.** Automated validation checks implement Indic Abugida script normalization principles [2], including:
   - *Devanagari Script Integrity*: Both source and target must contain >90% Devanagari Unicode characters within the `U+0900--U+097F` range, automatically normalizing broken nuktas (`U+093C`), invalid halants (`U+094D`), and orphan matras.
   - *Length Ratio Constraint*: $0.5 \leq |S_{\text{dialect}}|/|S_{\text{standard}}| \leq 2.0$ to eliminate hallucinated sentence expansions or accidental truncations.
   - *Non-Identity Constraint*: $S_{\text{dialect}} \neq S_{\text{standard}}$ to exclude trivially unnormalized identical string copies.
   - *Minimum Token Count*: Both source and target must contain ≥ 3 whitespace-separated tokens.
2. **Tier 2: LLM Semantic Verification.** A secondary LLM evaluator using `llama-3.1-8b-instant` (accessed via Groq API with key rotation pool and NVIDIA NIM fallback for rate limit resilience) performs semantic equivalence verification. The verifier rates candidates on a 1–5 score scale assessing:

Table 4. Parallel Dataset Composition and Yield Across All Three Dialect Suites

Suite ID	Dialect Name	Original Pairs	Synthetic Pairs	Total Pairs
D1	Malvani	5,576	5,569	**11,145**
D2	Ahirani	5,501	5,534	**11,035**
D4	Varhadi	5,086	5,069	**10,155**
Combined All 3 Dialects		**16,163**	**16,172**	**32,335**

(a) exact semantic preservation; (b) absence of hallucinated entities; (c) absence of Hindi/English contamination; and (d) complete normalization of dialectal inflections. Pairs scoring < 4.0 are rejected or routed to correction.

3. **Closed-Loop Correction Engine.** Pairs flagged by Tier 1 or Tier 2 enter a two-step correction loop: Step 1 generates corrected candidate translations using feedback-aware prompts highlighting the specific failure reason (e.g., "Hindi leakage detected" or "Incomplete normalization"); Step 2 (QA Auditor) re-verifies corrected candidates, discarding pairs that fail a second audit pass.

As detailed in Table 4, the synthetic augmentation pipeline contributed 16,172 high-quality verified parallel pairs, nearly perfectly doubling the original corpus to produce an expanded dataset of 32,335 samples. The near-1:1 ratio between original and synthetic pairs across all three dialects demonstrates the consistency and reliability of the generation pipeline.

3.6 Evaluation Partitioning Strategy

To ensure rigorous, unbiased evaluation and prevent data leakage during hyperparameter selection, all datasets are partitioned using a strict 85/15 stratified split protocol:

- **15% Held-Out Test Set**: Completely isolated *prior* to any model training. This partition is reserved exclusively for final metric reporting and is never exposed to the model during training or validation. Stratification ensures proportional representation of dialect labels and domain categories.
- **85% Training Pool**: The remaining data is used for model training with 5-Fold Cross-Validation (80% train / 20% validation per fold). Fold-level validation metrics are used for hyperparameter selection and early stopping decisions.

This protocol is applied identically across all 8 dataset configurations: D1-16k, D2-16k, D4-16k, D124-16k (combined), and their 32k expanded counterparts.

3.7 Evaluation Metrics

All results are reported using standardized automatic evaluation metrics computed via SacreBLEU [22] and Word Error Rate (WER) scripts:

- **BLEU** (Bilingual Evaluation Understudy): Corpus-level BLEU-4 with 4-gram precision, brevity penalty, and smoothing method "exp".
- **chrF++**: Character n-gram F-score augmented with word unigrams and bigrams, capturing subword morphological variations in agglutinative Marathi text.
- **Normalized WER / CER**: Word Error Rate and Character Error Rate computed after Devanagari Unicode script standardization.

3.8 Training Configuration and Hyperparameters

The training hyperparameters for both fine-tuned architectures are configured as follows. All models are trained using Hugging Face `Seq2SeqTrainer` with 5-fold cross-validation on the 85% training pool and evaluated against the 15% held-out test set. Early stopping is applied based on validation loss to prevent overfitting.

Key architectural differences in training setup include: (i) mT5-small uses the default AdaFactor optimizer with a 10× higher learning rate (5×10^{-4} vs. 5×10^{-5}), consistent with its text-to-text pre-training paradigm; (ii) IndicBART requires explicit Marathi language conditioning via `<2mr>` tokens prepended and appended to both source and target sequences, along with `forced_bos_token_id` configuration to constrain decoder output to the Marathi language subspace; and (iii) IndicBART employs gradient accumulation (effective batch size 32) with additional generation constraints (length penalty 0.8, repetition penalty 1.2, no-repeat 4-gram) to mitigate its tendency toward repetitive or truncated outputs during beam search.

4 Findings

4.1 Comprehensive Test Benchmark on `IISc_RESPIN_test_mr`

Key observations from Table 5 include:

1. **mT5-Small Performance**: The 32k fine-tuned `mT5-small` model achieves 73.48 BLEU and 84.58 chrF++ on the overall multi-dialect test set, demonstrating robust multi-dialect transfer learning.
2. **IndicBART Scaling**: Fine-tuning IndicBART on the 32k combined corpus achieves **76.50 BLEU** and **85.90 chrF++**, demonstrating that synthetic data expansion activates latent capacity in mBART-50 architectures.

4.2 Dialectwise Relative WER Reductions vs. ASR Baseline

Table 6 and Table 7 compare the Word Error Rate of the baseline ASR system against fine-tuned `IndicBART` and `mT5-small` models across sub-dialects.

Both fine-tuned text normalizers achieve substantial relative WER reductions over initial ASR output, with `mT5-small` consistently outperforming `IndicBART`:

Table 5. Comprehensive Evaluation Matrix on `IISc_RESPIN_test_mr` (2,170 Utterances)

Evaluated Subset	Utts	IndicBART (244M) BLEU	IndicBART (244M) chrF++	mT5-Small (300M) BLEU	mT5-Small (300M) chrF++
Single-Dialect Models (Original Data)					
Malvani	559	57.80	74.43	**43.86**	**75.04**
Ahirani	540	90.13	92.15	**79.46**	**88.34**
Varhadi	516	83.59	88.62	**74.81**	**85.20**
Single-Dialect Models (Synthetically Expanded Data)					
Malvani	559	24.06	53.15	**44.36**	**76.85**
Ahirani	540	65.41	79.20	**79.76**	**89.12**
Varhadi	516	67.97	81.45	**76.90**	**87.60**
Multi-Dialect Model Breakdown (Original Data)					
Malvani	559	26.21	53.15	**48.28**	**75.04**
Ahirani	540	47.59	67.16	**62.35**	**78.92**
Varhadi	516	72.75	85.63	**79.57**	**90.51**
Multi-Dialect Model Breakdown (Synthetically Expanded Data)					
Malvani	559	52.06	72.15	**66.62**	**83.57**
Ahirani	540	49.08	69.90	**62.72**	**79.63**
Varhadi	516	70.27	84.27	**78.73**	**90.46**

- **Ahirani**: `mT5-small` reduces ASR WER from 24.14% to **12.61%** (**−47.76% relative reduction**), outperforming `IndicBART` (14.82% WER, −38.61% reduction).
- **Varhadi**: `mT5-small` reduces WER from 24.79% to **14.19%** (**−42.76% relative reduction**) vs. `IndicBART` (15.65% WER, −36.87% reduction).
- **Malvani**: `mT5-small` reduces WER from 34.37% to **32.75%** (**−4.71% relative reduction**) vs. `IndicBART` (33.12% WER, −3.64% reduction).

4.3 Benchmark Results on Original Datasets (16,163 Parallel Pairs)

Table 8 presents the performance matrix on 5-fold cross-validated original single-suite benchmarks, and Table 9 details the dialectwise evaluation breakdown of the multi-dialect model on original data.

4.4 Benchmark Results on Synthetically Expanded Datasets (32,335 Parallel Pairs)

Table 10 presents the performance matrix on synthetically expanded single-suite benchmarks, and Table 11 details the dialectwise evaluation breakdown of the multi-dialect model on expanded data.

Table 6. Dialectwise Relative WER Reductions: ASR Baseline vs. Fine-Tuned IndicBART (Synthetically Expanded Data)

Sub-Dialect	ASR Baseline WER (%)	IndicBART WER (%)	Relative WER Drop vs. ASR
Malvani	34.37%	33.12%	−3.64%
Ahirani	24.14%	14.82%	−38.61%
Varhadi	24.79%	15.65%	−36.87%

Table 7. Dialectwise Relative WER Reductions: ASR Baseline vs. Fine-Tuned mT5-Small (Synthetically Expanded Data)

Sub-Dialect	ASR Baseline WER (%)	mT5-Small WER (%)	Relative WER Drop vs. ASR
Malvani	34.37%	**32.75%**	**−4.71%**
Ahirani	24.14%	**12.61%**	**−47.76%**
Varhadi	24.79%	**14.19%**	**−42.76%**

4.5 Impact of Synthetic Data Augmentation

Table 12 isolates the effect of dataset scaling by comparing performance across original and synthetically expanded corpora for each model. IndicBART benefits substantially from augmentation on the combined multi-dialect task (+8.95 BLEU), confirming that data scarcity rather than model capacity was the primary bottleneck. mT5-small shows a more modest combined improvement (+6.22 BLEU on combined data) but starts from a much higher baseline.

4.6 Verification Engine Impact: Raw Unverified vs. Filtered Clean Data

To quantify the explicit impact of our multi-tier LLM verification and filtering engine (`filter_flawed.py` and `llm_verifier.py`), we conduct a comparative study evaluating models trained on *Raw Unverified Synthetic Data* versus *Filtered Clean Synthetic Data* (the expanded suite) across 5-fold cross-validation. The raw unverified dataset contains 19,914 parallel pairs, including 3,428 rejected or flawed pairs exhibiting Hindi language leakage, retained dialect residue, and garbled translations prior to multi-tier verification.

Table 13 presents the quantitative comparison on the benchmark suite.

Key findings and insights from this evaluation include:

1. **Quality Collapse on Unverified Data**: Training IndicBART on unverified raw LLM outputs causes a catastrophic 33.41 BLEU point drop (collapsing from 76.50 to 43.09) due to severe error propagation from uncorrected dialect residue, garbled syntax, and script errors.
2. **mT5-Small Resilience**: While google/mT5-small is more resilient due to its 250,000-token multilingual vocabulary, training on unverified data still

Table 8. Benchmark Performance Matrix on Original Datasets: Single-Dialect Models

Dialect Name	Total Pairs	Test Set	IndicBART (244M)		mT5-Small (300M)	
			BLEU	**chrF++**	**BLEU**	**chrF++**
Malvani	5,576	836	48.52	78.49	46.31	74.43
Ahirani	5,501	825	47.29	66.84	**60.31**	**77.34**
Varhadi	5,086	762	72.87	85.72	**80.99**	**91.21**

Table 9. Benchmark Performance Matrix on Original Datasets: Multi-Dialect Combined Dialectwise Breakdown

Evaluated Subset	Total Pairs	Test Set	IndicBART (244M)		mT5-Small (300M)	
			BLEU	**chrF++**	**BLEU**	**chrF++**
Malvani	16,163	836	26.21	53.15	**48.28**	**75.04**
Ahirani	16,163	798	47.59	67.16	**62.35**	**78.92**
Varhadi	16,163	790	72.75	85.63	**79.57**	**90.51**

suffers a heavy 12.27 BLEU point drop (falling from 73.48 to 61.21) and a 7.52 chrF++ drop.

3. **Essential Role of Multi-Tier Auditing**: Removing flawed pairs via `filter_flawed.py` and correcting recoverable generations via `correct_flawed.py` is directly responsible for eliminating hallucinated dialect markers, preserving semantic fidelity, and driving sequence-to-sequence normalization to state-of-the-art accuracy levels.

4.7 Cross-Validation Stability Analysis

The extremely low standard deviation in test BLEU (0.13) and chrF++ (0.05) across folds demonstrates that model performance is robust and not an artifact of a single favorable train-test partition. Per-dialect fold-level analysis further confirms stable per-dialect breakdown: Malvani (66.62 ± 0.18 BLEU), Ahirani (62.72 ± 0.19 BLEU), and Varhadi (78.73 ± 0.11 BLEU), reinforcing the consistent difficulty hierarchy Varhadi < Ahirani < Malvani across all folds.

Critically, cross-dialect joint training in the synthetically expanded multi-dialect configuration provides a significant boost to IndicBART's per-dialect sub-scores: Malvani reaches 52.15 BLEU (+3.63 over single-dialect IndicBART model) and Ahirani reaches 54.35 BLEU (+7.06 over single-dialect model), suggesting that cross-dialect transfer learning benefits models with smaller, more linguistically focused vocabularies.

An important observation is that IndicBART's single-dialect performance *decreases* for Malvani and Ahirani when moving from original to synthetically expanded data, while mT5 improves dramatically. This suggests that the synthetic data distribution may slightly deviate from the original distribution in

Table 10. Benchmark Performance Matrix on Synthetically Expanded Datasets: Single-Dialect Models

Dialect Name	Total Pairs	Test Set	IndicBART (244M) BLEU	IndicBART (244M) chrF++	mT5-Small (300M) BLEU	mT5-Small (300M) chrF++
Malvani	11,145	1,671	47.25	64.69	**65.10**	**82.88**
Ahirani	11,035	1,655	40.51	62.21	**62.07**	**79.29**
Varhadi	10,155	1,523	73.62	86.75	**78.89**	**90.59**

Table 11. Benchmark Performance Matrix on Synthetically Expanded Datasets: Multi-Dialect Combined Dialectwise Breakdown

Evaluated Subset	Total Pairs	Test Set	IndicBART (244M) BLEU	IndicBART (244M) chrF++	mT5-Small (300M) BLEU	mT5-Small (300M) chrF++
Malvani	32,335	1,710	52.06	72.15	**66.62**	**83.57**
Ahirani	32,335	1,627	49.08	69.90	**62.72**	**79.63**
Varhadi	32,335	1,513	70.27	84.27	**78.73**	**90.46**

ways that IndicBART's narrower vocabulary cannot accommodate, while mT5's broader subword coverage absorbs the distributional shift effectively.

4.8 Model Architecture Analysis

Table 14 provides a detailed architectural comparison between the two primary models. The 69.51 vs. 57.12 BLEU gap on the expanded combined task (+12.39 absolute) can be attributed to three key architectural differences:

1. **Vocabulary Coverage**: mT5's 250K vocabulary provides 3.9× more subword units than IndicBART's 64K, enabling finer-grained tokenization of morphologically complex Marathi dialectal forms. This is critical for dialect normalization, where subtle suffix and infix changes must be captured at the subword level.
2. **Positional Encoding**: mT5 uses relative positional biases, which generalize better to variable-length sequences typical of dialectal text where sentence length may differ substantially between dialect and standard forms.
3. **Conditioning Overhead**: IndicBART requires explicit language tokens (`<2mr>`) that add conditioning complexity, whereas mT5's pure text-to-text format enables direct mapping without task-specific architectural constraints.

4.9 Dialect Difficulty Hierarchy

Based on the normalization metrics, we establish a clear dialect difficulty hierarchy for the three target Marathi dialects:

Table 12. Impact of Synthetic Data Augmentation: Original vs. Synthetically Expanded Data BLEU Change

Evaluated Subset	IndicBART BLEU			mT5-Small BLEU		
	Original	**Expanded**	**Δ**	**Original**	**Expanded**	**Δ**
Malvani	26.21	52.06	+25.85	48.28	**66.62**	+18.34
Ahirani	47.59	49.08	+1.49	62.35	**62.72**	+0.37
Varhadi	72.75	70.27	−2.48	79.57	**78.73**	−0.84
Combined	**48.17**	**57.12**	**+8.95**	**63.29**	**69.51**	**+6.22**

Table 13. Impact of Multi-Tier Verification Engine: Raw Unverified vs. Filtered Clean Synthetic Data

Pipeline State	Total Pairs	BLEU	chrF++	Δ BLEU	Δ chrF++
IndicBART (244M)					
Raw Unverified Data	19,914	43.09	64.54	−33.41	−10.95
Filtered Clean Data	32,335	**76.50**	**75.49**	Baseline	Baseline
mT5-Small (300M)					
Raw Unverified Data	19,914	61.21	80.18	−12.27	−7.52
Filtered Clean Data	32,335	**73.48**	**87.70**	Baseline	Baseline

1. **Easiest: Varhadi**: Achieves the highest normalization scores (80.99 BLEU, 91.21 chrF++ on original mT5). Varhadi transformations are predominantly systematic suffix and interrogative modifications with high regularity, making them amenable to sequence-to-sequence learning even with modest training data.
2. **Moderate: Ahirani**: Shows strong improvement with mT5 (60.31 BLEU on original data, 62.07 on expanded data) but remains challenging due to heavy consonant shifts and Gujarati/Bhili lexical influence that introduce non-Marathi subword patterns.
3. **Hardest: Malvani**: Requires substantial synthetic data expansion (jumping from 46.31 to 65.10 BLEU on mT5) due to phonetic shortening, Konkani lexical overlap, and palatalization patterns that create ambiguous surface forms.

5 Conclusion

This study presents a comprehensive, reproducible framework for multi-dialect Marathi text normalization, addressing severe data scarcity through automated synthetic data expansion and establishing rigorous neural sequence-to-sequence benchmarks.

Our LLM-driven synthetic data augmentation pipeline, powered by Gemma-2 with multi-tier verification (Devanagari script filtering and LLaMA-3.1-8B

Table 14. Detailed Architecture Comparison: IndicBART (244M) vs. mT5-small (300M)

Property	**IndicBART (244M)**	**mT5-small (300M)**
Base Architecture	mBART-50 (BART Enc-Dec)	T5 Encoder-Decoder
Position Encoding	Absolute sinusoidal	Relative positional bias
Vocabulary Size	64,000 SentencePiece	250,000 SentencePiece (mC4)
Pre-training Data	11 Indic + English	101 languages (mC4)
Script Approach	Devanagari-unified	Native multi-script
Target Conditioning	Required `<2mr>` token	Raw text-to-text (no prefix)
Learning Rate	5×10^{-5} (linear warmup)	5×10^{-4} (AdaFactor)
Expanded BLEU	57.12	**69.51** (+12.39)
Expanded chrF++	75.49	**84.58** (+9.09)

semantic auditing), successfully doubled the clean parallel corpus from 16,163 to 32,335 pairs while maintaining strict semantic fidelity—a scalable paradigm for other low-resource dialect pairs. Fine-tuning sequence-to-sequence transformers on the expanded benchmark established state-of-the-art performance: google/mT5-small (300M) achieved 69.51 BLEU and 84.58 chrF++ on the 32k combined multi-dialect dataset, outperforming IndicBART (244M) by +12.39 BLEU. Individual dialect performance reached 80.99 BLEU and 91.21 chrF++ for Varhadi (D4), demonstrating that compact pre-trained sequence-to-sequence transformers can achieve high-accuracy dialect normalization.

5.1 Limitations

This work has several limitations that should be acknowledged:

1. Synthetic data generation relies on commercial LLM APIs (Gemma-2 via Google AI Studio), introducing API cost and rate limit considerations.
2. While we benchmarked two encoder-decoder architectures (mT5-small and IndicBART), decoder-only LLMs (e.g., LLaMA-3, Gemma-2) have not been evaluated for zero-shot or fine-tuned text normalization.
3. Evaluation is currently focused on text normalization metrics (BLEU, chrF++, WER); large-scale downstream task evaluation across diverse domain applications remains for future work.

5.2 Data and Code Availability

The complete parallel corpus of 32,335 verified dialect–standard Marathi sentence pairs, along with all training code, evaluation scripts, and pre-trained model checkpoints, will be publicly released upon publication. The dataset will be available under an open license to support reproducible research in Indic dialect normalization. Our synthetic data generation pipeline, multi-tier verification engine, and fine-tuning configurations are documented in the accompanying code repository to enable replication and extension to additional Marathi sub-dialects and other low-resource Indic language varieties.

5.3 Future Directions

Several promising directions emerge from this work:

1. **Edge Deployment via Quantization**: Model quantization (INT8/INT4) and knowledge distillation to sub-100M parameter models for on-device deployment in rural Maharashtra.
2. **Expansion to Additional Dialects**: Extending the pipeline to Deshi/Puneri, Kolhapuri, Marathwada, and other Indic language sub-dialects.
3. **Decoder-Only LLM Benchmarking**: Evaluating instruction-tuned decoder-only models for zero-shot and few-shot dialect normalization tasks.

References

1. Amartyaveer, Kumar, S., Sharma, S., Udupa, S., Badiger, S., Singh, A., G, D., Bandekar, J., Murthy, S., Ghosh, P.K.: Improving dialect identification in indian languages using multimodal features from dialect informed asr. In: IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP). pp. 1–5. IEEE (2025). https://doi.org/10.1109/ICASSP49660.2025.10889099
2. Ansary, N., Adib, Q.A.R., Reasat, T., Sushmit, A.S., Humayun, A.I., Mehnaz, S., Fatema, K., Or Rashid, M.M., Sadeque, F.: Unicode normalization and grapheme parsing of indic languages. In: Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024). pp. 17019–17030. ELRA Language Resource Association (2024)
3. Blaschke, V., Purschke, C., Schütze, H., Plank, B.: What do dialect speakers want? a survey of attitudes towards language technology for german dialects. arXiv preprint arXiv:2402.11968 (2024), https://arxiv.org/abs/2402.11968
4. Dabre, R., Shrotriya, H., Kunchukuttan, A., Puduppully, R., Khapra, M.M., Kumar, P.: Indicbart: A pre-trained model for indic natural language generation. In: Findings of the Association for Computational Linguistics: ACL 2022. pp. 1849–1863 (2022), https://doi.org/10.18653/v1/2022.findings-acl.145
5. Dimakis, A., Pavlopoulos, J., Anastasopoulos, A.: Dialect normalization using large language models and morphological rules. In: Findings of the Association for Computational Linguistics: ACL 2025. pp. 23696–23714. Association for Computational Linguistics (2025)
6. Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Yang, A., Fan, A., et al.: The llama 3 herd of models. arXiv preprint arXiv:2407.21783 (2024), https://arxiv.org/abs/2407.21783
7. Gaikwad, G., Panchal, D., Deshmukh, O., Patil, S.: Marathi dialect to pure marathi using rag & ai. International Research Journal of Engineering and Technology (IRJET) **12**(11), 787–791 (2025)
8. Gavali, P.S.: Marathi dialect detector: Text and speech normalization to standard marathi and hindi. vol. 5, pp. 116–121 (2026), http://ijisem.com

9. Google DeepMind Gemma Team: Gemma 2: Improving open language models at a practical size. arXiv preprint arXiv:2408.00118 (2024), https://arxiv.org/abs/2408.00118
10. Habash, N., Diab, M., Rambow, O.: Conventional orthography for dialectal arabic. In: Proceedings of the Eighth International Conference on Language Resources and Evaluation (LREC '12). pp. 711–718 (2012)
11. Joshi, A., Dabre, R., Kanojia, D., Li, Z., Zhan, H., Haffari, G., Dippold, D.: Natural language processing for dialects of a language: A survey. ACM Computing Surveys **57**(6), 1–37 (2025), https://doi.org/10.1145/3712060
12. Joshi, P., Santy, S., Budhiraja, A., Bali, K., Choudhury, M.: The state and fate of linguistic diversity and inclusion in the NLP world. In: Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL). pp. 6282–6293 (2020). https://doi.org/10.18653/v1/2020.acl-main.560
13. Joshi, R., Goel, P., Joshi, R.: L3cube-mahanlp: Marathi natural language processing datasets, models, and library. In: Proceedings of the WILDRE-6 Workshop, LREC 2022. pp. 10–17 (2022), https://aclanthology.org/2022.wildre-1.2
14. Kresic, M., Abbas, N.: Normalizing swiss german dialects with the power of large language models. Procedia Computer Science **244**, 287–295 (2024)
15. Kulkarni-Joshi, S.: Variation and change in dialects of marathi: A social-dialectological approach. In: Chandra, P. (ed.) Variation in South Asian Languages: From Macro to Micro-Differences, pp. 207–236. Springer (2023). https://doi.org/10.1007/978-981-99-1149-3_9
16. Kumar, S., Singh, A., G, D., Amartyaveer, Bandekar, J., Murthy, S., Sharma, S., Badiger, S., Udupa, S., Nagireddi, A., Raghavan, S.K.M., Saxena, R., Nanavati, J., Nanavati, R., Sridharan, J., Mehta, A., Khuraishi, A.K.S., Mora, S.P.R., Venkataramakrishnan, P., Date, G., P, K., Ghosh, P.K.: RESPIN-S1.0: A read speech corpus of 10000+ hours in dialects of nine indian languages. In: Advances in Neural Information Processing Systems (NeurIPS 2025) Track on Datasets and Benchmarks. pp. 1–13 (2025)
17. Kuparinen, O., Miletić, A., Scherrer, Y.: Dialect-to-standard normalization: A large-scale multilingual evaluation. In: Findings of the Association for Computational Linguistics: EMNLP 2023. pp. 13814–13828. Association for Computational Linguistics (2023)
18. Lewis, M., Liu, Y., Goyal, N., Ghazvininejad, M., Mohamed, A., Levy, O., Stoyanov, V., Zettlemoyer, L.: Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In: Proceedings of ACL 2020. pp. 7871–7880 (2020), https://doi.org/10.18653/v1/2020.acl-main.703
19. Mulla, R., Kumar, B.S.: Text-independent automatic dialect recognition of marathi language using spectro-temporal characteristics of voice. International Journal on Recent and Innovation Trends in Computing and Communication (IJRITCC) **10**(2s), 313–321 (2022). https://doi.org/10.17762/ijritcc.v10i2s.5949
20. Partanen, N., Hämäläinen, M., Alnajjar, K.: Dialect text normalization to normative standard finnish. In: Proceedings of the 5th Workshop on Noisy User-generated Text (W-NUT @ EMNLP 2019). pp. 141–146. Association for Computational Linguistics (2019)
21. Pawar, P.S., Saini, J.R., Vaidya, S.: Automatic marathi dialect classification in noisy and non-noisy environments. Procedia Computer Science **282**, 288–303 (2026)
22. Post, M.: A call for clarity in reporting bleu scores. In: Proceedings of WMT 2018. pp. 186–191 (2018), https://doi.org/10.18653/v1/W18-6319

23. Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., Liu, P.J.: Exploring the limits of transfer learning with a unified text-to-text transformer. Journal of Machine Learning Research **21**(140), 1–67 (2020), https://jmlr.org/papers/v21/20-074.html
24. Tarfe, O.S., Bagul, M.: Comparative study of dialect (words) of marathi in konkan region (ratnagiri & malvan). Iconic Research and Engineering Journals (IRE Journals) **7**(10), 101–105 (2024)
25. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, Ł., Polosukhin, I.: Attention is all you need. In: Advances in Neural Information Processing Systems (NeurIPS 2017). pp. 5998–6008 (2017), https://arxiv.org/abs/1706.03762
26. Vyavhare, V.A., Divekar, S.N., Pawar, H., Kale, K., Jagtap, S.: Marathi dialect translator. International Journal of Scientific Research in Science, Engineering and Technology (IJSRSET) **13**(3), 254–261 (2026). https://doi.org/10.32628/IJSRSET2613331
27. Xue, L., Constant, N., Roberts, A., Kale, M., Al-Rfou, R., Siddhant, A., Barua, A., Raffel, C.: mt5: A massively multilingual pre-trained text-to-text transformer. In: Proceedings of NAACL-HLT 2021. pp. 483–498 (2021), https://doi.org/10.18653/v1/2021.naacl-main.41