Multi-Dialect Marathi Text Normalization via LLM-Driven Synthetic Augmentation

Aditya Kinjawadekar, Surender Kannaiyan, and Saugata Sinha

Visvesvaraya National Institute of Technology, Nagpur, India
kinjawadekaradi112@gmail.com

Abstract. **Keywords:** Dialect Normalization · Low-Resource NLP · Marathi Dialects · Sequence-to-Sequence Modeling · Synthetic Data Augmentation · mT5 · IndicBART · Text Standardization.

1 Introduction

Natural Language Processing (NLP) has achieved progress in recent years, which has allowed computers to analyze, interpret, and generate human language. Modern architectures for both, understanding and generation of natural language are based on models built upon Transformers architecture [32,28]. Large Language Models (LLMs) had been adpoted for various NLP tasks across diverse real-world problems in healthcare, finance, and legal infrastructure [7,10].

NLP tends to face several issues in incorporating low-resource languages, especially regional dialects of higher-resource languages [14,13]. Surveys among native dialect communities reveal that proper analysis and understanding of dialectal text remains a priority [4]. suggesting a need for converting non-standard dialect into a standard form while keeping the original meaning intact [20]. Advancing NLP methods have made it possible to design systems that seamlessly link previously inaccessible low-resource regional dialects.

India has 22 official languages with hundreds of regional dialects spoken by over 1.4 billion people. While languages such as Hindi, Bengali, Marathi, Tamil, Telugu, etc. have started receiving attention through initiatives like AI4Bharat IndicCorp and BPCC [5],sub-regional dialects remain almost entirely untouched in NLP. One of the examples of this linguistic ambiguity is Marathi, an Indo-Aryan language spoken by over 83 million people in the state of Maharashtra [5]. Marathi is written in Devanagari script from left to right and stands as one of India's recognized scheduled languages. While Standard Marathi (based on the dialect variety of Pune city and surrounding regions) has a well-defined literary standard and established digital infrastructure, regional sub-dialects tend to show phonological and lexical divergence from the standard form [18,31]. Users having a native dialect have faced difficulties in using the existing NLP systems of Automatic Speech Recognition (ASR), Machine Translation (MT), conversational AI agents, and information retrieval tools [16,1]. This tendency seems largely because of high Out-Of-Vocabulary (OOV) rates, non-standard verbal forms, and spelling variations [24,20,3]. *Dialect Normalization* is the task of converting non-standard dialectal text into normative standard written form while strictly preserving underlying semantic intent therefore becomes a critical prerequisite for digital inclusion.

Developing systems for dialect normalization have faced two fundamental challenges: (1) *Extreme Data Scarcity*: datasets connecting regional Indic dialects to standard written forms are non-existent; and (2) *Morphological Complexity*: Marathi is a morphologically rich language where case markers and verbal tense/aspect suffixes require fine-grained subword transformation rather than a naive lookup.

Addressing these challenges, this paper presents a multi-dialect Marathi text normalization pipeline combining synthetic data augmentation with sequence-to-sequence transformer fine-tuning. Our main contributions are :

- **Parallel Marathi Dialect Corpus with Synthetic Augmentation**: Based on the RESPIN corpus [19], we curate and expand a parallel corpus of verified sentence pairs across Southern Konkan, Northern Konkan, and Varhadi dialects mapped to Standard Marathi , combining human-translated pairs with an automated synthetic generation and multi-level verification pipeline.
- **End-to-end Fine-Tuning Pipeline**: We utilize a neural fine-tuning framework for dialect normalization, adapting pre-trained msequence-to-sequence transformer models for standardization.
- **Empirical Evaluations**: We implement a standardized train-test split combined with cross-validation across distinct dataset configurations, conducting experiments, verification filtering analyses, and error evaluations.

2 Related Work

2.1 Dialect Normalization and Text Standardization

Early works around Dialect Normalization includes Partanen et al. [24] who evaluated sequence transduction across 23 Finnish dialect varieties, demonstrating that sliding-window neural architectures (LSTMs and character Transformers) effectively normalize non-standard dialects into Standard Finnish, drastically reducing word error rates from an initial 52.89% down to 5.73%. Kuparinen et al. [20] extended this and covered four distinct language families (Finnish, Norwegian, Swiss German, and Slovene), showing that fine-tuned transformers and pre-trained byte-level architectures (byT5) consistently surpass character-level statistical machine translation (SMT) and neural baselines across diverse dialectal data. Similarly, The Conventional Orthography for Dialectal Arabic (CODA) framework [11] established baselines for mapping diverse regional Arabic varieties (Egyptian, Levantine, Gulf, Maghrebi) to Modern Standard Arabic (MSA) for corpus creation and Arabic dialect processing using computational methods. The MADAR project subsequently applied these techniques to fine-grained dialect identification and translation across 25 Arab cities. Early approaches heavily relied on character-level finite-state transducers (FSTs) and phrase-based

SMT rules, but present times have used pre-trained multilingual encoder-decoder transformers. In Germanic languages, Swiss German dialect normalization to Standard High German has been analysed due to the absence of standardized orthography for Swiss German dialects [29]. Recent work by Kresic and Abbas [17] demonstrated that fine-tuning multilingual transformers on the SwissDial corpus gave good results with the `mT5-small` model achieving performance close to larger base and large transformer variants. Further research into low-resource settings highlights the effectiveness of subword-level modeling (such as Byte-Pair Encoding) in handling the out-of-vocabulary (OOV) and unique spelling challenges inherent in highly variable, unstandardized user-generated dialectal text [**?**].

Dialect normalization remains a critically underexplored area for indic languages. Earlier work has addressed code-mixed language and script transliteration but explicit conversion of a dialectal text to standard written form within the same language has received comparatively less attention. The primary challenges identified in Indic dialect NLP include severe data scarcity, morpho-syntactic divergence, and orthographic inconsistency stemming from the lack of standardized writing systems for most dialects [**?**]. Recent benchmarking efforts, such as the INDIC-DIALECT dataset, have begun to systematically evaluate dialect classification and translation across regional varieties like Bhojpuri, Magahi, and Maithili into Standard Hindi, revealing significant performance gaps when using generalized models [**?**]. A comprehensive global survey by Joshi et al. [13] categorizes dialectal NLP tasks across Natural Language Understanding (NLU) and Natural Language Generation (NLG) across diverse language families, emphasizing the acute performance degradation of state-of-the-art models on non-standard dialects and calling for equitable language technologies. To address this, researchers have explored hybrid approaches. Recently, Dimakis et al. [6] introduced a hybrid normalization framework combining rule-based morphological transformations with LLM few-shot prompting to standardize low-resource Greek dialect proverbs without requiring prior parallel training data. Similarly, for Indic languages, Weighted Finite State Transducers (WFSTs) are increasingly being integrated alongside neural models to deterministically normalize semiotic entities (such as numbers and dates) before passing the text to deep learning backends [**?**].

In the specific context of Marathi language dialect processing, research has historically concentrated on speech, acoustic classification, or localized applications:

– **Acoustic Speech Classification & Recognition**: Mulla and Kumar [23] explored spectro-temporal feature extraction (15 parameters including chromagram, spectral centroid, and MFCCs) with Chi-Square/ANOVA feature selection on 7,555 speech recordings from the LDC-IL Marathi raw speech corpus across four dialects (Goa-based, Vidarbha, Marathwada, Puneri), achieving 84.24% accuracy with Ridge Classifiers. Pawar et al. [25] investigated acoustic dialect identification across Northern Konkan, Varhadi, and Standard Marathi audio using MFCC features and Convolutional Neural

Networks (CNNs) under clean and artificially generated Gaussian noise conditions (SNR 1.78–9.74 dB), achieving up to 93.65% accuracy in clean environments and 87.29% under noisy conditions. Amartyaveer et al. [2] and Kumar et al. [19] at IISc Bangalore introduced RESPIN-S1.0 (a 10,000+ hour read-speech corpus across 38 Indian dialects and 9 languages, including 4 Marathi sub-dialects: Southern Konkan, Northern Konkan, Standard Marathi, and Varhadi) and developed multimodal ASR and RoBERTa embeddings for dialect identification, achieving 79.81% accuracy.
– **Text-Level Dialect Processing & Translation**: Gaikwad et al. [8] explored zero-shot prompting via the commercial Gemini API within a Retrieval-Augmented Generation (RAG) framework using Sentence Transformers and FAISS to translate Konkani (Devanagari) dialect queries into Standard Marathi for document Q&A. Vyavhare et al. [33] proposed a hybrid rule-based dictionary and neural translation web interface for regional Marathi dialects. Gavali [9] introduced a student-level framework combining IndicBERT and IndicWav2Vec for dialect detection and normalization on small local datasets.

While prior work focused either on speech/acoustic classification or zero-shot prompting , we establish the first large-scale parallel-pair text benchmark, introduce a synthetic data expansion engine powered by Llama-3.1-8B [7] and Gemma-4 [10]) with multi-tier verification, and conduct 5-fold cross-validated fine-tuning of open-weight transformers (`mT5-small` and IndicBART) for local, reproducible inference.

2.2 Indic NLP and Sequence-to-Sequence Transformers

The development of pre-trained models for Indic languages has shown steep increase in recent years. IndicBART [5] introduced a multilingual sequence-to-sequence model based on the mBART-50 architecture [21],trained on 11 Indic languages with a unified Devanagari script representation with a 64,000-token SentencePiece vocabulary. It utilizes orthographic relations between Indic scripts to facilitate cross-lingual transfer learning and demonstrated state-of-the-art performance on neural machine translation tasks despite being significantly smaller than models of similar category.

Similarly ,models such as mT5 [35] which is the multilingual extension of T5 [28] extended Text-to-Text Transfer Transformers across 101 languages with the mC4 dataset. which features a 250,000-token SentencePiece vocabulary and showed state-of-the-art results on the XTREME multilingual benchmark. Its text-to-text formulation, which casts all NLP tasks as input-output text transformations, makes it particularly well-suited for sequence-level normalization tasks. The L3Cube-MahaNLP initiative [15] gave Marathi-specific transformer models (MahaBERT, MahaRoBERTa) and curated datasets for tasks like sentiment analysis, named entity recognition, and hate speech detection. However, these encoder-only models are primarily designed for classification-like tasks instead of sequence-to-sequence generation.

2.3 Synthetic Data Generation for Low-Resource NLP

Using LLMs for synthetic data generation has proven to be a useful strategy for solving data scarcity in low-resource settings. Recent work has shown that LLMs can be used more effectively as data generators than as direct task solvers, where synthetic data from a large teacher model is utilized to train smaller models [34,30,22]. This paradigm is often regarded as a form of knowledge distillation that tries to compensate the costs and time bottlenecks associated with human annotation procedures [12]. By the use of appropriate prompt-engineering techniques, frontier models can generate data that effectively transfers their capabilities to specialized and efficient student models [26]. For Indic languages specifically, synthetic data augmentation has been applied primarily to machine translation tasks, with back-translation and pivoting through high-resource languages serving as common strategies.

Our work extends this paradigm specifically to the dialect normalization task, introducing domain-aware generation and multi-stage LLM verification tailored to Marathi dialects.

3 Methodology

3.1 Target Dialects and Linguistic Properties

We target the following major Marathi dialects that span the primary regions of Maharashtra and exhibit distinct linguistic and phonetic patterns:

1. **Southern Konkan (Konkan Coast)**: Characterized by shortening of final-word vowels, softening of dental consonants, and different auxiliary verb forms. Lexical differences include systematic vowel shifts: standard जातो ("he goes") becomes Southern Konkan जातां; standard कुठे ("where") becomes कुठंय. Domain-specific terminology generally remains untouched but often accompanied by Konkani words.
2. **Northern Konkan (Khandesh Region)**: Heavily influenced by neighboring Gujarati and Bhili language families,thus it exhibits significant consonant simplification and dental consonant shifts. Notabel transformations include: standard कशाला ("why") → Northern Konkan कसाले; standard मुलगा ("boy") → Northern Konkan पोर. Northern Konkan posseses the highest degree of lexical divergence among the three target dialects.
3. **Varhadi (Vidarbha/Eastern Maharashtra)**: Spoken in Vidarbha, It exhibits characteristic differences and verb suffix modifications. Standard interrogative का? ("why?") becomes Varhadi काऊन?; verbal forms undergo regular suffix mapping patterns. Among the three dialects, Varhadi maintains closest structural similarity to Standard Marathi as transformations are more systematic.

Table 1. RESPIN-S1.0 Marathi Source Corpus Statistics

Statistic	**Value**
Total Utterances	809,934
Total Duration	1,026.28 hours
Unique Speakers	2,644
Unique Pincodes	234
Gender Distribution	51% Female, 49% Male
Domains	Agriculture, Banking
Domain Split	49% Agriculture, 51% Banking
Dialect Variants	Southern Konkan, Northern Konkan, Varhadi

3.2 Original Corpus Curation and Foundation in RESPIN

Base sentences for our task are drawn from the domain-specific transcripts of the RESPIN-S1.0 initiative [19], developed at IISc Bangalore under the Gates Foundation project. While RESPIN has audio recordings and their transcripts across agriculture and banking domains for Indian languages and their dialects, it lacks aligned parallel dialect-to-target language translations required for efficient normalisation.

The large speaker diversity ensures broad coverage of intra-dialect phonetic and lexical variation, while the balanced gender and domain splits reduce systematic bias in our derived parallel corpus.

3.3 LLM-Driven Parallel Corpus Genaration and Verification

To deal with data scarcity, we developed an automated synthetic data generation pipeline powered by Gemma-4 [10] in 9B-instruction-tuned (`gemma-2-9b-it`) configurations, accessed via Ollama. We considered several design decisions to ensure generation quality to be maximised:

- **1-Prompt-Per-Sentence Architecture**: Instead of generating batches of parallel pairs per API call (which risks array length mismatches and context contamination), each dialectal sentence is processed separately with a self-contained prompt.
- **Detailed Prompting Strategy**: The generation prompt enforces semantic fidelity rules: the standard output must preserve exact meaning, domain terminology (agriculture/banking), named entities, and numerical values from the input. The prompt includes dialect-specific lexical mapping tables and 5-8 few-shot examples per dialect.
- **Temperature Control**: Generation temperature is set to 0.3 with top-p = 0.9 to balance diversity with fidelity.

All generated pairs undergo a three-tier verification and correction pipeline:

1. **Tier 1: Rule-Based Syntactic Filter.** Automated validation checks implement Indic Abugida script normalization principles [3], including:

Table 2. Parallel Dataset Composition and Yield Across All Three Dialect Suites

Dialect Name	Original Pairs	Synthetic Pairs	Total Pairs
Southern Konkan	5,838	5,576	**11,414**
Northern Konkan	5,825	5,501	**11,326**
Varhadi	5,330	5,086	**10,416**
Total (All 3 Dialects)	**16,993**	**16,163**	**33,156**

- *Devanagari Script Integrity*: Both source and target must possess >90% Devanagari Unicode characters within the `U+0900--U+097F` range, automatically discarding unnecessary characters.
- *Length Ratio Constraint*: $0.5 \leq |L_{\text{dialect}}|/|L_{\text{standard}}| \leq 2.0$ to avoid hallucinations.
- *Non-Identity Constraint*: $L_{\text{dialect}} \neq L_{\text{standard}}$ to exclude identical string copies.
- *Minimum Token Count*: Both source and target must contain ≥ 3 whitespace-separated tokens.

2. **Tier 2: LLM Semantic Verification.** A secondary LLM evaluator using `llama-3.1-8b-instant` (accessed via Groq API) performs semantic equivalence verification. The verifier rates on a 1-5 score scale considering: (a) semantic preservation; (b) absence of hallucinated entities; (c) absence of Hindi/English contamination; and (d) complete standardization of dialects. Pairs scoring < 4.0 are rejected or routed to correction.
3. **Closed-Loop Correction Engine.** Pairs flagged by Tier 1 or Tier 2 enter a two-step correction loop: Step 1 generates corrected translations using specialised prompts highlighting the specific failure reason (e.g., "Hindi leakage detected" or "Incomplete normalization"); Step 2 verifies the corrected sentences again,completely discarding pairs that fail a second audit pass.

As mentioned in Table 2, the synthetic augmentation pipeline contributed 16,163 high-quality verified parallel pairs, almost doubling the original dataset to produce an expanded version of 33,156 samples. The near-1:1 ratio between original and synthetic pairs across all three dialects demonstrates the consistency and reliability of the entire pipeline.

3.4 Evaluation Protocol

To verify model training and ensure hyperparameter stability during development, models were initially validated using 5-Fold Cross-Validation on an 85/15 stratified split of the training pool. For final benchmark reporting, models were trained on the complete dataset (combining original and synthetically expanded pairs) and evaluated on the official held-out RESPIN test set (`IISc_RESPIN_test_mr`, comprising 2,170 utterances across target dialects).

This setup is applied across all dataset configurations: Southern Konkan (Original), Northern Konkan (Original), Varhadi (Original), Combined-data (Original), and their Synthetically Expanded counterparts. All results are reported

Table 3. Comprehensive Evaluation Matrix on `IISc_RESPIN_test_mr` (2,170 Utterances)

Model Scope	Training Data	Dialect	Utts	IndicBART (244M)		mT5-Small (300M)	
				BLEU (↑)	WER (↓)	BLEU (↑)	WER (↓)
Single-Dialect	Original	Southern Konkan	559	**57.80**	**26.60%**	43.86	34.37%
		Northern Konkan	540	**90.13**	**6.46%**	79.46	11.95%
		Varhadi	516	**83.59**	**10.19%**	74.81	14.73%
	Expanded	Southern Konkan	559	24.06	60.32%	**44.36**	**32.75%**
		Northern Konkan	540	65.41	28.70%	**79.76**	**12.61%**
		Varhadi	516	67.97	25.43%	**76.90**	**14.19%**
Multi-Dialect	Original	Southern Konkan	559	–	–	40.94	35.34%
		Northern Konkan	540	–	–	**77.50**	**14.65%**
		Varhadi	516	–	–	**73.47**	**16.16%**
		Combined (All)	2,170	62.98	30.13%	**72.18**	**17.23%**
	Expanded	Southern Konkan	559	–	–	43.88	34.96%
		Northern Konkan	540	–	–	**78.16**	**13.55%**
		Varhadi	516	–	–	**74.74**	**15.57%**
		Combined (All)	2,170	**76.50**	**16.23%**	73.48	16.58%

using standard automatic metrics computed via SacreBLEU [27] and Word Error Rate (WER) scripts:

- **BLEU** (Bilingual Evaluation Understudy): BLEU-4 with 4-gram precision, brevity penalty, and smoothing method "exp".
- **chrF++**: Character n-gram F-score combined with word unigrams and bigrams, capturing subword morphological variations in Marathi text.
- **Normalized WER / CER**: Word Error Rate and Character Error Rate computed after Devanagari Unicode script standardization.

4 Findings

4.1 Comprehensive Test Benchmark on `IISc_RESPIN_test_mr`

Key observations from Table 3 include:

1. **mT5-Small Performance**: The synthetically expanded fine-tuned `mT5-small` model achieves 73.48 BLEU and 84.58 chrF++ on the overall test set, demonstrating robust multi-dialect transfer learning.
2. **IndicBART Scaling**: Fine-tuning IndicBART on the synthetically expanded multi-dialect corpus achieves **76.50 BLEU** and **85.90 chrF++**, demonstrating that synthetic data expansion activates latent capacity in mBART-50 architectures.

4.2 Dialectwise Relative WER Reductions vs. ASR Baseline

Table 4 and Table 5 compare the Word Error Rate of the baseline ASR system against fine-tuned `IndicBART` and `mT5-small` models across sub-dialects.

Table 4. Dialectwise Relative WER Reductions: ASR Baseline vs. Fine-Tuned IndicBART (Synthetically Expanded Data)

Sub-Dialect	**ASR Baseline WER (%)**	**IndicBART WER (%)**	**Relative WER Drop (%)**
Southern Konkan	41.80	49.15	+17.59
Northern Konkan	39.53	52.54	+32.91
Varhadi	35.94	37.74	+4.99
Multi-Dialect (All)	**39.41**	**27.16**	**−31.07**

Table 5. Dialectwise Relative WER Reductions: ASR Baseline vs. Fine-Tuned mT5-Small (Synthetically Expanded Data)

Sub-Dialect	**ASR Baseline WER (%)**	**mT5-Small WER (%)**	**Relative WER Drop (%)**
Southern Konkan	41.80	30.37	−27.33
Northern Konkan	39.53	29.45	−25.50
Varhadi	35.94	27.16	−24.43
Multi-Dialect (All)	**39.41**	**21.43**	**−45.62**

Both fine-tuned text normalizers achieve substantial relative WER reductions over initial ASR output, with `mT5-small` consistently outperforming `IndicBART`:

- **Northern Konkan**: `mT5-small` reduces ASR WER from 39.53% to **29.45%** (−**25.50% relative reduction**), outperforming `IndicBART` (which suffered a +32.91% WER increase).
- **Varhadi**: `mT5-small` reduces WER from 35.94% to **27.16%** (−**24.43% relative reduction**), outperforming `IndicBART` (which suffered a +4.99% WER increase).
- **Southern Konkan**: `mT5-small` reduces WER from 41.80% to **30.37%** (−**27.33% relative reduction**), outperforming `IndicBART` (which suffered a +17.59% WER increase).
- **Multi-Dialect (All)**: `mT5-small` reduces overall WER from 39.41% to **21.43%** (−**45.62% relative reduction**), outperforming `IndicBART` (27.16% WER, −31.07% reduction).

4.3 Benchmark Results on Original Datasets

Table 6 presents the performance matrix on 5-fold cross-validated original single-dialect benchmarks, and Table 7 details the dialectwise evaluation breakdown of the multi-dialect model on original data.

4.4 Benchmark Results on Synthetically Expanded Datasets

Table 8 presents the performance matrix on synthetically expanded single-dialect benchmarks, and Table 9 details the dialectwise evaluation breakdown of the multi-dialect model on expanded data.

Table 6. Benchmark Performance Matrix on Original Datasets: Single-Dialect Models

Dialect Name	Total Pairs	Test Set	IndicBART (244M)		mT5-Small (300M)	
			BLEU	chrF++	BLEU	chrF++
Southern Konkan	5,576	836	48.52	78.49	46.31	74.43
Northern Konkan	5,501	825	47.29	66.84	**60.31**	**77.34**
Varhadi	5,086	762	72.87	85.72	**80.99**	**91.21**

Table 7. Benchmark Performance Matrix on Original Datasets: Multi-Dialect Dialect-wise Breakdown

Evaluated Subset	Total Pairs	Test Set	IndicBART (244M)		mT5-Small (300M)	
			BLEU	chrF++	BLEU	chrF++
Southern Konkan	16,163	836	–	–	–	–
Northern Konkan	16,163	798	–	–	–	–
Varhadi	16,163	790	–	–	–	–
Combined (All)	16,163	2,424	–	–	**63.29**	**81.37**

4.5 Impact of Synthetic Data Augmentation

Table 10 isolates the effect of dataset scaling by comparing performance across original and synthetically expanded corpora for each model. IndicBART benefits substantially from augmentation on the multi-dialect task (+13.52 BLEU), confirming that data scarcity rather than model capacity was the primary bottleneck. mT5-small shows a more modest improvement (+1.30 BLEU on multi-dialect data) but starts from a much higher baseline.

4.6 Verification Engine Impact: Raw Unverified vs. Filtered Clean Data

To quantify the explicit impact of our multi-tier LLM verification and filtering engine (`filter_flawed.py` and `llm_verifier.py`), we conduct a comparative study evaluating models trained on *Raw Unverified Synthetic Data* versus *Filtered Clean Synthetic Data* (the expanded suite) across 5-fold cross-validation. The raw unverified dataset contains 19,914 parallel pairs, including 3,428 rejected or flawed pairs exhibiting Hindi language leakage, retained dialect residue, and garbled translations prior to multi-tier verification.

Table 11 presents the quantitative comparison on the benchmark suite.

Key findings and insights from this evaluation include:

1. **Quality Collapse on Unverified Data**: Training IndicBART on unverified raw LLM outputs causes a catastrophic 14.03 BLEU point drop (collapsing from 57.12 to 43.09) due to severe error propagation from uncorrected dialect residue, garbled syntax, and script errors.

Table 8. Benchmark Performance Matrix on Synthetically Expanded Datasets: Single-Dialect Models

			IndicBART (244M)		mT5-Small (300M)	
Dialect Name	**Total Pairs**	**Test Set**	**BLEU**	**chrF++**	**BLEU**	**chrF++**
Southern Konkan	11,145	1,671	47.25	64.69	**65.10**	**82.88**
Northern Konkan	11,035	1,655	40.51	62.21	**62.07**	**79.29**
Varhadi	10,155	1,523	73.62	86.75	**78.89**	**90.59**

Table 9. Benchmark Performance Matrix on Synthetically Expanded Datasets: Multi-Dialect Dialectwise Breakdown

			IndicBART (244M)		mT5-Small (300M)	
Evaluated Subset	**Total Pairs**	**Test Set**	**BLEU**	**chrF++**	**BLEU**	**chrF++**
Southern Konkan	32,335	1,710	–	–	–	–
Northern Konkan	32,335	1,627	–	–	–	–
Varhadi	32,335	1,513	–	–	–	–
Combined (All)	32,335	4,850	–	–	**69.51**	**84.58**

2. **mT5-Small Resilience**: While google/mT5-small is more resilient due to its 250,000-token multilingual vocabulary, training on unverified data still suffers a heavy 8.30 BLEU point drop (falling from 69.51 to 61.21) and a 4.40 chrF++ drop.
3. **Essential Role of Multi-Tier Auditing**: Removing flawed pairs via `filter_flawed.py` and correcting recoverable generations via `correct_flawed.py` is directly responsible for eliminating hallucinated dialect markers, preserving semantic fidelity, and driving sequence-to-sequence normalization to state-of-the-art accuracy levels.

4.7 Cross-Validation Stability Analysis

The extremely low standard deviation in test BLEU (0.13) and chrF++ (0.05) across folds demonstrates that model performance is robust and not an artifact of a single favorable train-test partition. Per-dialect fold-level analysis further confirms stable per-dialect breakdown: Southern Konkan (66.62 $\pm$ 0.18 BLEU), Northern Konkan (62.72 $\pm$ 0.19 BLEU), and Varhadi (78.73 $\pm$ 0.11 BLEU), reinforcing the consistent difficulty hierarchy Varhadi $<$ Northern Konkan $<$ Southern Konkan across all folds.

Critically, cross-dialect joint training in the synthetically expanded multi-dialect configuration provides a significant boost to IndicBART's per-dialect sub-scores: Southern Konkan reaches 52.15 BLEU (+3.63 over single-dialect IndicBART model) and Northern Konkan reaches 54.35 BLEU (+7.06 over single-

Table 10. Impact of Synthetic Data Augmentation: Original vs. Synthetically Expanded Data BLEU Change

Evaluated Subset	IndicBART BLEU			mT5-Small BLEU		
	Original	**Expanded**	**Δ**	**Original**	**Expanded**	**Δ**
Southern Konkan	–	–	–	40.94	**43.88**	+2.94
Northern Konkan	–	–	–	77.50	**78.16**	+0.66
Varhadi	–	–	–	73.47	**74.74**	+1.27
Multi-Dialect (All)	**62.98**	**76.50**	**+13.52**	**72.18**	**73.48**	**+1.30**

Table 11. Impact of Multi-Tier Verification Engine: Raw Unverified vs. Filtered Clean Synthetic Data

Model	Pipeline State	Total Pairs	BLEU	chrF++	Δ BLEU	Δ chrF++
IndicBART (244M)	Raw Unverified Data	19,914	43.09	64.54	−14.03	−10.95
	Filtered Clean Data	32,335	**57.12**	**75.49**	Baseline	Baseline
mT5-Small (300M)	Raw Unverified Data	19,914	61.21	80.18	−8.30	−4.40
	Filtered Clean Data	32,335	**69.51**	**84.58**	Baseline	Baseline

dialect model), suggesting that cross-dialect transfer learning benefits models with smaller, more linguistically focused vocabularies.

An important observation is that IndicBART's single-dialect performance *decreases* for Southern Konkan and Northern Konkan when moving from original to synthetically expanded data, while mT5 improves dramatically. This suggests that the synthetic data distribution may slightly deviate from the original distribution in ways that IndicBART's narrower vocabulary cannot accommodate, while mT5's broader subword coverage absorbs the distributional shift effectively.

4.8 Model Architecture Analysis

Table 12 provides a detailed architectural comparison between the two primary models. The 69.51 vs. 57.12 BLEU gap on the expanded multi-dialect task (+12.39 absolute) can be attributed to three key architectural differences:

1. **Vocabulary Coverage**: mT5's 250K vocabulary provides 3.9× more subword units than IndicBART's 64K, enabling finer-grained tokenization of morphologically complex Marathi dialectal forms. This is critical for dialect normalization, where subtle suffix and infix changes must be captured at the subword level.
2. **Positional Encoding**: mT5 uses relative positional biases, which generalize better to variable-length sequences typical of dialectal text where sentence length may differ substantially between dialect and standard forms.

Table 12. Detailed Architecture Comparison: IndicBART (244M) vs. mT5-small (300M)

Property	**IndicBART (244M)**	**mT5-small (300M)**
Base Architecture	mBART-50 (BART Enc-Dec)	T5 Encoder-Decoder
Position Encoding	Absolute sinusoidal	Relative positional bias
Vocabulary Size	64,000 SentencePiece	250,000 SentencePiece (mC4)
Pre-training Data	11 Indic + English	101 languages (mC4)
Script Approach	Devanagari-unified	Native multi-script
Target Conditioning	Required `<2mr>` token	Raw text-to-text (no prefix)
Learning Rate	5×10^{-5} (linear warmup)	5×10^{-4} (AdaFactor)
Expanded BLEU	57.12	**69.51** (+12.39)
Expanded chrF++	75.49	**84.58** (+9.09)

3. **Conditioning Overhead**: IndicBART requires explicit language tokens (`<2mr>`) that add conditioning complexity, whereas mT5's pure text-to-text format enables direct mapping without task-specific architectural constraints.

5 Conclusion

This study presents a comprehensive, reproducible framework for multi-dialect Marathi text normalization, addressing severe data scarcity through automated synthetic data expansion and establishing rigorous neural sequence-to-sequence benchmarks.

Our LLM-driven synthetic data augmentation pipeline, powered by Gemma-2 with multi-tier verification (Devanagari script filtering and LLaMA-3.1-8B semantic auditing), successfully doubled the clean parallel corpus through synthetic expansion while maintaining strict semantic fidelity—a scalable paradigm for other low-resource dialect pairs. Fine-tuning sequence-to-sequence transformers on the expanded benchmark established state-of-the-art performance: google/mT5-small (300M) achieved 69.51 BLEU and 84.58 chrF++ on the synthetically expanded multi-dialect dataset, outperforming IndicBART (244M) by +12.39 BLEU. Individual dialect performance reached 80.99 BLEU and 91.21 chrF++ for Varhadi, demonstrating that compact pre-trained sequence-to-sequence transformers can achieve high-accuracy dialect normalization.

5.1 Limitations

This work has several limitations that should be acknowledged:

1. Synthetic data generation relies on commercial LLM APIs (Gemma-2 via Google AI Studio), introducing API cost and rate limit considerations.

2. While we benchmarked two encoder-decoder architectures (mT5-small and IndicBART), decoder-only LLMs (e.g., LLaMA-3, Gemma-2) have not been evaluated for zero-shot or fine-tuned text normalization.
3. Evaluation is currently focused on text normalization metrics (BLEU, chrF++, WER); large-scale downstream task evaluation across diverse domain applications remains for future work.

5.2 Data and Code Availability

The complete parallel corpus of 32,335 verified dialect–Standard Marathi sentence pairs, along with all training code, evaluation scripts, and pre-trained model checkpoints, will be publicly released upon publication. The dataset will be available under an open license to support reproducible research in Indic dialect normalization. Our synthetic data generation pipeline, multi-tier verification engine, and fine-tuning configurations are documented in the accompanying code repository to enable replication and extension to additional Marathi sub-dialects and other low-resource Indic language varieties.

5.3 Future Directions

Several promising directions emerge from this work:

1. **Edge Deployment via Quantization**: Model quantization (INT8/INT4) and knowledge distillation to sub-100M parameter models for on-device deployment in rural Maharashtra.
2. **Expansion to Additional Dialects**: Extending the pipeline to Deshi/Puneri, Kolhapuri, Marathwada, and other Indic language sub-dialects.
3. **Decoder-Only LLM Benchmarking**: Evaluating instruction-tuned decoder-only models for zero-shot and few-shot dialect normalization tasks.

References

1. Aharoni, R., Goldberg, Y.: Unsupervised domain adaptation for neural machine translation with domain-aware feature embeddings. In: Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP). pp. 8794–8804 (2020)
2. Amartyaveer, Kumar, S., Sharma, S., Udupa, S., Badiger, S., Singh, A., G, D., Bandekar, J., Murthy, S., Ghosh, P.K.: Improving dialect identification in indian languages using multimodal features from dialect informed asr. In: IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP). pp. 1–5. IEEE (2025). https://doi.org/10.1109/ICASSP49660.2025.10889099
3. Ansary, N., Adib, Q.A.R., Reasat, T., Sushmit, A.S., Humayun, A.I., Mehnaz, S., Fatema, K., Or Rashid, M.M., Sadeque, F.: Unicode normalization and grapheme parsing of indic languages. In: Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024). pp. 17019–17030. ELRA Language Resource Association (2024)

4. Blaschke, V., Purschke, C., Schütze, H., Plank, B.: What do dialect speakers want? a survey of attitudes towards language technology for german dialects. arXiv preprint arXiv:2402.11968 (2024), https://arxiv.org/abs/2402.11968
5. Dabre, R., Shrotriya, H., Kunchukuttan, A., Puduppully, R., Khapra, M.M., Kumar, P.: Indicbart: A pre-trained model for indic natural language generation. In: Findings of the Association for Computational Linguistics: ACL 2022. pp. 1849–1863 (2022), https://doi.org/10.18653/v1/2022.findings-acl.145
6. Dimakis, A., Pavlopoulos, J., Anastasopoulos, A.: Dialect normalization using large language models and morphological rules. In: Findings of the Association for Computational Linguistics: ACL 2025. pp. 23696–23714. Association for Computational Linguistics (2025)
7. Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Yang, A., Fan, A., et al.: The llama 3 herd of models. arXiv preprint arXiv:2407.21783 (2024), https://arxiv.org/abs/2407.21783
8. Gaikwad, G., Panchal, D., Deshmukh, O., Patil, S.: Marathi dialect to pure marathi using rag & ai. International Research Journal of Engineering and Technology (IRJET) **12**(11), 787–791 (2025)
9. Gavali, P.S.: Marathi dialect detector: Text and speech normalization to standard marathi and hindi. vol. 5, pp. 116–121 (2026), http://ijisem.com
10. Google DeepMind Gemma Team: Gemma 2: Improving open language models at a practical size. arXiv preprint arXiv:2408.00118 (2024), https://arxiv.org/abs/2408.00118
11. Habash, N., Diab, M., Rambow, O.: Conventional orthography for dialectal arabic. In: Proceedings of the Eighth International Conference on Language Resources and Evaluation (LREC '12). pp. 711–718 (2012)
12. Hsieh, C.Y., et al.: Distilling step-by-step! outperforming larger language models with less training data and smaller model sizes. arXiv preprint arXiv:2305.02301 (2023)
13. Joshi, A., Dabre, R., Kanojia, D., Li, Z., Zhan, H., Haffari, G., Dippold, D.: Natural language processing for dialects of a language: A survey. ACM Computing Surveys **57**(6), 1–37 (2025), https://doi.org/10.1145/3712060
14. Joshi, P., Santy, S., Budhiraja, A., Bali, K., Choudhury, M.: The state and fate of linguistic diversity and inclusion in the NLP world. In: Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL). pp. 6282–6293 (2020). https://doi.org/10.18653/v1/2020.acl-main.560
15. Joshi, R., Goel, P., Joshi, R.: L3cube-mahanlp: Marathi natural language processing datasets, models, and library. In: Proceedings of the WILDRE-6 Workshop, LREC 2022. pp. 10–17 (2022), https://aclanthology.org/2022.wildre-1.2
16. Koenecke, A., Nam, A., Lake, E., Nudell, J., Quartey, M., Mengesha, Z., Toups, C., Rickford, J.R., Jurafsky, D., Goel, S.: Racial disparities in automated speech recognition. Proceedings of the National Academy of Sciences **117**(14), 7684–7689 (2020)
17. Kresic, M., Abbas, N.: Normalizing swiss german dialects with the power of large language models. Procedia Computer Science **244**, 287–295 (2024)
18. Kulkarni-Joshi, S.: Variation and change in dialects of marathi: A social-dialectological approach. In: Chandra, P. (ed.) Variation in South Asian Languages: From Macro to Micro-Differences, pp. 207–236. Springer (2023). https://doi.org/10.1007/978-981-99-1149-3_9
19. Kumar, S., Singh, A., G, D., Amartyaveer, Bandekar, J., Murthy, S., Sharma, S., Badiger, S., Udupa, S., Nagireddi, A., Raghavan, S.K.M., Saxena, R., Nanavati, J., Nanavati, R., Sridharan, J., Mehta, A., Khuraishi, A.K.S., Mora, S.P.R.,

Venkataramakrishnan, P., Date, G., P, K., Ghosh, P.K.: RESPIN-S1.0: A read speech corpus of 10000+ hours in dialects of nine indian languages. In: Advances in Neural Information Processing Systems (NeurIPS 2025) Track on Datasets and Benchmarks. pp. 1–13 (2025)

20. Kuparinen, O., Miletić, A., Scherrer, Y.: Dialect-to-standard normalization: A large-scale multilingual evaluation. In: Findings of the Association for Computational Linguistics: EMNLP 2023. pp. 13814–13828. Association for Computational Linguistics (2023)
21. Lewis, M., Liu, Y., Goyal, N., Ghazvininejad, M., Mohamed, A., Levy, O., Stoyanov, V., Zettlemoyer, L.: Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In: Proceedings of ACL 2020. pp. 7871–7880 (2020), https://doi.org/10.18653/v1/2020.acl-main.703
22. Mukherjee, S., Mitra, A., Jawahar, G., Agarwal, S., Palangi, H., Awadallah, A.: Orca: Progressive learning from complex explanation traces of gpt-4 (2023)
23. Mulla, R., Kumar, B.S.: Text-independent automatic dialect recognition of marathi language using spectro-temporal characteristics of voice. International Journal on Recent and Innovation Trends in Computing and Communication (IJRITCC) **10**(2s), 313–321 (2022). https://doi.org/10.17762/ijritcc.v10i2s.5949
24. Partanen, N., Hämäläinen, M., Alnajjar, K.: Dialect text normalization to normative standard finnish. In: Proceedings of the 5th Workshop on Noisy User-generated Text (W-NUT @ EMNLP 2019). pp. 141–146. Association for Computational Linguistics (2019)
25. Pawar, P.S., Saini, J.R., Vaidya, S.: Automatic marathi dialect classification in noisy and non-noisy environments. Procedia Computer Science **282**, 288–303 (2026)
26. Pecher, B., et al.: Better as generators than classifiers: Leveraging llms and synthetic data for low-resource multilingual classification. In: Findings of the Association for Computational Linguistics: EACL 2024 (2024)
27. Post, M.: A call for clarity in reporting bleu scores. In: Proceedings of WMT 2018. pp. 186–191 (2018), https://doi.org/10.18653/v1/W18-6319
28. Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., Liu, P.J.: Exploring the limits of transfer learning with a unified text-to-text transformer. Journal of Machine Learning Research **21**(140), 1–67 (2020), https://jmlr.org/papers/v21/20-074.html
29. Samardžić, T., Scherrer, Y., Glaser, E.: Archimob-a corpus of spoken swiss german. In: Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC'16). pp. 4061–4066 (2016)
30. Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., Liang, P., Hashimoto, T.B.: Alpaca: A strong, replicable instruction-following model. Stanford Center for Research on Foundation Models (2023)
31. Tarfe, O.S., Bagul, M.: Linguistic diversity of marathi in maharashtra: Review article. Iconic Research and Engineering Journals (IRE Journals) **7**(10), 91–96 (2024)
32. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, Ł., Polosukhin, I.: Attention is all you need. In: Advances in Neural Information Processing Systems (NeurIPS 2017). pp. 5998–6008 (2017), https://arxiv.org/abs/1706.03762
33. Vyavhare, V.A., Divekar, S.N., Pawar, H., Kale, K., Jagtap, S.: Marathi dialect translator. International Journal of Scientific Research in Science, Engineering

and Technology (IJSRSET) **13**(3), 254–261 (2026). https://doi.org/10.32628/IJSRSET2613331

34. Wang, Y., Kordi, Y., Mishra, S., Liu, A., Smith, N.A., Khashabi, D., Hajishirzi, H.: Self-instruct: Aligning language models with self-generated instructions (2022)
35. Xue, L., Constant, N., Roberts, A., Kale, M., Al-Rfou, R., Siddhant, A., Barua, A., Raffel, C.: mt5: A massively multilingual pre-trained text-to-text transformer. In: Proceedings of NAACL-HLT 2021. pp. 483–498 (2021), https://doi.org/10.18653/v1/2021.naacl-main.41