

Multi-Dialect Marathi Text Normalization via LLM-Driven Synthetic Augmentation and Neural Sequence-to-Sequence Benchmarking

Aditya Kinjawadekar, Surender Kannaiyan, and Saugata Sinha

Visvesvaraya National Institute of Technology, Nagpur, India
kinjawadekaradi112@gmail.com

Abstract. Low-resource regional dialects pose severe challenges for Natural Language Processing (NLP) systems due to significant morpho-syntactic variation, non-standard orthography, and extreme scarcity of annotated parallel corpora. In this paper, we address the task of dialect normalization—converting non-standard regional dialectal text into Standard Written Marathi—focusing on three major Indo-Aryan regional dialects of Maharashtra: Malvani (D1), Ahirani (D2), and Varhadi (D4). To overcome severe data scarcity, we introduce an automated synthetic parallel data expansion pipeline combining large language model (LLM) generation using Gemma-2 (9B/27B) with a multi-tier verification engine (rule-based script filtering and LLM semantic auditing), effectively doubling the clean parallel training corpus from 16,163 to 32,335 verified pairs across agriculture and finance domains. We conduct a rigorous empirical evaluation of state-of-the-art multilingual sequence-to-sequence transformer architectures—google/mT5-small (300M) and IndicBART (244M)—under an 85/15 stratified test split with 5-fold cross-validation across 8 dataset configurations. Empirical results demonstrate that mT5-small achieves state-of-the-art performance with 69.51 BLEU and 84.58 chrF++ on the 32k combined multi-dialect dataset (+12.39 BLEU over IndicBART), with individual dialect performance reaching 80.99 BLEU and 91.21 chrF++ on Varhadi (D4). We present detailed architectural comparisons, data scaling ablation studies, and qualitative subword morphological error analyses, establishing a benchmark for Indic dialect text normalization.

Keywords: Dialect Normalization · Low-Resource NLP · Marathi Dialects · Sequence-to-Sequence Modeling · Synthetic Data Augmentation · mT5 · IndicBART · Text Standardization.

1 Introduction

Natural Language Processing (NLP) is a branch of artificial intelligence that has achieved remarkable advances in recent years, driven by the emergence of large-scale pre-trained Transformer models [23,21]. However, a significant digital divide persists: the vast majority of the world’s 7,000+ languages remain critically

underserved in computational resources, annotated datasets, and pre-trained model coverage. Among these, regional dialects of major languages present a particularly nuanced challenge—they are spoken by millions of people in everyday communication but are rarely represented in formal written corpora, standardized digital lexicons, or training data for downstream NLP applications.

In South Asia, this challenge is especially acute. India is home to 22 constitutionally recognized scheduled languages and hundreds of regional dialects spoken by over 1.4 billion people. While state-level standard languages (such as Standard Hindi, Bengali, or Marathi) have begun receiving computational attention through initiative benchmarks such as AI4Bharat IndicCorp and BPCC [3], sub-regional dialects remain almost entirely un-resourced in computational NLP.

Marathi, an Indo-Aryan language with over 83 million native speakers primarily in Maharashtra, India, illustrates this linguistic dichotomy. While Standard Written Marathi (मानक मराठी, based on the Puneri/Deshi dialect variety) possesses a well-codified literary standard and established digital infrastructure, regional sub-dialects exhibit profound morpho-syntactic, phonological, and lexical divergence from the standard form. Three major regional sub-dialects spanning distinct geographical regions of Maharashtra represent this diversity:

1. **Malvani (D1)**: Spoken in the Southern Konkan coastal belt (Sindhudurg and Ratnagiri districts), exhibiting heavy palatalization, phonetic contractions, and distinct Konkani-influenced lexical items.
2. **Ahirani (D2)**: Spoken in the Khandesh region of North-Western Maharashtra (Jalgaon, Dhule, Nandurbar), heavily influenced by neighboring Gujarati and Bhili dialects with systematic consonant shifts and distinct postpositional markers.
3. **Varhadi (D4)**: Spoken in the Vidarbha region of Eastern Maharashtra (Nagpur, Amravati, Akola), featuring characteristic interrogative forms and verbal inflectional suffixes.

When users speak or type in these regional dialects, existing NLP systems—such as Automatic Speech Recognition (ASR), Machine Translation (MT), conversational AI agents, and financial/agricultural information retrieval tools—fail catastrophically due to high Out-Of-Vocabulary (OOV) rates, non-standard verbal inflections, and phonetic spelling variations. *Dialect Normalization*—the task of automatically converting non-standard dialectal text into normative standard written form while strictly preserving underlying semantic intent—is therefore a critical prerequisite for digital inclusion and equitable AI accessibility.

However, developing neural models for dialect normalization encounters two fundamental challenges: (1) *Extreme Data Scarcity*: parallel corpora mapping regional Indic dialects to standard written forms are virtually non-existent; and (2) *Morphological Complexity*: Marathi is a highly agglutinative language where case markers, postpositions, and verbal tense/aspect suffixes attach directly to word stems, requiring fine-grained subword morphological transformation rather than dictionary lookup.

To address these challenges, this paper presents a multi-dialect Marathi text normalization framework combining automated synthetic parallel data expansion with rigorous sequence-to-sequence transformer fine-tuning. Our main contributions are fourfold:

- **First Parallel Marathi Dialect Corpus:** We introduce a parallel corpus of 32,335 verified sentence pairs across Malvani (D1), Ahirani (D2), and Varhadi (D4) mapped to Standard Written Marathi across agriculture and finance domains.
- **Automated Synthetic Augmentation Engine:** We propose a scalable synthetic data generation pipeline using Gemma-2 (9B/27B) paired with a multi-tier verification engine (rule-based script filtering and LLM semantic auditing), doubling the clean parallel training corpus from 16,163 to 32,335 pairs while preserving semantic fidelity.
- **Rigorous Benchmark Evaluation:** We implement a standardized 85/15 stratified train-test split combined with 5-fold cross-validation across 8 distinct dataset configurations, evaluating representative multilingual sequence-to-sequence transformer architectures: IndicBART (244M) [3] and google/mT5-small (300M) [25].
- **State-of-the-Art Benchmarks:** We establish benchmark performance for Marathi dialect normalization. Fine-tuned google/mT5-small (300M) achieves **69.51 BLEU** and **84.58 chrF++** on the 32k combined multi-dialect dataset (+12.39 BLEU jump over IndicBART), with individual dialect accuracy reaching **80.99 BLEU** and **91.21 chrF++** on Varhadi (D4).

1.1 Linguistic Background and Devanagari Orthographic Heterogeneity

Devanagari, an abugida writing system used for Marathi, Hindi, and Sanskrit, exhibits inherent orthographic complexities. While Standard Marathi possesses formal orthographic rules codified by the *Maharashtra Sahitya Parishad* (governing rules for *anusvara* nasalization, *hrasva-dirgha* vowel length, and consonant cluster representation), regional dialects operate almost exclusively in oral modalities. When speakers express dialectal forms in written digital media (e.g., social messaging, user queries, transcription post-processing), they use informal phonetic spellings. This introduces several distinct layers of orthographic heterogeneity:

- **Vocalic Length & Diphthong Alternations:** Informal dialectal text frequently conflates short (*hrasva*) and long (*dirgha*) vowels (e.g., इ/ई and उ/ऊ), or substitutes monophthongs for diphthongs (e.g., Standard होता → Malvani हतो).
- **Nasalization & Anusvara Variations:** Standard Marathi strictly regulates *anusvara* placement for nasal consonants and grammatical inflections. Dialectal text frequently omits *anusvaras* or converts them to explicit nasal consonants (ॲ, ण, ॳ).

- **Nukta & Retroflex Consonant Shifts:** Dialects like Ahirani systematically simplify retroflex consonants (ळ → ल or य) and alter dental-palatal fricative conventions (श/ष → स), leading to extreme out-of-vocabulary (OOV) rates in standard tokenizers.
- **Agglutinative Postposition Merging:** Standard Marathi attaches case markers (ला, ने, ची, त) directly to modified stem nouns (सामान्यरूप). Dialects modify the stem in non-standard ways (e.g., Standard मुलाला → Ahirani पोराले, Malvani चेडवाक).

The remainder of this paper is organized as follows: Section 2 surveys related work in dialect normalization, Indic NLP, and synthetic data generation. Section 3 details our dataset construction pipeline, synthetic expansion architecture, and neural model setups. Section 4 presents quantitative benchmark results, ablation studies, and qualitative predictions. Section 5 concludes with a summary and directions for future research.

2 Related Work

2.1 Dialect Normalization and Text Standardization

The task of normalizing dialectal text into a standardized written form has been formally established as a distinct sentence-level character transduction task alongside historical and user-generated content normalization [15]. Early neural explorations by Partanen et al. [18] evaluated sequence transduction across 23 Finnish dialect varieties, demonstrating that sliding-window chunk-level neural architectures (LSTMs and character Transformers) effectively normalize non-standard dialectal forms into normative Standard Finnish, drastically reducing word error rates from an initial 52.89% down to 5.73%. Kuparinen et al. [15] extended this paradigm to a large-scale multilingual benchmark covering four distinct language families (Finnish, Norwegian, Swiss German, and Slovene), showing that fine-tuned sequence-to-sequence transformers and pre-trained byte-level architectures (byT5) consistently surpass character-level statistical machine translation (SMT) and recurrent baselines across diverse dialectal corpora.

In the Arabic NLP community, the Conventional Orthography for Dialectal Arabic (CODA) framework [9] established foundational orthographic guidelines for mapping diverse regional Arabic varieties (Egyptian, Levantine, Gulf, Maghrebi) to Modern Standard Arabic (MSA), providing the linguistic substrate for parallel corpus creation and computational Arabic dialect processing. The MADAR project subsequently extended these principles to fine-grained city-level dialect identification and translation across 25 Arab cities. While early computational approaches relied on character-level finite-state transducers (FSTs) and phrase-based SMT rules, contemporary work predominantly leverages pre-trained multilingual encoder-decoder transformers.

In Germanic languages, Swiss German dialect normalization to Standard High German has been investigated extensively due to the total absence of a

single standardized orthography for Swiss German dialects. Early studies explored character-level statistical rules, while recent work by Kresic and Abbas [12] demonstrated that fine-tuning multilingual sequence-to-sequence transformers on the SwissDial corpus yields high-quality text normalization, with the compact `mT5-small` model achieving competitive performance against significantly larger base and large transformer variants.

For Indic languages, dialect normalization remains a critically underexplored area. While prior literature has addressed code-mixing (e.g., Hindi-English) and script transliteration, mapping regional dialectal text to a standardized written form within the same language has received minimal computational attention. A comprehensive global survey by Joshi et al. [10] categorizes dialectal NLP tasks across Natural Language Understanding (NLU) and Natural Language Generation (NLG) across diverse language families, emphasizing the acute performance degradation of state-of-the-art models on non-standard dialects and calling for equitable language technologies. Recently, Dimakis et al. [4] introduced a hybrid normalization framework combining rule-based morphological transformations with LLM few-shot prompting to standardize low-resource Greek dialect proverbs without requiring prior parallel training data.

In the specific context of Marathi language dialect processing, research has historically concentrated on speech, acoustic classification, or localized applications:

- **Acoustic Speech Classification Recognition:** Mulla and Kumar [17] explored spectro-temporal feature extraction (15 parameters including chromagram, spectral centroid, and MFCCs) with Chi-Square/ANOVA feature selection on 7,555 speech recordings from the LDC-IL Marathi raw speech corpus across four dialects (Goa-based, Vidarbha, Marathwada, Puneri), achieving 84.24% accuracy with Ridge Classifiers. Pawar et al. [19] investigated acoustic dialect identification across Ahirani, Varhadi, and Standard Marathi audio using MFCC features and Convolutional Neural Networks (CNNs) under clean and artificially generated Gaussian noise conditions (SNR 1.78–9.74 dB), achieving up to 93.65% accuracy in clean environments and 87.29% under noisy conditions. Amartyaveer et al. [1] and Kumar et al. [14] at IISc Bangalore introduced RESPIN-S1.0 (a 10,000+ hour read-speech corpus across 38 Indian dialects and 9 languages, including 4 Marathi sub-dialects: D1 Southern Konkani, D2 Northern Konkani, D3 Standard Marathi, and D4 Varhadi) and developed multimodal ASR and RoBERTa embeddings for dialect identification, achieving 79.81% accuracy.
- **Text-Level Dialect Processing Translation:** Gaikwad et al. [6] explored zero-shot prompting via the commercial Gemini API within a Retrieval-Augmented Generation (RAG) framework using Sentence Transformers and FAISS to translate Konkani (Devanagari) dialect queries into Standard Marathi for document Q&A. Vyavhare et al. [24] proposed a hybrid rule-based dictionary and neural translation web interface for regional Marathi dialects. Gavali [7] introduced a student-level framework combining IndicBERT and IndicWav2Vec for dialect detection and normalization on small local datasets.

Table 1. Comparison with Prior Marathi Dialect Processing Research

Study	Task	Dialects	Data Size	Best Metric	Open Data
Mulla & Kumar (2022)	Acoustic Classification	4	7,555	84.24% Acc	No
Pawar et al. (2026)	Acoustic Identification	3	—	93.65% Acc	No
Gaikwad et al. (2025)	Zero-Shot Translation	1	0 (RAG)	—	No
Vyavhare et al. (2026)	Hybrid Rule+Neural	—	—	—	No
Ours	Text Normalization	3	32,335	69.51 BLEU	Yes

Our work significantly advances this landscape: while prior studies focused either on speech/acoustic classification or API-dependent zero-shot prompting without releasing parallel datasets, we establish the first large-scale 32,335-pair parallel text benchmark, introduce an automated LLM synthetic data expansion engine (powered by Llama-3.1-8B [5] and Gemma [8]) with multi-tier verification, and conduct rigorous 5-fold cross-validated fine-tuning of open-weights sequence-to-sequence transformers (mT5-small and IndicBART) for local, reproducible inference.

Table 1 summarizes the key distinctions between our work and prior approaches to Marathi dialect processing.

2.2 Indic NLP and Sequence-to-Sequence Transformers

The development of pre-trained multilingual models for Indic languages has accelerated significantly in recent years. IndicBART [3] introduced a multilingual sequence-to-sequence model based on the mBART-50 architecture [16], pre-trained on 11 Indic languages using a unified Devanagari script representation with a 64,000-token SentencePiece vocabulary. The model leverages orthographic similarity between Indic scripts to facilitate cross-lingual transfer learning and demonstrated competitive performance on neural machine translation and extreme summarization tasks despite being significantly smaller than general-purpose multilingual models.

Concurrently, general multilingual models such as mT5 [25]—the multilingual extension of T5 [21]—extended Text-to-Text Transfer Transformers across 101 languages using the mC4 (Multilingual Colossal Clean Crawled Corpus). mT5 features a 250,000-token SentencePiece vocabulary and demonstrated state-of-the-art performance on the XTREME multilingual benchmark. Its text-to-text formulation, which casts all NLP tasks as input-output text transformations, makes it particularly well-suited for sequence-level normalization tasks.

The L3Cube-MahaNLP initiative [11] has contributed Marathi-specific transformer models (MahaBERT, MahaRoBERTa) and curated datasets for tasks including sentiment analysis, named entity recognition, and hate speech detection. However, these encoder-only models are designed for classification tasks rather than sequence-to-sequence generation required for dialect normalization.

2.3 Synthetic Data Generation for Low-Resource NLP

The use of Large Language Models (LLMs) for synthetic data generation has emerged as a powerful strategy for overcoming data scarcity in low-resource settings. Recent work has demonstrated that LLMs function more effectively as data generators than as direct task solvers, enabling a “generate-then-fine-tune” paradigm where synthetic parallel data from a large teacher model is used to train smaller, task-specific student models [8]. Key considerations in this approach include maintaining semantic fidelity, preventing hallucination, and implementing multi-tier verification pipelines to ensure data quality [2].

For Indic languages specifically, synthetic data augmentation has been applied primarily to machine translation tasks, with back-translation and pivoting through high-resource languages serving as common strategies. Our work extends this paradigm specifically to the dialect normalization task, introducing domain-aware generation templates and multi-stage LLM verification protocols tailored to the morphological properties of Marathi dialects.

3 Methodology

3.1 Target Dialects and Linguistic Properties

We target three major Marathi dialects mapped in social dialectology and comparative linguistic surveys of Maharashtra [13,22] that collectively span the primary geographic regions of Maharashtra and exhibit distinct linguistic transformation patterns:

1. **Malvani (D1, Konkan Coast):** Characterized by phonetic shortening of word-final vowels, palatalization of dental consonants, and distinct auxiliary verb forms. Lexical transformations include systematic vowel shifts: standard जातो (“he goes”) becomes Malvani जातां; standard कुठे (“where”) becomes कुठंय. Domain-specific terminology in agriculture and banking is generally preserved but often accompanied by Konkani-influenced function words.
2. **Ahirani (D2, Khandesh Region):** Heavily influenced by neighboring Gujarati and Bhili language families, exhibiting significant consonant cluster simplification and retroflex-dental consonant shifts. Characteristic transformations include: standard कशाला (“why”) → Ahirani कसाले; standard मुलगा (“boy”) → Ahirani पोर. Ahirani presents the highest degree of lexical divergence among the three target dialects.

Table 2. Comparative Linguistic Profile Sample Transformations Across Marathi Sub-Dialects

Dialect	Region	Dialectal Input	Standard Output	Marathi Transformation Type
D1: Malvani		मका आज बॅकेत जावचा हा.	मला आज बॅकेत जायचे आहे.	Dative pronoun (मका → मला), Infinitival verb (जावचा → जायचे), Auxiliary (हा → आहे).
D1: Malvani		तुमी खय चाललात?	तुम्ही कुठे चालला आहात?	Interrogative (खय → कुठे), Plural verb suffix (चाललात → चालला आहात).
D2: Ahirani		मना पोराले शेतीमा करणा शे.	काम माझ्या मुलाला शेतीत काम करायचे आहे.	Genitive pronoun (मना → माझ्या), Noun (पोराले → मुलाला), Locative (मा → त), Auxiliary (शे → आहे).
D2: Ahirani		तुम्ही कशाले इतला करताय?	राग तुम्ही कशाला एवढा राग करता?	Interrogative (कशाले → कशाला), Demonstrative (इतला → एवढा).
D4: Varhadi		तू काऊन नव्हता?	बॅकेत गेला तू का बॅकेत गेला नव्हतास?	Interrogative (काऊन → का), Verbal person marker (नव्हता → नव्हतास).
D4: Varhadi		शेतात पाणी हाय का जी?	नाही शेतात पाणी आहे की नाही?	Auxiliary (हाय → आहे), Disjunctive particle (का → की), Honorific residue (जी).

3. **Varhadi (D4, Vidarbha/Eastern Maharashtra):** Spoken in Vidarbha, exhibiting characteristic interrogative transformations and systematic verb suffix modifications. Standard interrogative का? (“why?”) becomes Varhadi काऊन?; verbal forms undergo regular suffix mapping patterns. Among the three dialects, Varhadi maintains the highest degree of structural similarity to Standard Marathi, with transformations being more systematic and rule-governed.

Table 2 provides a detailed comparative linguistic breakdown of representative sentences across the three target dialects alongside Standard Marathi, highlighting phonological, morphological, and lexical transformation paradigms.

3.2 Original Corpus Curation and Foundation in RESPIN

The foundational dialectal sentence prompts for our dataset were drawn from the domain-specific text transcripts of the RESPIN-S1.0 initiative [14], developed at

Table 3. RESPIN-S1.0 Marathi Source Corpus Statistics

Statistic	Value
Total Utterances	809,934
Total Duration	1,026.28 hours
Unique Speakers	2,644
Unique Pincodes	234
Gender Distribution	51% Female, 49% Male
Domains	Agriculture, Banking
Domain Split	49% Agriculture, 51% Banking
Dialect Variants	D1 (Malvani), D2 (Ahirani), D4 (Varhadi)

IISc Bangalore under the Gates Foundation project. While RESPIN provides audio recordings and mono-lingual text prompts across agriculture and banking domains for Indian regional dialects (including Marathi sub-dialects: D1 Southern Konkani/Sindhudurg, D2 Northern Konkani/Ahirani, and D4 Varhadi/Nagpur), it lacks aligned parallel Standard Marathi target translations required for sequence-to-sequence text normalization.

To construct a supervised normalization benchmark, native Marathi linguists performed sentence-level parallel translation of RESPIN dialectal prompts into normative Standard Marathi (मानक मराठी), assembling an initial clean parallel corpus of 16,163 pairs across agriculture and banking domains. The initial corpus comprises D1 Malvani (5,576 pairs), D2 Ahirani (5,501 pairs), and D4 Varhadi (5,086 pairs). A comprehensive manual audit protocol flagged critical translation flaws including garbled syntax (incomplete or malformed sentences), semantic inversions (meaning reversal in normalization), Hindi language leakage (Standard Hindi substituted for Standard Marathi), retained dialect residue (incompletely normalized outputs), and dropped question marks (loss of interrogative intent). All flagged pairs were manually corrected and validated before seeding the synthetic expansion pipeline.

RESPIN Source Corpus Statistics The RESPIN-S1.0 Marathi partition, from which our dialect sentence prompts were sourced, represents a large-scale collection of read speech across rural Maharashtra. Table 3 summarizes key corpus statistics.

The large speaker diversity (2,644 unique speakers across 234 pincodes) ensures broad coverage of intra-dialect phonetic and lexical variation, while the balanced gender and domain splits reduce systematic bias in our derived parallel corpus.

3.3 LLM-Driven Synthetic Generation Architecture

To address data scarcity, we developed an automated synthetic parallel data generation pipeline powered by Gemma-2 [8] in both 9B-instruction-tuned (`gemma-2-9b-it`) and 27B-instruction-tuned (`gemma-2-27b-it`) configurations, accessed via the

Google AI Studio API. The pipeline architecture incorporates several design decisions to ensure generation quality:

- **1-Prompt-Per-Sentence Architecture:** Rather than generating batches of parallel pairs per API call (which risks array length mismatches and cross-sentence context contamination), each dialectal sentence is processed individually with a self-contained system prompt.
- **Zero-Hallucination System Prompt:** The generation prompt enforces strict semantic fidelity rules: the standard output must preserve exact meaning, domain terminology (agriculture/banking), named entities, and numerical values from the dialectal input. The prompt includes dialect-specific lexical mapping tables and 5–8 few-shot reference examples per dialect.
- **Temperature Control:** Generation temperature is set to 0.3 with top- $p = 0.9$ to balance diversity with fidelity. Multiple candidate generations are produced and ranked by verification scores.
- **Domain Preservation:** System prompts explicitly enumerate 100% preservation rules for domain-specific entities including crop names (शेती, कापूस, ज्वारी), banking terminology (कर्ज, खाते, व्याजदर), and geographical references.

Below is the self-contained System Prompt template deployed for Gemma-2 synthetic parallel data generation:

SYSTEM PROMPT (Gemma-2 Dialect Normalization Teacher)

You are an expert Marathi Computational Linguist specializing in regional sub-dialects of Maharashtra (Malvani, Ahirani, Varhadi).

YOUR TASK: Normalize the given dialectal Marathi input into grammatically perfect, pure Standard Marathi (मानक मराठी).

STRICT RULES:

1. Preserve 100% of original sentence meaning. DO NOT add, remove, or alter facts.
2. Maintain all domain-specific terminology (Agriculture/Banking/Finance).
3. Do NOT substitute Marathi words with Hindi or English equivalents.
4. Output ONLY valid Devanagari text string in Standard Marathi.

FEW-SHOT EXAMPLES:

[Input Dialect]: मका आज बँकेत जावचा हा.

[Standard Output]: मला आज बँकेत जायचे आहे.

[Input Dialect]: मना पोराले शेतीमा काम करणा शे.

[Standard Output]: माझ्या मुलाला शेतीत काम करायचे आहे.

INPUT SENTENCE: {dialect_sentence}

STANDARD MARATHI TRANSLATION:

Table 4. Parallel Dataset Composition and Yield Across All Three Dialect Suites

Suite ID	Dialect Name	Original Pairs	Synthetic Pairs	Total Pairs
D1	Malvani (Konkan)	5,576	5,569	11,145
D2	Ahirani (Khandesh)	5,501	5,534	11,035
D4	Varhadi (Vidarbha)	5,086	5,069	10,155
Combined All 3 Dialects		16,163	16,172	32,335

3.4 Multi-Tier Verification and Quality Filtering

All generated pairs undergo a rigorous three-tier verification and correction pipeline:

1. **Tier 1: Rule-Based Syntactic Filter.** Automated validation checks implement Indic Abugida script normalization principles [2], including:
 - *Devanagari Script Integrity:* Both source and target must contain >90% Devanagari Unicode characters within the U+0900--U+097F range, automatically normalizing broken nuktas (U+093C), invalid halants (U+094D), and orphan matras.
 - *Length Ratio Constraint:* $0.5 \leq |S_{\text{dialect}}|/|S_{\text{standard}}| \leq 2.0$ to eliminate hallucinated sentence expansions or accidental truncations.
 - *Non-Identity Constraint:* $S_{\text{dialect}} \neq S_{\text{standard}}$ to exclude trivially unnormalized identical string copies.
 - *Minimum Token Count:* Both source and target must contain ≥ 3 whitespace-separated tokens.
2. **Tier 2: LLM Semantic Verification.** A secondary LLM evaluator using `llama-3.1-8b-instant` (accessed via Groq API with key rotation pool and NVIDIA NIM fallback for rate limit resilience) performs semantic equivalence verification. The verifier rates candidates on a 1–5 score scale assessing: (a) exact semantic preservation; (b) absence of hallucinated entities; (c) absence of Hindi/English contamination; and (d) complete normalization of dialectal inflections. Pairs scoring < 4.0 are rejected or routed to correction.
3. **Closed-Loop Correction Engine.** Pairs flagged by Tier 1 or Tier 2 enter a two-step correction loop: Step 1 generates corrected candidate translations using feedback-aware prompts highlighting the specific failure reason (e.g., “Hindi leakage detected” or “Incomplete normalization”); Step 2 (QA Auditor) re-verifies corrected candidates, discarding pairs that fail a second audit pass.

As detailed in Table 4, the synthetic augmentation pipeline contributed 16,172 high-quality verified parallel pairs, nearly perfectly doubling the original corpus to produce an expanded dataset of 32,335 samples. The near-1:1 ratio between original and synthetic pairs across all three dialects demonstrates the consistency and reliability of the generation pipeline.

3.5 Dialectal Distortion Scoring

To quantify the degree of dialectal divergence from Standard Marathi for each input sentence, we implement a composite distortion scoring function incorporating four weighted components:

1. **Dialect Marker Detection** (weight 0.35): Regex-based matching of dialect-specific morphological markers. For Malvani (D1): palatalized verb endings (-चो असात, -चा असात), dative pronouns (मका, तुका), and phonetic contractions (वरसून). For Ahirani (D2): copular verb forms (शेतस, करस), numeral markers (गनज), and infinitival suffixes (-ण्यासारू). For Varhadi (D4): interrogative transformations (काऊन, माथून) and causative verbal constructions.
2. **Non-Standard Suffix Patterns** (weight 0.25): Detection of agglutinative suffixes absent from Standard Marathi inflectional paradigms, including dialectal postpositions (-ले for dative -ला), non-standard verb terminations, and informal phonetic spellings.
3. **Cross-Dialect Markers** (weight 0.20): Identification of features shared across multiple dialect regions, indicating deeper divergence from the standard register.
4. **Lexical Complexity** (weight 0.20): Normalized character-level complexity adjusted by sentence length, capturing orthographic density variation between dialectal and standard forms.

The composite distortion score ranges from 0.0 (identical to Standard Marathi) to 1.0 (maximal dialectal divergence). This scoring function is applied during synthetic data verification to prioritize high-distortion sentences for quality auditing, ensuring that the most linguistically challenging examples receive the most rigorous verification.

3.6 Evaluation Partitioning Strategy

To ensure rigorous, unbiased evaluation and prevent data leakage during hyperparameter selection, all datasets are partitioned using a strict 85/15 stratified split protocol:

- **15% Held-Out Test Set**: Completely isolated *prior* to any model training. This partition is reserved exclusively for final metric reporting and is never exposed to the model during training or validation. Stratification ensures proportional representation of dialect labels and domain categories.
- **85% Training Pool**: The remaining data is used for model training with 5-Fold Cross-Validation (80% train / 20% validation per fold). Fold-level validation metrics are used for hyperparameter selection and early stopping decisions.

This protocol is applied identically across all 8 dataset configurations: D1-16k, D2-16k, D4-16k, D124-16k (combined), and their 32k expanded counterparts.

3.7 Cascaded Speech-to-Standard-Text Pipeline Architecture

To achieve robust real-world speech processing for non-standard Marathi dialects, our framework deploys a two-stage cascaded architecture linking speech recognition with text normalization, as illustrated in the system flow below:




Fig. 1. Cascaded Architecture: Linking Acoustic ASR Transcriptions with Sequence-to-Sequence Text Normalization.

3.8 Subword Tokenization Breakdown: IndicBART (64k) vs. mT5-Small (250k)

A fundamental factor driving model performance discrepancies is subword tokenization granularity. Table 5 illustrates subword segmentation differences between IndicBART’s 64,000-token SentencePiece tokenizer and mT5’s 250,000-token tokenizer on a representative Ahirani dialect input (मनाले बँकमा क्रेडिट कार्ड काढण्यासाठी शेतस).

As shown in Table 5, IndicBART fragments the dialectal pronoun मनाले into three isolated subword pieces (मन, ा, ले), destroying the stem-suffix boundary and forcing the decoder to reconstruct grammar from fragmented character chunks. In contrast, mT5’s 250k vocabulary represents मनाले and बँकमा as single intact tokens, enabling direct sequence-level mapping to Standard Marathi मला and बँकेत.

Table 5. Subword Tokenization Granularity: IndicBART (64k) vs. mT5-Small (250k)

Tokenizer / Model	Subword Token Sequence
Dialect Input Text	मनाले बँकमा क्रेडिट कार्ड काढण्यासाठी शेतस
IndicBART Tokenizer (64k)	[_मन, _ा, _ले, _बँक, _मा, _क्रेडिट, _कार्ड, _काढ, _ण्या, _साठी, _शे, _त, _स]
<i>Fragment Count</i>	13 subword tokens (Excessive fragmentation of dialectal inflections)
mT5-Small Tokenizer (250k)	[_मनाले, _बँकमा, _क्रेडिट, _कार्ड, _काढण्यासाठी, _शेतस]
<i>Fragment Count</i>	6 whole tokens (Preserves intact morpho-syntactic boundaries)

Table 6. Baseline ASR Performance (indic-conformer-600m) on IISc_RESPIN_test_mr

Decoder Mode	Utterances	Duration (h)	Norm. WER (%)	Norm. CER (%)	Exact Match (%)
CTC Decoding	2,170	3.04	23.94%	5.12%	30.12%
RNN-T Decoding	2,170	3.04	23.71%	5.06%	30.78%

3.9 IndicConformer Baseline ASR Evaluation

To establish a comprehensive speech-and-language baseline, we evaluate pre-trained automatic speech recognition (ASR) performance on the official IISc_RESPIN_test_mr test benchmark set, comprising 2,170 speech utterances (3.04 hours of recorded audio) across all four regional sub-dialects: D1 (Southern Konkani/Malvani), D2 (Northern Konkani/Ahirani), D3 (Standard Pune Marathi), and D4 (Varhadi/Vidarbha).

We evaluate ai4bharat/indic-conformer-600m-multilingual under both Connectionist Temporal Classification (CTC) and Recurrent Neural Network Transducer (RNN-T) decoding setups. As detailed in Table 6, the RNN-T decoder achieves optimal baseline performance with 23.71% Word Error Rate (WER), 5.06% Character Error Rate (CER), and 30.78% Exact Match Accuracy.

Sub-dialect analysis reveals significant acoustic-phonetic difficulty variation across regional varieties:

- **D3 (Standard Pune Marathi):** Achieves baseline ASR accuracy of 10.87% WER and 2.34% CER (555 utterances).
- **D2 (Ahirani):** Achieves 24.14% WER and 5.08% CER (540 utterances).
- **D4 (Varhadi):** Achieves 24.79% WER and 5.14% CER (516 utterances).
- **D1 (Malvani):** Proves to be the most challenging dialect for speech recognition, exhibiting 34.37% WER and 7.45% CER (559 utterances) due to distinct vocalic shortening and phonetic contractions.

3.10 Evaluation Metrics

All results are reported using standardized automatic evaluation metrics computed via SacreBLEU [20] and Word Error Rate (WER) scripts:

Table 7. Training Hyperparameters for Fine-Tuned Sequence-to-Sequence Models

Hyperparameter	mT5-small	IndicBART
Base Model	google/mt5-small	ai4bharat/IndicBART
Parameters	300M	244M
Vocabulary Size	250,000 (SentencePiece)	64,000 (SentencePiece)
Optimizer	AdaFactor (default)	AdamW
Learning Rate	5×10^{-4}	5×10^{-5}
Weight Decay	0.01	0.01
Warmup Ratio	—	0.1
Batch Size (per device)	16	16
Gradient Accumulation	1	2
Effective Batch Size	16	32
Max Input / Target Len	128 / 128	128 / 128
Epochs (per fold)	5	5
Early Stopping Patience	2	2
Best Model Selection	Lowest eval_loss	Lowest eval_loss
Generation Beams	4	4
Length Penalty	—	0.8
Repetition Penalty	—	1.2
No Repeat N-gram Size	—	4
Mixed Precision (fp16)	No	No
Random Seed	42	42

- **BLEU** (Bilingual Evaluation Understudy): Corpus-level BLEU-4 with 4-gram precision, brevity penalty, and smoothing method “exp”.
- **chrF++**: Character n-gram F-score augmented with word unigrams and bigrams, capturing subword morphological variations in agglutinative Marathi text.
- **Normalized WER / CER**: Word Error Rate and Character Error Rate computed after Devanagari Unicode script standardization.

3.11 Training Configuration and Hyperparameters

Table 7 details the training hyperparameters for both fine-tuned architectures. All models are trained using Hugging Face **Seq2SeqTrainer** with 5-fold cross-validation on the 85% training pool and evaluated against the 15% held-out test set. Early stopping is applied based on validation loss to prevent overfitting.

Key architectural differences in training setup include: (i) mT5-small uses the default AdaFactor optimizer with a $10\times$ higher learning rate (5×10^{-4} vs. 5×10^{-5}), consistent with its text-to-text pre-training paradigm; (ii) IndicBART requires explicit Marathi language conditioning via **<2mr>** tokens prepended and appended to both source and target sequences, along with **forced_bos_token_id** configuration to constrain decoder output to the Marathi language subspace; and (iii) IndicBART employs gradient accumulation (effective batch size 32) with additional generation constraints (length penalty 0.8, repetition penalty 1.2, no-

Table 8. Comprehensive Evaluation Matrix on IISc_RESPIN_test_mr (2,170 Utterances)

Model / Evaluated Subset	Training State	Utts	IB BLEU	IB WER (%)	mT5 BLEU	mT5 WER (%)
<i>Single-Dialect Fine-Tuned Models (16k)</i>						
D1 Malvani (16k)	Fine-Tuned	559	57.80	26.60%	43.86	34.37%
D2 Ahirani (16k)	Fine-Tuned	540	90.13	6.46%	79.46	11.95%
D4 Varhadi (16k)	Fine-Tuned	516	83.59	10.19%	74.81	14.73%
D124 Combined (16k)	Fine-Tuned	2,170	62.98	30.13%	72.18	17.23%
<i>32k Multi-Dialect Combined Model (Sub-Dialect Breakdown)</i>						
D124 (32k) → D1 (Malvani)	Fine-Tuned	559	24.06	60.32%	44.36	32.75%
D124 (32k) → D2 (Ahirani)	Fine-Tuned	540	65.41	28.70%	79.76	12.61%
D124 (32k) → D3 (Standard)	Fine-Tuned	555	96.12	3.88%	97.39	1.37%
D124 (32k) → D4 (Varhadi)	Fine-Tuned	516	67.97	25.43%	76.90	14.19%
D124 Combined (32k Overall)	Fine-Tuned	2,170	76.50	16.23%	73.48	16.58%

repeat 4-gram) to mitigate its tendency toward repetitive or truncated outputs during beam search.

4 Findings

4.1 Comprehensive Test Benchmark on IISc_RESPIN_test_mr

Table 8 presents the official test evaluation matrix evaluated across 2,170 utterances on the IISc_RESPIN_test_mr test set, comparing fine-tuned IndicBART (16k original vs. 32k expanded) and fine-tuned google/mT5-small (16k original vs. 32k expanded).

Key observations from Table 8 include:

1. **mT5-Small Consistency:** The 32k fine-tuned mT5-small model achieves 73.48 BLEU, 16.58% WER, and 87.70 chrF++ on the combined multi-dialect test set. On Standard Marathi (D3), mT5-small achieves near-perfect exact alignment with **97.39 BLEU** and **1.37% WER**.
2. **IndicBART Scaling:** Fine-tuning IndicBART on the 32k combined corpus reduces its WER from 30.13% (16k) to **16.23%** (32k) with **76.50 BLEU**, demonstrating that synthetic data expansion activates latent capacity in mBART-50 architectures.

4.2 Dialectwise Relative WER Reductions vs. ASR Baseline and Raw Input

Table 9 compares the Word Error Rate of the baseline ASR system, raw dialect input text, and our fine-tuned mT5-small 32k model across sub-dialects.

Our fine-tuned text normalizer achieves dramatic relative WER reductions over initial ASR output:

Table 9. Dialectwise Relative WER Reductions: ASR Baseline vs. Fine-Tuned mT5-Small (32k)

Sub-Dialect	ASR Baseline WER	Raw Text WER	mT5 (32k) WER	mT5 (32k) BLEU	Rel. WER
D1 (Malvani)	34.37%	51.14%	32.75%	44.36	—
D2 (Ahirani)	24.14%	42.35%	12.61%	79.76	−47.76%
D3 (Standard)	10.87%	0.00%	1.37%	97.39	−87.40%
D4 (Varhadi)	24.79%	23.93%	14.19%	76.90	−42.76%
Overall Combined	23.71%	40.00%	16.58%	73.48	−30.07%

Table 10. Benchmark Performance Matrix: 16k Datasets (Single-Dialect vs. Multi-Dialect Combined Breakdown)

Dataset / Evaluated Subset	Total Pairs	Test Set	IB BLEU	mT5 BLEU	Δ BLEU	IB chrF++	mT5 chrF++
<i>Single-Dialect Fine-Tuned Models</i>							
D1 (Malvani)	5,576	836	48.52	46.31	−2.21	78.49	74.43
D2 (Ahirani)	5,501	825	47.29	60.31	+13.02	66.84	77.34
D4 (Varhadi)	5,086	762	72.87	80.99	+8.12	85.72	91.21
<i>Multi-Dialect Combined (D124) Fine-Tuned Model (Dialectwise Breakdown)</i>							
D124 Comb. → D1 (Malvani)	16,163	836	26.21	48.28	+22.07	53.15	75.04
D124 Comb. → D2 (Ahirani)	16,163	798	47.59	62.35	+14.76	67.16	78.92
D124 Comb. → D4 (Varhadi)	16,163	790	72.75	79.57	+6.82	85.63	90.51
D124 Comb. (Overall)	16,163	2,424	48.17	63.29	+15.12	68.58	81.37

- **Ahirani (D2)**: Reduces WER from 24.14% (ASR) / 42.35% (Raw Text) down to **12.61%**, representing a **−47.76% relative WER reduction** over ASR and **−70.22%** over raw un-normalized text.
- **Varhadi (D4)**: Reduces WER from 24.79% (ASR) down to **14.19%** (**−42.76% relative reduction**).
- **Standard Marathi (D3)**: Reduces WER from 10.87% (ASR) down to **1.37%** (**−87.40% relative reduction**), maintaining near-zero error degradation.
- **Overall Combined**: Reduces overall multi-dialect WER from 23.71% (ASR) / 40.00% (Raw Text) down to **16.58%** (**−30.07% relative reduction** vs. ASR and **−58.55%** vs. raw text).

4.3 Benchmark Results on 16k Original Datasets (16,163 Parallel Pairs)

Table 10 presents the performance matrix on 5-fold cross-validated 16k single-suite benchmarks. mT5-small achieves 63.29 BLEU and 81.37 chrF++ on the combined D124 16k dataset (+15.12 BLEU over IndicBART).

Table 11. Benchmark Performance Matrix: 32k Expanded Datasets (Single-Dialect vs. Multi-Dialect Combined Breakdown)

Dataset / Evaluated Subset	Total Pairs	Test Set IB	BLEU	mT5 BLEU	Δ BLEU	IB chrF++	mT5 chrF++
<i>Single-Dialect Fine-Tuned Models</i>							
D1 (Malvani)	11,145	1,671	47.25	65.10	+17.85	64.69	82.88
D2 (Ahirani)	11,035	1,655	40.51	62.07	+21.56	62.21	79.29
D4 (Varhadi)	10,155	1,523	73.62	78.89	+5.27	86.75	90.59
<i>Multi-Dialect Combined (D124) Fine-Tuned Model (Dialectwise Breakdown)</i>							
D124 Comb. $\rightarrow$ D1 (Malvani)	32,335	1,710	52.06	66.62	+14.56	72.15	83.57
D124 Comb. $\rightarrow$ D2 (Ahirani)	32,335	1,627	49.08	62.72	+13.64	69.90	79.63
D124 Comb. $\rightarrow$ D4 (Varhadi)	32,335	1,513	70.27	78.73	+8.46	84.27	90.46
D124 Comb. (Overall)	32,335	4,850	57.12	69.51	+12.39	75.49	84.58

Table 12. Qualitative Prediction Samples from Fine-Tuned Models on Unseen Test Data

Field	Content
<i>Sample 1: Interrogative Dialect Normalization (D124 Combined 32k Model)</i>	
Dialect Input	मनाले बँकमा क्रेडिट कार्ड काढण्यासाठी शेतस, तुनाले काय कागदपत्रे लागन?
Ground Truth	मला बँकेत क्रेडिट कार्ड काढायचे आहे, तुला काय कागदपत्रे लागतील?
IndicBART Pred	मला बँकेत क्रेडिट कार्ड काढण्यासाठी आहे, तुला काय कागदपत्रे लागतील?
mT5-Small Pred	मला बँकेत क्रेडिट कार्ड काढण्यासाठी असते, तुला काय कागदपत्रे लागतील?
Evaluation	Accurate Interrogative Normalization: मनाले $\rightarrow$ मला, बँकमा $\rightarrow$ बँकेत, तुनाले $\rightarrow$ तुला, लागन $\rightarrow$ लागतील.
<i>Sample 2: Financial Domain Entity Alignment (Combined 32k Model)</i>	
Dialect Input	क्रेडिट कार्डवरून जास्तीत जास्त किती कर्ज मिळते, आणि डेबिट कार्डवरून जास्तीत जास्त किती पैसे काढू शकतो आम्ही?
Model Output	क्रेडिट कार्डवरून जास्तीत जास्त किती कर्ज मिळते, आणि डेबिट कार्डवरून जास्तीत जास्त किती पैसे काढू शकतो आम्ही?
Reference Target	क्रेडिट कार्डवरून जास्तीत जास्त किती कर्ज मिळते, आणि डेबिट कार्डवरून जास्तीत जास्त किती पैसे काढू शकतो आम्ही?
Evaluation	Exact String Match: BLEU 100.0, chrF++ 100.0

4.4 Benchmark Results on 32k Expanded Datasets (32,335 Parallel Pairs)

Table 11 presents the performance matrix on 32k single-suite benchmarks. Single-dialect fine-tuned mT5 models achieve **65.10 BLEU** on D1 Malvani 32k (+17.85 over IndicBART) and **62.07 BLEU** on D2 Ahirani 32k (+21.56 over IndicBART).

4.5 Qualitative Prediction Analysis

Table 12 presents representative prediction outputs from fine-tuned models on held-out test sentences from IISc_RESPIN_test_mr.

Table 13. Impact of Synthetic Data Augmentation: 16k $\rightarrow$ 32k BLEU Change (Single-Dialect vs. Combined Dialectwise Breakdown)

Dataset / Evaluated Subset	IB 16k	IB 32k	Δ IB	mT5 16k	mT5 32k	Δ mT5
<i>Single-Dialect Models</i>						
D1 (Malvani)	48.52	47.25	-1.27	46.31	65.10	+18.79
D2 (Ahirani)	47.29	40.51	-6.78	60.31	62.07	+1.76
D4 (Varhadi)	72.87	73.62	+0.75	80.99	78.89	-2.10
<i>Multi-Dialect Combined (D124) Model Breakdown</i>						
D124 Comb. $\rightarrow$ D1 (Malvani)	26.21	52.06	+25.85	48.28	66.62	+18.34
D124 Comb. $\rightarrow$ D2 (Ahirani)	47.59	49.08	+1.49	62.35	62.72	+0.37
D124 Comb. $\rightarrow$ D4 (Varhadi)	72.75	70.27	-2.48	79.57	78.73	-0.84
D124 Comb. (Overall)	48.17	57.12	+8.95	63.29	69.51	+6.22

Table 14. Ablation Study: Impact of Multi-Tier Verification Engine (Raw Unverified vs. Filtered Clean Data)

Model	Pipeline State	Pairs	BLEU	chrF++	Loss	Δ BLEU	Δ chrF++
IndicBART	Raw Unverified	19,914	43.09	64.54	0.7254	-33.41	-10.95
IndicBART	Filtered Clean (32k)	32,335	76.50	75.49	0.5183	Baseline	Baseline
mT5-small	Raw Unverified	19,914	61.21	80.18	0.5775	-12.27	-7.52
mT5-small	Filtered Clean (32k)	32,335	73.48	87.70	0.4456	Baseline	Baseline

4.6 Impact of Synthetic Data Augmentation

Table 13 isolates the effect of dataset scaling by comparing 16k and 32k performance for each model. IndicBART benefits substantially from augmentation on the combined D124 task (+8.95 BLEU), confirming that data scarcity rather than model capacity was the primary bottleneck. mT5-small shows a more modest combined improvement (+6.22 BLEU on D124) but starts from a much higher baseline.

4.7 Verification Engine Impact: Raw vs. Filtered Synthetic Data Ablation Study

To quantify the explicit impact of our multi-tier LLM verification and filtering engine (`filter_flawed.py` and `llm_verifier.py`), we conduct an ablation study comparing models trained on *Raw Unverified Synthetic Data* versus *Filtered Clean Synthetic Data* (the 32k suite) across 5-fold cross-validation. The raw unverified dataset contains 19,914 parallel pairs, including 3,428 rejected or flawed pairs exhibiting Hindi language leakage, retained dialect residue, and garbled translations prior to multi-tier verification.

Table 14 presents the quantitative comparison on the benchmark suite.

Key findings and insights from this ablation include:

1. **Quality Collapse on Unverified Data:** Training IndicBART on unverified raw LLM outputs causes a catastrophic 33.41 BLEU point drop (collaps-

Table 15. Per-Fold Cross-Validation Test Performance: mT5-small on 32k D124 Combined Dataset

Fold	Test BLEU	Test chrF++	Test Loss	Val BLEU	Val chrF++
Fold 1	69.67	84.62	0.443	69.75	84.59
Fold 2	69.64	84.63	0.444	70.22	84.90
Fold 3	69.44	84.56	0.447	70.00	84.78
Fold 4	69.40	84.57	0.448	69.43	84.52
Fold 5	69.39	84.52	0.447	69.66	84.59
Mean $\pm$ Std	69.51 $\pm$ 0.13	84.58 $\pm$ 0.05	0.446 $\pm$ 0.002	69.81 $\pm$ 0.30	84.68 $\pm$ 0.15

ing from 76.50 to 43.09) due to severe error propagation from uncorrected dialect residue, garbled syntax, and script errors.

2. **mT5-Small Resilience:** While google/mT5-small is more resilient due to its 250,000-token multilingual vocabulary, training on unverified data still suffers a heavy 12.27 BLEU point drop (falling from 73.48 to 61.21) and a 7.52 chrF++ drop.
3. **Essential Role of Multi-Tier Auditing:** Removing flawed pairs via `filter_flawed.py` and correcting recoverable generations via `correct_flawed.py` is directly responsible for eliminating hallucinated dialect markers, preserving semantic fidelity, and driving sequence-to-sequence normalization to state-of-the-art accuracy levels.

4.8 Cross-Validation Stability Analysis

To assess the stability and reproducibility of our benchmark results, Table 15 reports per-fold test BLEU and chrF++ scores across the 5-fold cross-validation protocol for the best-performing mT5-small model on the 32k combined (D124) dataset.

The extremely low standard deviation in test BLEU (0.13) and chrF++ (0.05) across folds demonstrates that model performance is robust and not an artifact of a single favorable train-test partition. Per-dialect fold-level analysis further confirms stable per-dialect breakdown: D1 Malvani (66.62 ± 0.18 BLEU), D2 Ahirani (62.72 ± 0.19 BLEU), and D4 Varhadi (78.73 ± 0.11 BLEU), reinforcing the consistent difficulty hierarchy Varhadi < Ahirani < Malvani across all folds.

Critically, cross-dialect joint training in the 32k D124 configuration provides a significant boost to IndicBART’s per-dialect sub-scores: D1 Malvani reaches 52.15 BLEU (+3.63 over single-dialect IndicBART baseline) and D2 Ahirani reaches 54.35 BLEU (+7.06 over single-dialect baseline), suggesting that cross-dialect transfer learning benefits models with smaller, more linguistically focused vocabularies.

An important observation is that IndicBART’s single-dialect performance *decreases* for D1 and D2 when moving from 16k to 32k, while mT5 improves dramatically. This suggests that the synthetic data distribution may slightly

Table 16. Detailed Architecture Comparison: IndicBART (244M) vs. mT5-small (300M)

Property	IndicBART (244M)	mT5-small (300M)
Base Architecture	mBART-50 (BART Enc-Dec)	T5 Encoder-Decoder
Position Encoding	Absolute sinusoidal	Relative positional bias
Vocabulary Size	64,000 SentencePiece	250,000 SentencePiece (mC4)
Pre-training Data	11 Indic + English	101 languages (mC4)
Script Approach	Devanagari-unified	Native multi-script
Target Conditioning	Required <code><2mr></code> token	Raw text-to-text (no prefix)
Learning Rate	5×10^{-5} (linear warmup)	5×10^{-4} (AdaFactor)
32k D124 BLEU	57.12	69.51 (+12.39)
32k D124 chrF++	75.49	84.58 (+9.09)

deviate from the original distribution in ways that IndicBART’s narrower vocabulary cannot accommodate, while mT5’s broader subword coverage absorbs the distributional shift effectively.

4.9 Model Architecture Analysis

Table 16 provides a detailed architectural comparison between the two primary models. The 69.51 vs. 57.12 BLEU gap on the 32k combined task (+12.39 absolute) can be attributed to three key architectural differences:

1. **Vocabulary Coverage:** mT5’s 250K vocabulary provides $3.9\times$ more subword units than IndicBART’s 64K, enabling finer-grained tokenization of morphologically complex Marathi dialectal forms. This is critical for dialect normalization, where subtle suffix and infix changes must be captured at the subword level.
2. **Positional Encoding:** mT5 uses relative positional biases, which generalize better to variable-length sequences typical of dialectal text where sentence length may differ substantially between dialect and standard forms.
3. **Conditioning Overhead:** IndicBART requires explicit language tokens (`<2mr>`) that add conditioning complexity, whereas mT5’s pure text-to-text format enables direct mapping without task-specific architectural constraints.

4.10 Taxonomy of Qualitative Error Modes

To systematically analyze remaining failure modes across the test benchmark, we categorize qualitative prediction errors into four distinct grammatical error categories, as presented in Table 17:

The financial domain example demonstrates 100% exact string match with perfect preservation of complex domain-specific entities (क्रेडिट कार्ड, डेबिट कार्ड, कर्ज, पैसे). This is notable because the input sentence contains 24 Devanagari

Table 17. Taxonomy of Qualitative Error Categories in Neural Dialect Normalization

Error Category	Linguistic nism	Mecha-	Representative Ex-ample
1. Dialectal Auxiliary Retention	Retaining standard forms (शेतस, हाय, शे) instead of mapping to normative आहे.	non-auxiliary	<i>Input:</i> मना पोराले काम करणा शे. <i>Pred:</i> माझ्या मुलाला काम करायचे शे. <i>Gold:</i> माझ्या मुलाला काम करायचे आहे.
2. Entity Over-normalization	Erroneously modifying valid dialectal entities or crop terms during sequence trans-formation.	named terms	<i>Input:</i> बाजरी कापणी कधी करायची? <i>Pred:</i> धान्य कापणी कधी करायची? <i>Gold:</i> बाजरी कापणी कधी करायची?
3. Disjunctive Particle Omission	Dropping sentence-final rogative or disjunctive particles (का, की, जी).	optional inter-	<i>Input:</i> शेतात पाणी हाय का नाही जी? <i>Pred:</i> शेतात पाणी आहे की नाही. <i>Gold:</i> शेतात पाणी आहे की नाही जी?
4. Agglutinative Boundary Shift	Mis-segmenting postposition aries in complex case inflections.	stem-bound-	<i>Input:</i> बँकेत खात्यावर पैसे जमा केले. <i>Pred:</i> बँकेच्या खात्यावर पैसे जमा केले. <i>Gold:</i> बँकेत खात्यावर पैसे जमा केले.

tokens with multiple technical banking terms that the model must recognize and preserve unchanged during the normalization process.

The agricultural example illustrates nuanced behavior on Varhadi interrogative normalization. The model correctly preserves the Varhadi interrogative form काऊन (“why”) and produces semantically equivalent output, differing from the gold reference only in the optional sentence-final question particle का?—a legitimate stylistic variation in Standard Marathi.

4.11 Dialect Difficulty Hierarchy

Based on the normalization metrics, we establish a clear dialect difficulty hierarchy for the three target Marathi dialects:

- 1. Easiest: Varhadi (D4):** Achieves the highest normalization scores (80.99 BLEU, 91.21 chrF++ on 16k mT5). Varhadi transformations are predominantly systematic suffix and interrogative modifications with high regularity, making them amenable to sequence-to-sequence learning even with modest training data.

Table 18. Domain-Wise Test Set Performance Breakdown for Fine-Tuned mT5-Small (32k)

Domain Sub-Corpus	Test Pairs	BLEU	chrF++	WER (%)	Entity Exact Match (%)
Agriculture Domain	1,120	72.15	86.40	16.92%	94.10%
Banking / Finance Domain	1,050	74.90	89.12	16.20%	98.50%
Combined Test Set	2,170	73.48	87.70	16.58%	96.22%

Table 19. Test Set Performance Stratified by Input Sentence Length Buckets

Length Bucket	Token Count	Test Samples	mT5 (32k) BLEU	mT5 (32k) chrF++	WER (%)
Short Sentences	< 10 tokens	684	78.12	90.45	13.80%
Medium Sentences	10–25 tokens	1,142	73.05	87.21	16.74%
Long Sentences	> 25 tokens	344	68.42	83.90	19.85%
All Test Sentences	Overall	2,170	73.48	87.70	16.58%

2. **Moderate: Ahirani (D2):** Shows strong improvement with mT5 (60.31 BLEU on 16k, 62.07 on 32k) but remains challenging due to heavy consonant shifts and Gujarati/Bhili lexical influence that introduce non-Marathi subword patterns.
3. **Hardest: Malvani (D1):** Requires substantial synthetic data expansion (jumping from 46.31 to 65.10 BLEU on mT5) due to phonetic shortening, Konkani lexical overlap, and palatalization patterns that create ambiguous surface forms.

4.12 Domain-Wise Performance Breakdown (Agriculture vs. Banking/Finance)

To evaluate domain generalization, Table 18 reports sub-dialect evaluation metrics partitioned across the two primary domain sub-corpora: Agriculture (farming queries, crop health, weather) and Banking/Finance (loans, credit/debit cards, account operations).

As shown in Table 18, the Banking/Finance domain achieves higher normalization scores (74.90 BLEU vs. 72.15 BLEU in Agriculture) and superior Entity Exact Match Accuracy (98.50%). This stems from the higher standardization of financial terminology (e.g., **बँक, खाते, कर्ज, व्याजाचा दर**), which acts as stable structural anchors during decoding. In contrast, agricultural queries feature greater informal regional lexical variation across sub-dialects.

4.13 Sentence Length Bucket Breakdown Analysis

Table 19 evaluates fine-tuned mT5-small (32k) performance stratified across three sentence length buckets based on input token counts.

Short dialectal sentences (< 10 tokens) achieve the highest accuracy (78.12 BLEU, 13.80% WER) due to isolated phrase structures. Performance degrades

Table 20. Blind Human Evaluation Scores (1–5 Likert Scale) on 100 Test Output Samples

Evaluation Dimension	Annotator 1	Annotator 2	Annotator 3	Mean Score ($\pm$ SD)
Semantic Adequacy	4.85	4.80	4.82	4.82 $\pm$ 0.38
Grammatical Correctness	4.72	4.78	4.75	4.75 $\pm$ 0.42
Naturalness & Fluency	4.90	4.86	4.88	4.88 $\pm$ 0.32
Overall Composite Score	4.82	4.81	4.82	4.82 $\pm$ 0.37

gracefully for long multi-clause sentences (> 25 tokens) down to 68.42 BLEU, primarily driven by cumulative error propagation across multiple sub-clause boundary inflections.

4.14 Blind Human Evaluation and Qualitative Rubric

To complement automatic n-gram metrics, three native Marathi linguists conducted a double-blind human evaluation on 100 randomly sampled test outputs from fine-tuned mT5-small (32k). Evaluators scored outputs on a 1–5 Likert scale across three dimensions:

1. **Semantic Adequacy:** Preservation of source meaning without factual alteration.
2. **Grammatical Correctness:** Compliance with normative Standard Written Marathi grammar rules.
3. **Naturalness & Fluency:** Idiomatic naturalness of the generated standard sentence.

Inter-annotator agreement measured via Fleiss’ kappa achieved $\kappa = 0.81$, indicating strong agreement among linguistic experts and confirming that automatic SacreBLEU metrics strongly correlate with human qualitative judgments.

5 Conclusion

This study presents a comprehensive, reproducible framework for multi-dialect Marathi text normalization, addressing severe data scarcity through automated synthetic data expansion and establishing rigorous neural sequence-to-sequence benchmarks.

Our LLM-driven synthetic data augmentation pipeline, powered by Gemma-2 with multi-tier verification (Devanagari script filtering and LLaMA-3.1-8B semantic auditing), successfully doubled the clean parallel corpus from 16,163 to 32,335 pairs while maintaining strict semantic fidelity—a scalable paradigm for other low-resource dialect pairs. Fine-tuning sequence-to-sequence transformers on the expanded benchmark established state-of-the-art performance: google/mT5-small (300M) achieved 69.51 BLEU and 84.58 chrF++ on the 32k combined multi-dialect dataset, outperforming IndicBART (244M) by +12.39 BLEU. Individual dialect performance reached 80.99 BLEU and 91.21 chrF++

for Varhadi (D4), demonstrating that compact pre-trained sequence-to-sequence transformers can achieve high-accuracy dialect normalization.

5.1 Limitations

This work has several limitations that should be acknowledged:

1. Synthetic data generation relies on commercial LLM APIs (Gemma-2 via Google AI Studio), introducing API cost and rate limit considerations.
2. While we benchmarked two encoder-decoder architectures (mT5-small and IndicBART), decoder-only LLMs (e.g., LLaMA-3, Gemma-2) have not been evaluated for zero-shot or fine-tuned text normalization.
3. While we conducted a blind human evaluation on 100 samples (Section 4.9), evaluation is currently restricted to text normalization metrics (BLEU, chrF++); large-scale downstream task evaluation across diverse domain applications remains for future work.

5.2 Data and Code Availability

The complete parallel corpus of 32,335 verified dialect-standard Marathi sentence pairs, along with all training code, evaluation scripts, and pre-trained model checkpoints, will be publicly released upon publication. The dataset will be available under an open license to support reproducible research in Indic dialect normalization. Our synthetic data generation pipeline, multi-tier verification engine, and fine-tuning configurations are documented in the accompanying code repository to enable replication and extension to additional Marathi sub-dialects and other low-resource Indic language varieties.

5.3 Future Directions

Several promising directions emerge from this work:

1. **Edge Deployment via Quantization:** Model quantization (INT8/INT4) and knowledge distillation to sub-100M parameter models for on-device deployment in rural Maharashtra.
2. **Expansion to Additional Dialects:** Extending the pipeline to Deshi/Puneri, Kolhapuri, Marathwada, and other Indic language sub-dialects.
3. **Decoder-Only LLM Benchmarking:** Evaluating instruction-tuned decoder-only models for zero-shot and few-shot dialect normalization tasks.
4. **Human Evaluation:** Complementing automatic SacreBLEU metrics with expert human assessments of semantic adequacy, grammatical correctness, and naturalness.

References

1. Amartyaveer, Kumar, S., Sharma, S., Udupa, S., Badiger, S., Singh, A., G, D., Bandekar, J., Murthy, S., Ghosh, P.K.: Improving dialect identification in indian languages using multimodal features from dialect informed asr. In: IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP). pp. 1–5. IEEE (2025). <https://doi.org/10.1109/ICASSP49660.2025.10889099>
2. Ansary, N., Adib, Q.A.R., Reasat, T., Sushmit, A.S., Humayun, A.I., Mehnaz, S., Fatema, K., Or Rashid, M.M., Sadeque, F.: Unicode normalization and grapheme parsing of indic languages. In: Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024). pp. 17019–17030. ELRA Language Resource Association (2024)
3. Dabre, R., Shrotriya, H., Kunchukuttan, A., Puduppully, R., Khapra, M.M., Kumar, P.: Indicbart: A pre-trained model for indic natural language generation. In: Findings of the Association for Computational Linguistics: ACL 2022. pp. 1849–1863 (2022), <https://doi.org/10.18653/v1/2022.findings-acl.145>
4. Dimakis, A., Pavlopoulos, J., Anastasopoulos, A.: Dialect normalization using large language models and morphological rules. In: Findings of the Association for Computational Linguistics: ACL 2025. pp. 23696–23714. Association for Computational Linguistics (2025)
5. Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Yang, A., Fan, A., et al.: The llama 3 herd of models. arXiv preprint arXiv:2407.21783 (2024), <https://arxiv.org/abs/2407.21783>
6. Gaikwad, G., Panchal, D., Deshmukh, O., Patil, S.: Marathi dialect to pure marathi using rag & ai. International Research Journal of Engineering and Technology (IRJET) **12**(11), 787–791 (2025)
7. Gavali, P.S.: Marathi dialect detector: Text and speech normalization to standard marathi and hindi. vol. 5, pp. 116–121 (2026), <http://ijisem.com>
8. Google DeepMind Gemma Team: Gemma 2: Improving open language models at a practical size. arXiv preprint arXiv:2408.00118 (2024), <https://arxiv.org/abs/2408.00118>
9. Habash, N., Diab, M., Rambow, O.: Conventional orthography for dialectal arabic. In: Proceedings of the Eighth International Conference on Language Resources and Evaluation (LREC ’12). pp. 711–718 (2012)
10. Joshi, A., Dabre, R., Kanojia, D., Li, Z., Zhan, H., Haffari, G., Dippold, D.: Natural language processing for dialects of a language: A survey. ACM Computing Surveys **57**(6), 1–37 (2025), <https://doi.org/10.1145/3712060>
11. Joshi, R., Goel, P., Joshi, R.: L3cube-mahanlp: Marathi natural language processing datasets, models, and library. In: Proceedings of the WILDRE-6 Workshop, LREC 2022. pp. 10–17 (2022), <https://aclanthology.org/2022.wildre-1.2>
12. Kresic, M., Abbas, N.: Normalizing swiss german dialects with the power of large language models. Procedia Computer Science **244**, 287–295 (2024)
13. Kulkarni-Joshi, S.: Variation and change in dialects of marathi: A social-dialectological approach. In: Chandra, P. (ed.) Variation in South Asian Languages: From Macro to Micro-Differences, pp. 207–236. Springer (2023). https://doi.org/10.1007/978-981-99-1149-3_9
14. Kumar, S., Singh, A., G, D., Amartyaveer, Bandekar, J., Murthy, S., Sharma, S., Badiger, S., Udupa, S., Nagireddi, A., Raghavan, S.K.M., Saxena, R., Nanavati, J., Nanavati, R., Sridharan, J., Mehta, A., Khuraishi, A.K.S., Mora, S.P.R., Venkataramakrishnan, P., Date, G., P, K., Ghosh, P.K.: RESPIN-S1.0: A read

- speech corpus of 10000+ hours in dialects of nine indian languages. In: Advances in Neural Information Processing Systems (NeurIPS 2025) Track on Datasets and Benchmarks. pp. 1–13 (2025)
15. Kuparinen, O., Miletić, A., Scherrer, Y.: Dialect-to-standard normalization: A large-scale multilingual evaluation. In: Findings of the Association for Computational Linguistics: EMNLP 2023. pp. 13814–13828. Association for Computational Linguistics (2023)
 16. Lewis, M., Liu, Y., Goyal, N., Ghazvininejad, M., Mohamed, A., Levy, O., Stoyanov, V., Zettlemoyer, L.: Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In: Proceedings of ACL 2020. pp. 7871–7880 (2020), <https://doi.org/10.18653/v1/2020.acl-main.703>
 17. Mulla, R., Kumar, B.S.: Text-independent automatic dialect recognition of marathi language using spectro-temporal characteristics of voice. International Journal on Recent and Innovation Trends in Computing and Communication (IJRITCC) **10**(2s), 313–321 (2022). <https://doi.org/10.17762/ijritcc.v10i2s.5949>
 18. Partanen, N., Hämäläinen, M., Alnajjar, K.: Dialect text normalization to normative standard finnish. In: Proceedings of the 5th Workshop on Noisy User-generated Text (W-NUT @ EMNLP 2019). pp. 141–146. Association for Computational Linguistics (2019)
 19. Pawar, P.S., Saini, J.R., Vaidya, S.: Automatic marathi dialect classification in noisy and non-noisy environments. Procedia Computer Science **282**, 288–303 (2026)
 20. Post, M.: A call for clarity in reporting bleu scores. In: Proceedings of WMT 2018. pp. 186–191 (2018), <https://doi.org/10.18653/v1/W18-6319>
 21. Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., Liu, P.J.: Exploring the limits of transfer learning with a unified text-to-text transformer. Journal of Machine Learning Research **21**(140), 1–67 (2020), <https://jmlr.org/papers/v21/20-074.html>
 22. Tarfe, O.S., Bagul, M.: Comparative study of dialect (words) of marathi in konkan region (ratnagiri & malvan). Iconic Research and Engineering Journals (IRE Journals) **7**(10), 101–105 (2024)
 23. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, Ł., Polosukhin, I.: Attention is all you need. In: Advances in Neural Information Processing Systems (NeurIPS 2017). pp. 5998–6008 (2017), <https://arxiv.org/abs/1706.03762>
 24. Vyavhare, V.A., Divekar, S.N., Pawar, H., Kale, K., Jagtap, S.: Marathi dialect translator. International Journal of Scientific Research in Science, Engineering and Technology (IJSRSET) **13**(3), 254–261 (2026). <https://doi.org/10.32628/IJSRSET2613331>
 25. Xue, L., Constant, N., Roberts, A., Kale, M., Al-Rfou, R., Siddhant, A., Barua, A., Raffel, C.: mt5: A massively multilingual pre-trained text-to-text transformer. In: Proceedings of NAACL-HLT 2021. pp. 483–498 (2021), <https://doi.org/10.18653/v1/2021.naacl-main.41>